

Computer Organization

計算機組織

Pipelining

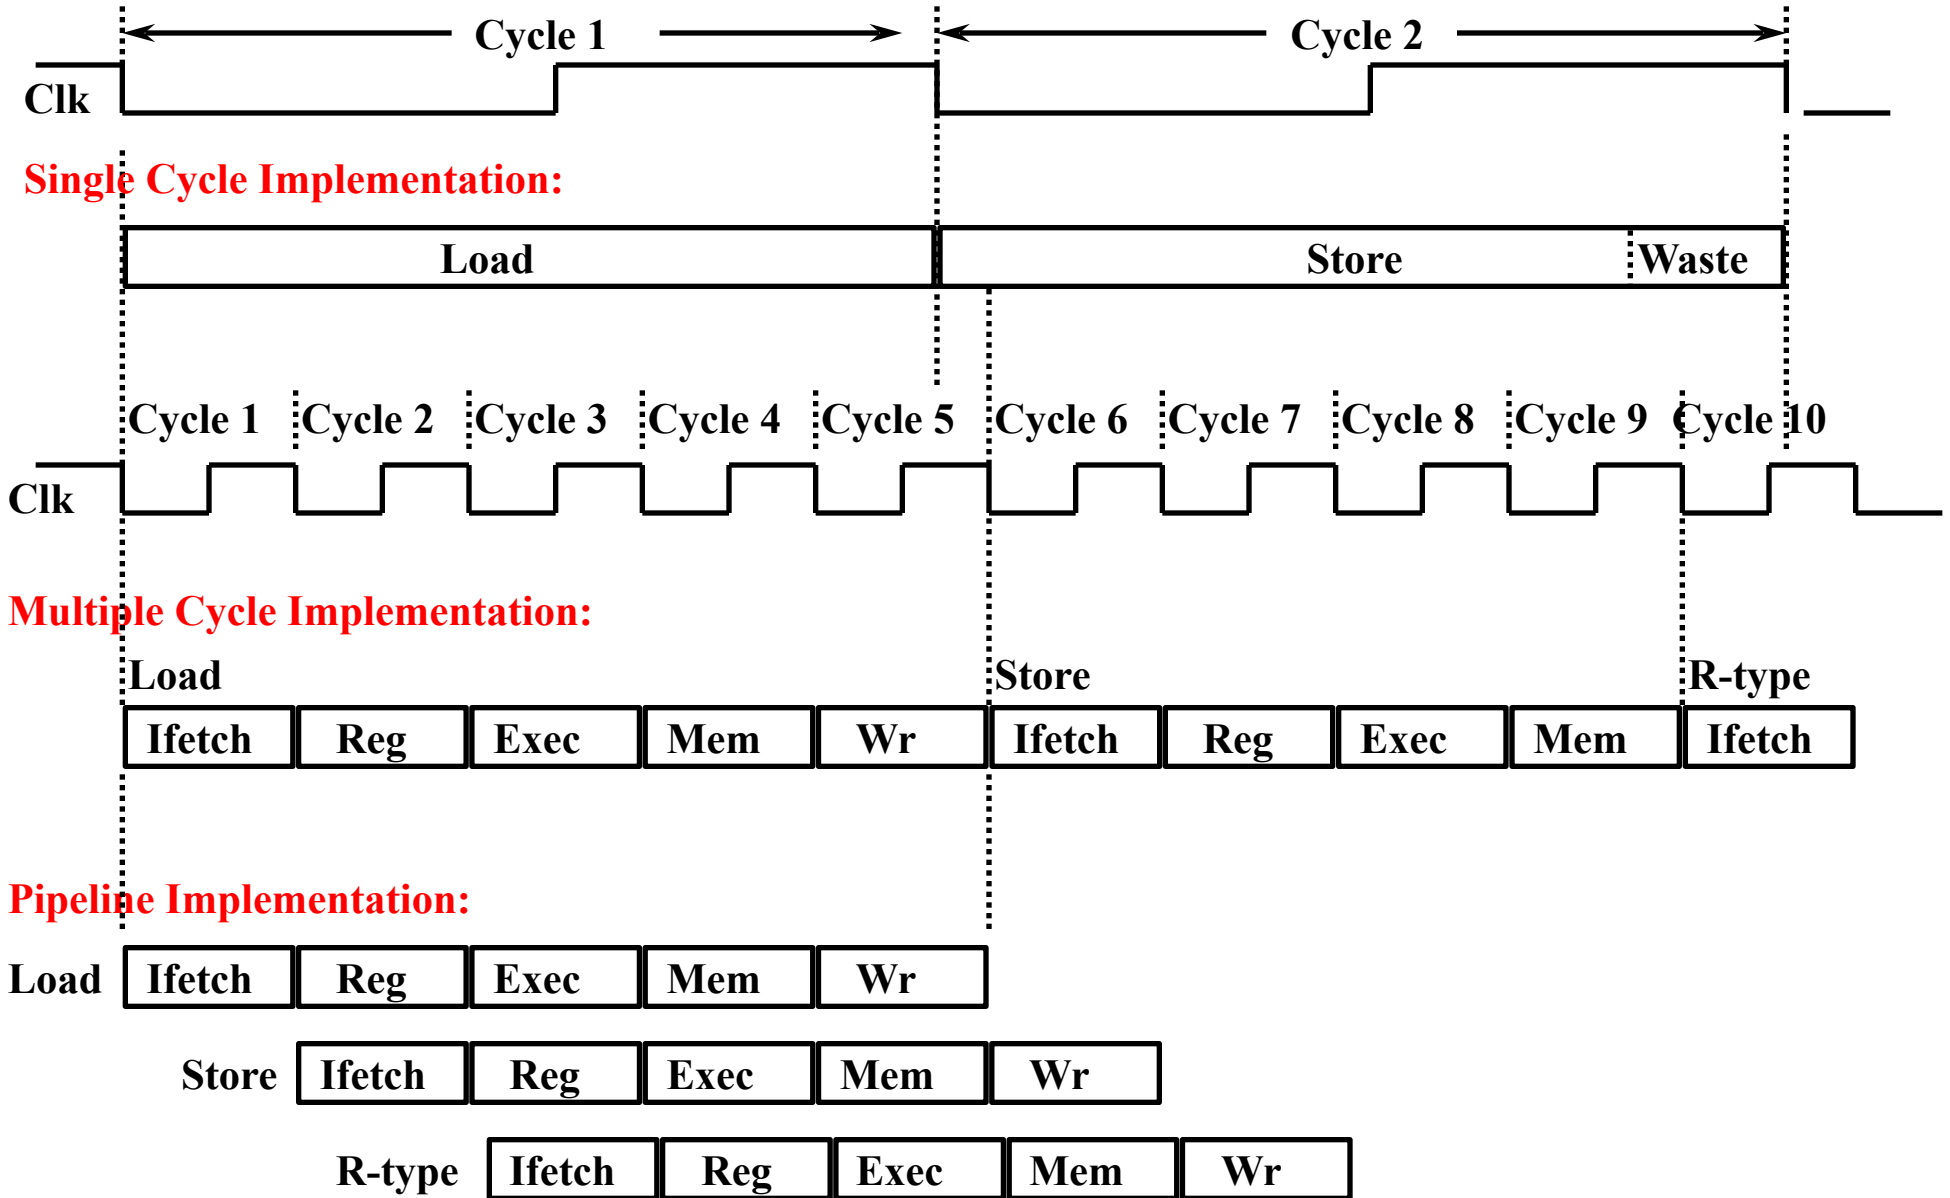


國立成功大學資訊工程學系
105年度第二學期

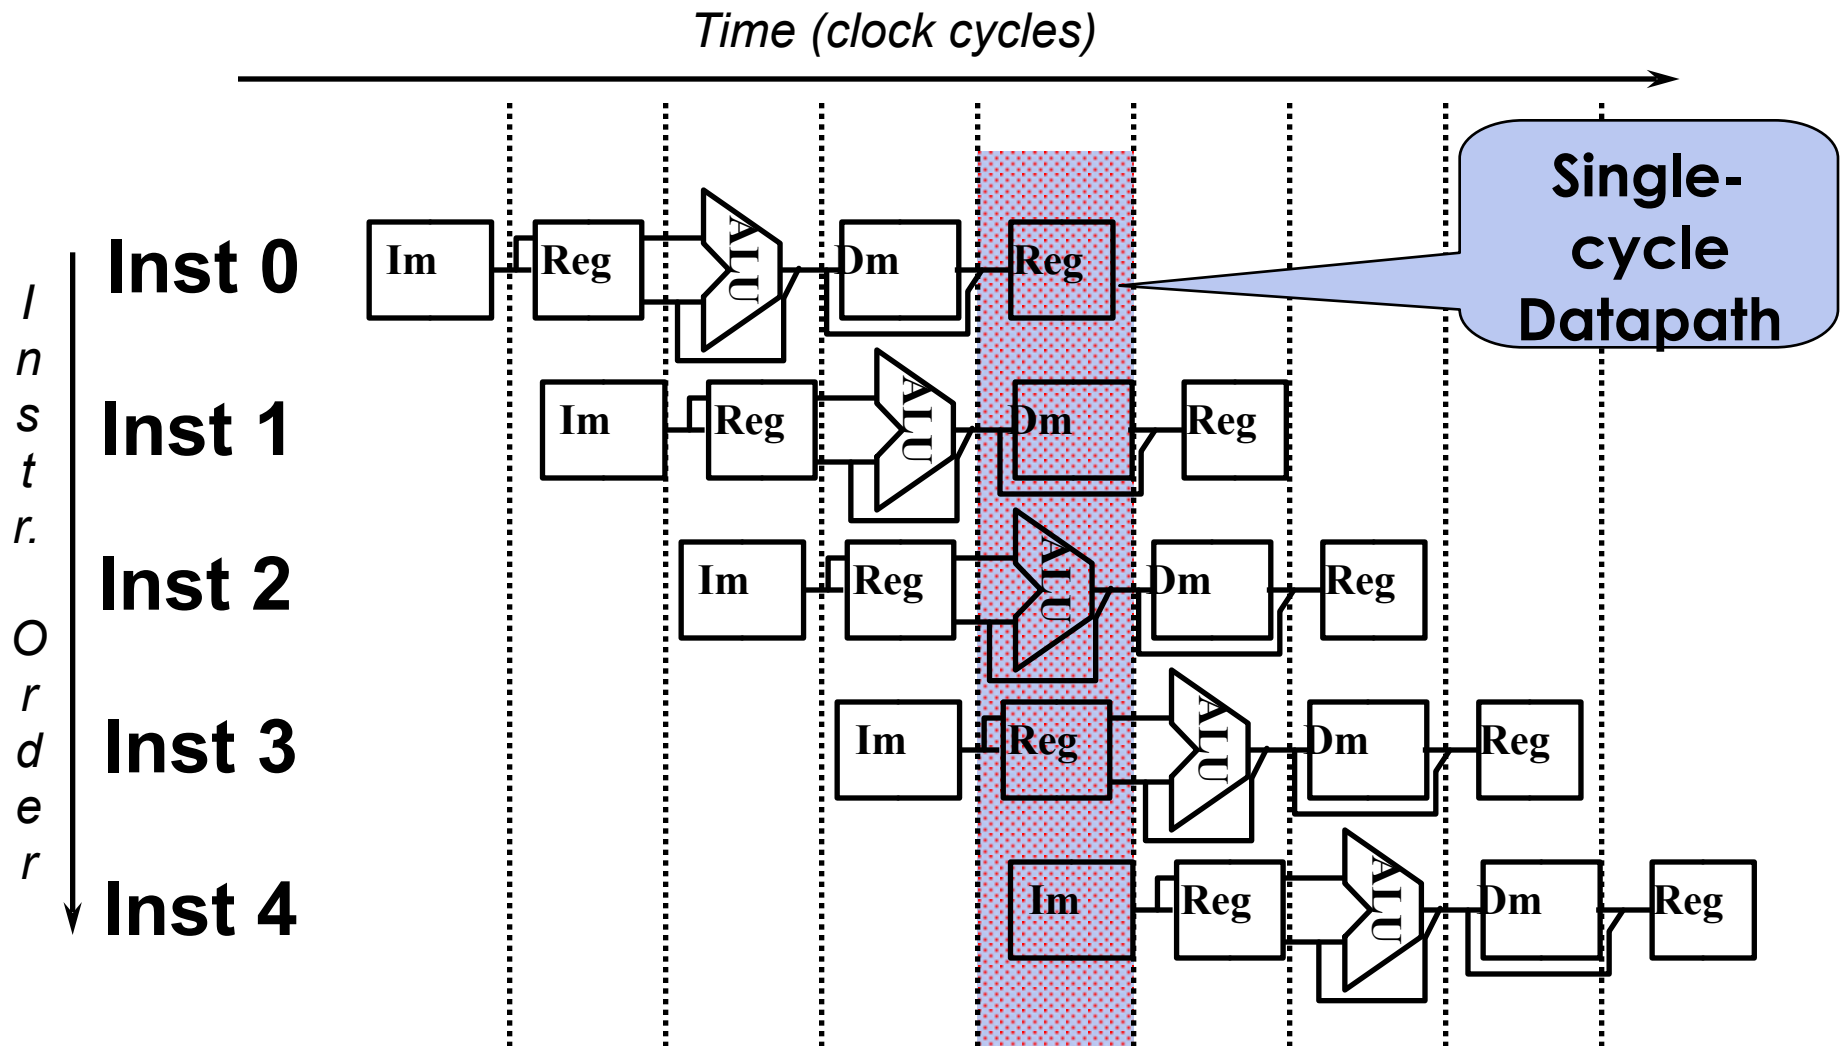
Outline

- ♦ **Pipelined processor design**
 - An overview of pipelining
 - A pipelined datapath
 - Pipelined control
- ♦ **Problems with pipelined processor**
 - Data hazards and forwarding
 - Data hazards and stalls
 - Branch hazards

Single-, Multi-Cycle, vs. Pipeline



Why Pipeline? Because the Resources are There!



Outline

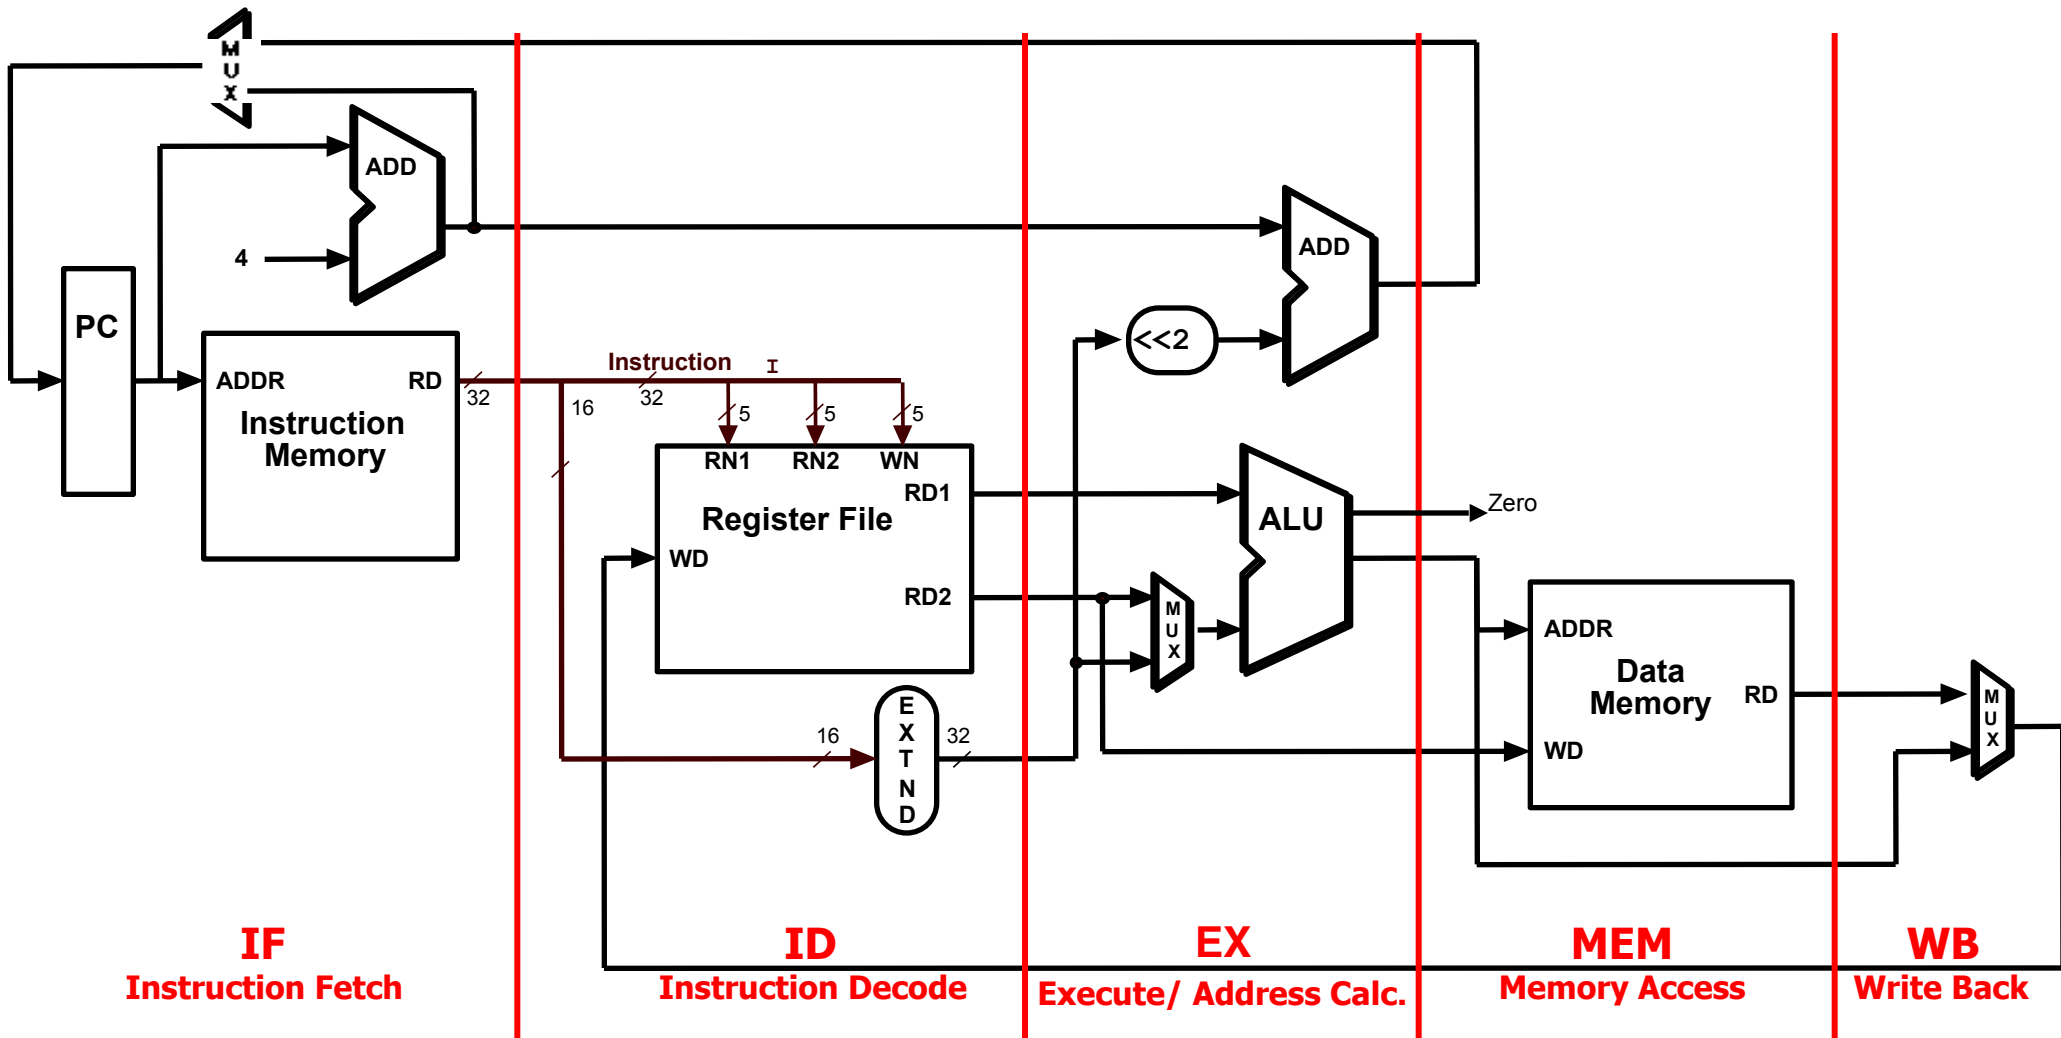
- ♦ **A pipelined datapath**
- ♦ **Pipelined control**
- ♦ **Data hazards and forwarding**
- ♦ **Data hazards and stalls**
- ♦ **Branch hazards**

Recall the 5 steps in Instruction Execution



1. Instruction Fetch & PC Increment (**IF**)
2. Instruction Decode and Register Read (**ID**)
3. Execution or calculate address (**EX**)
4. Memory access (**MEM**)
5. Write result into register (**WB**)

Recall “Single-Cycle” Datapath

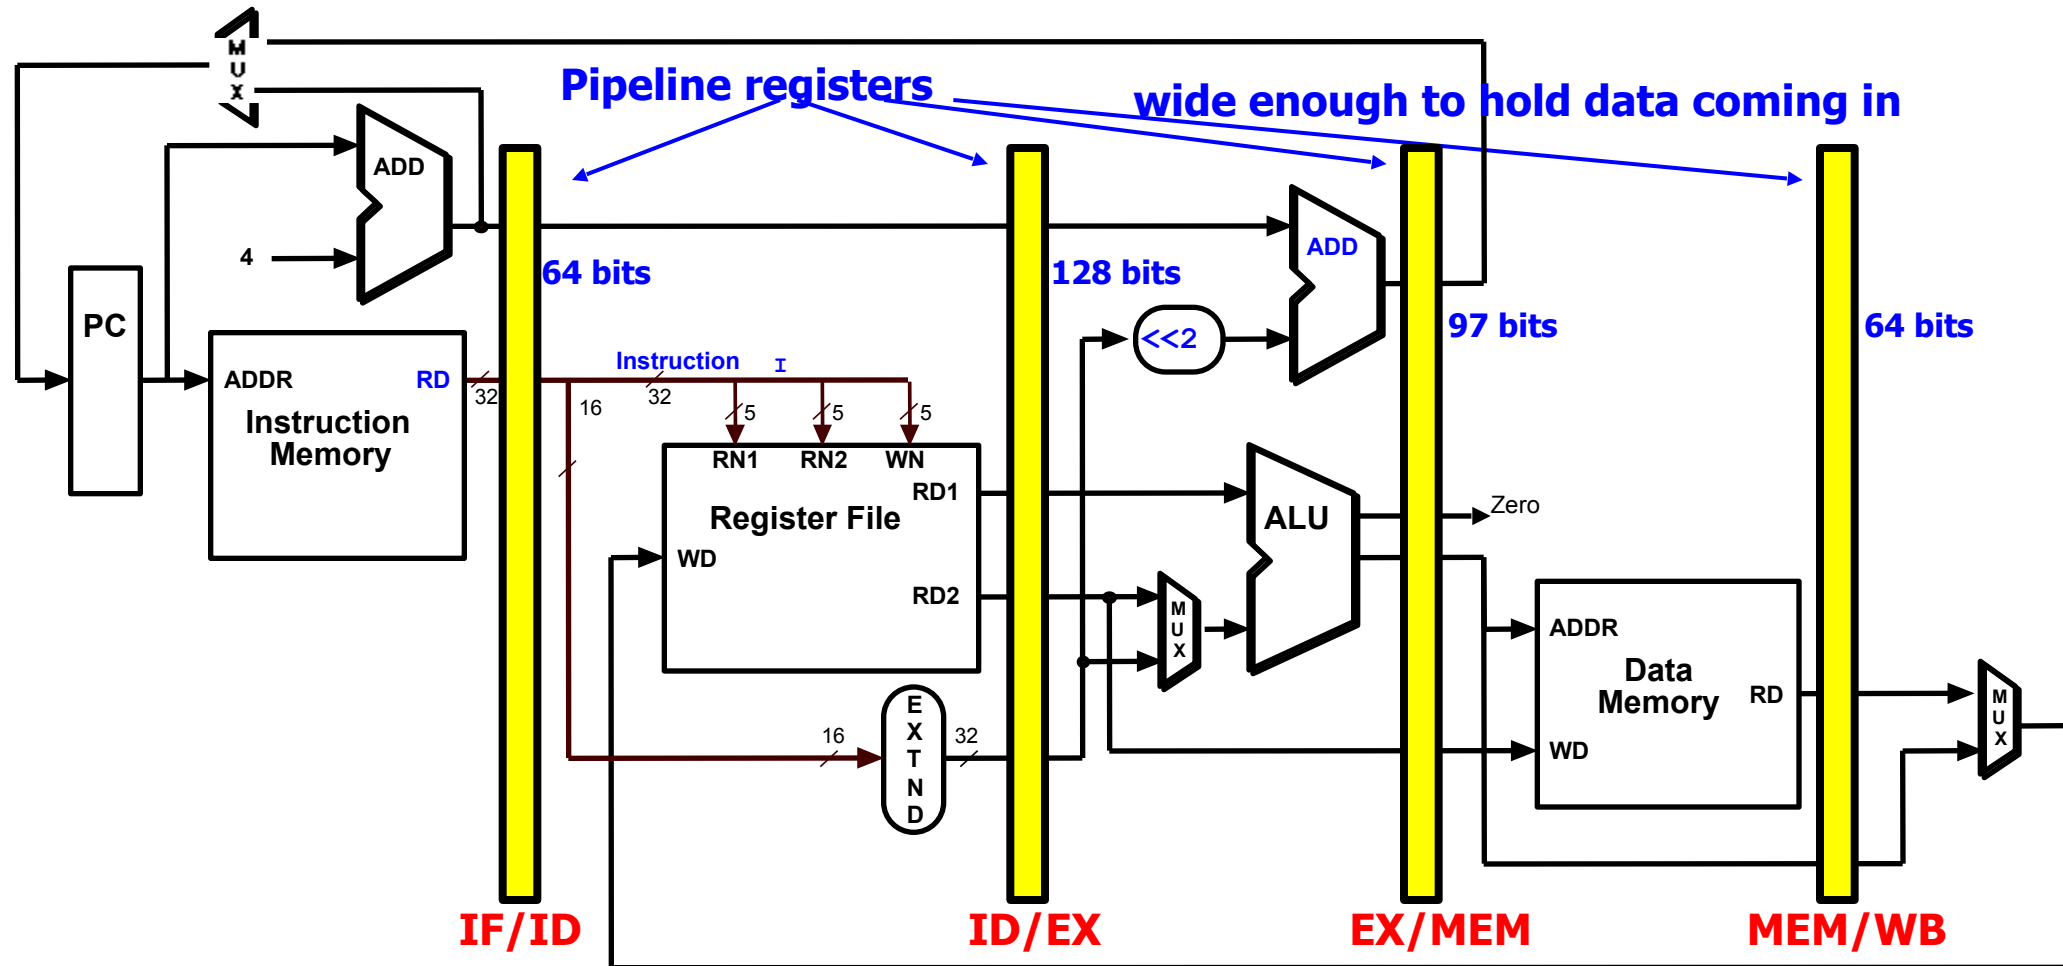


Pipelined Datapath – Key Idea

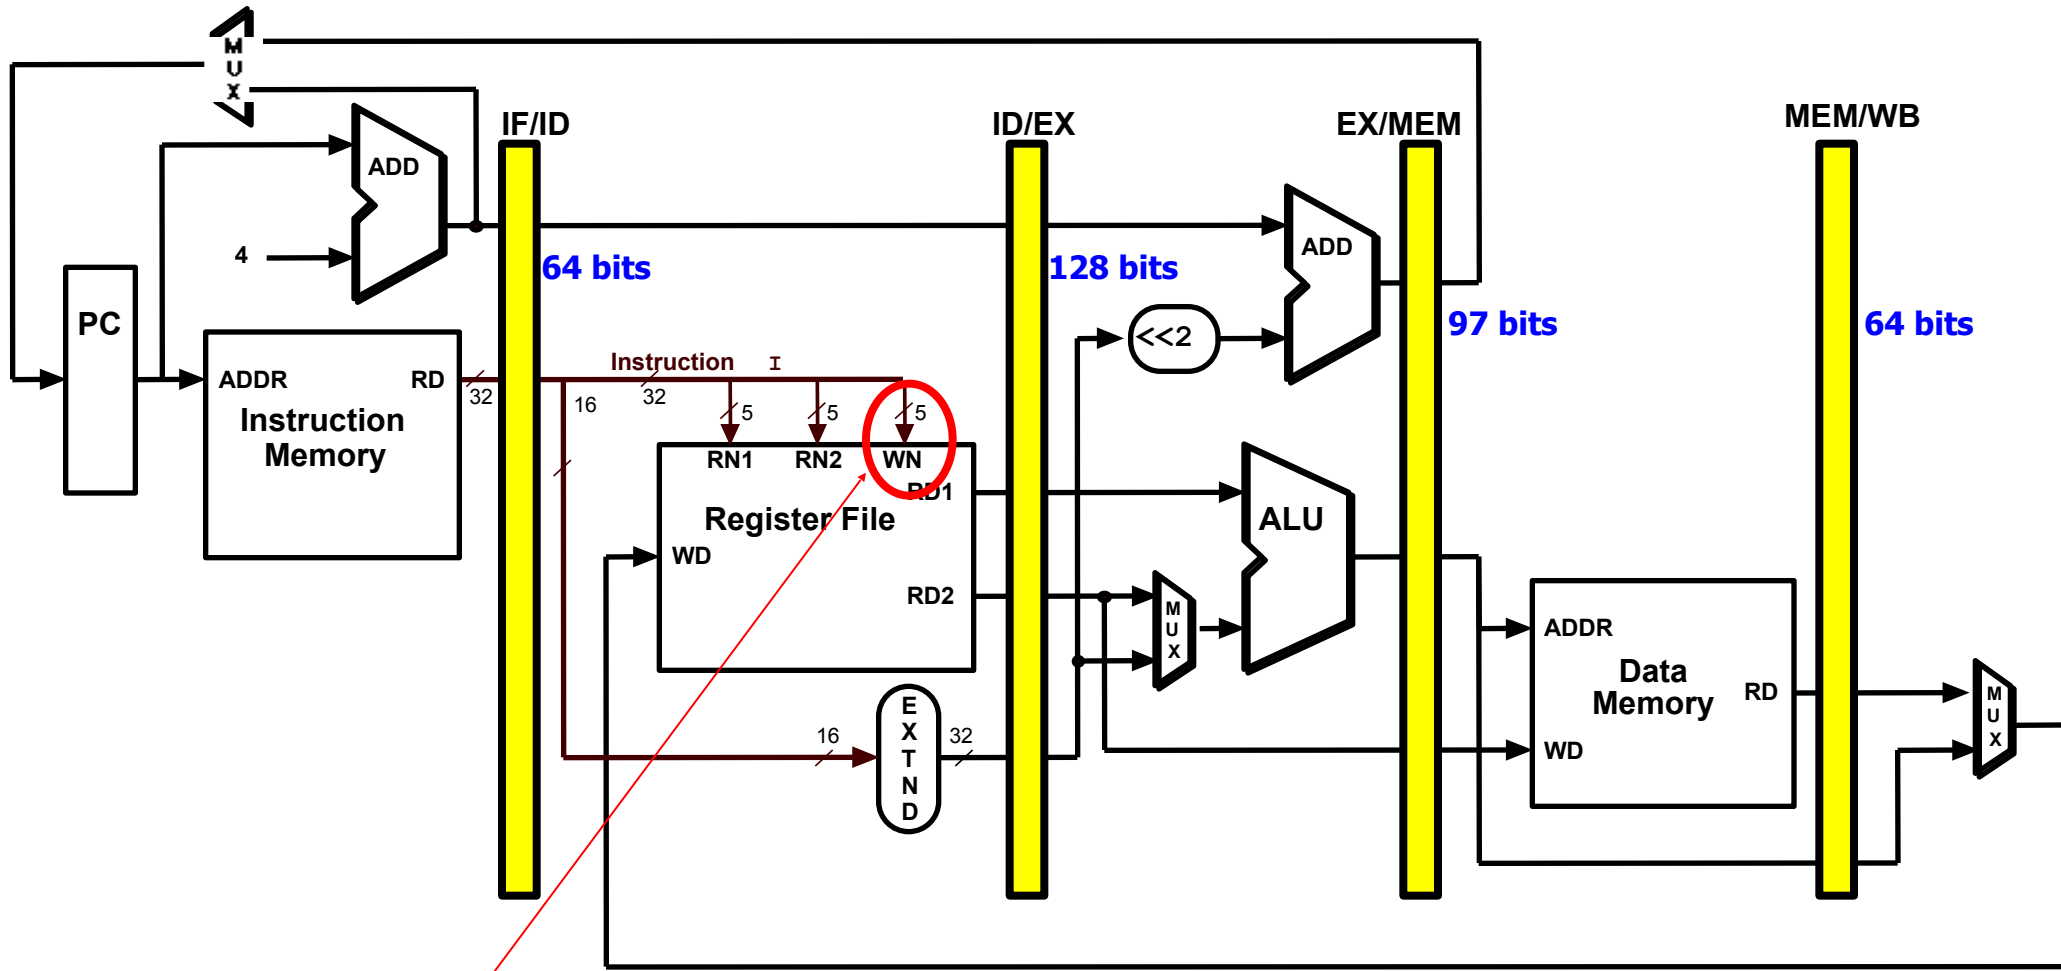


- ♦ We will have several instructions that execute simultaneously in the processor
- ♦ Similar to multicycle design
 - Revising the single-cycle design by introducing *extra* registers---*pipeline registers*---to hold “relevant state” between cycles

Pipelined Datapath

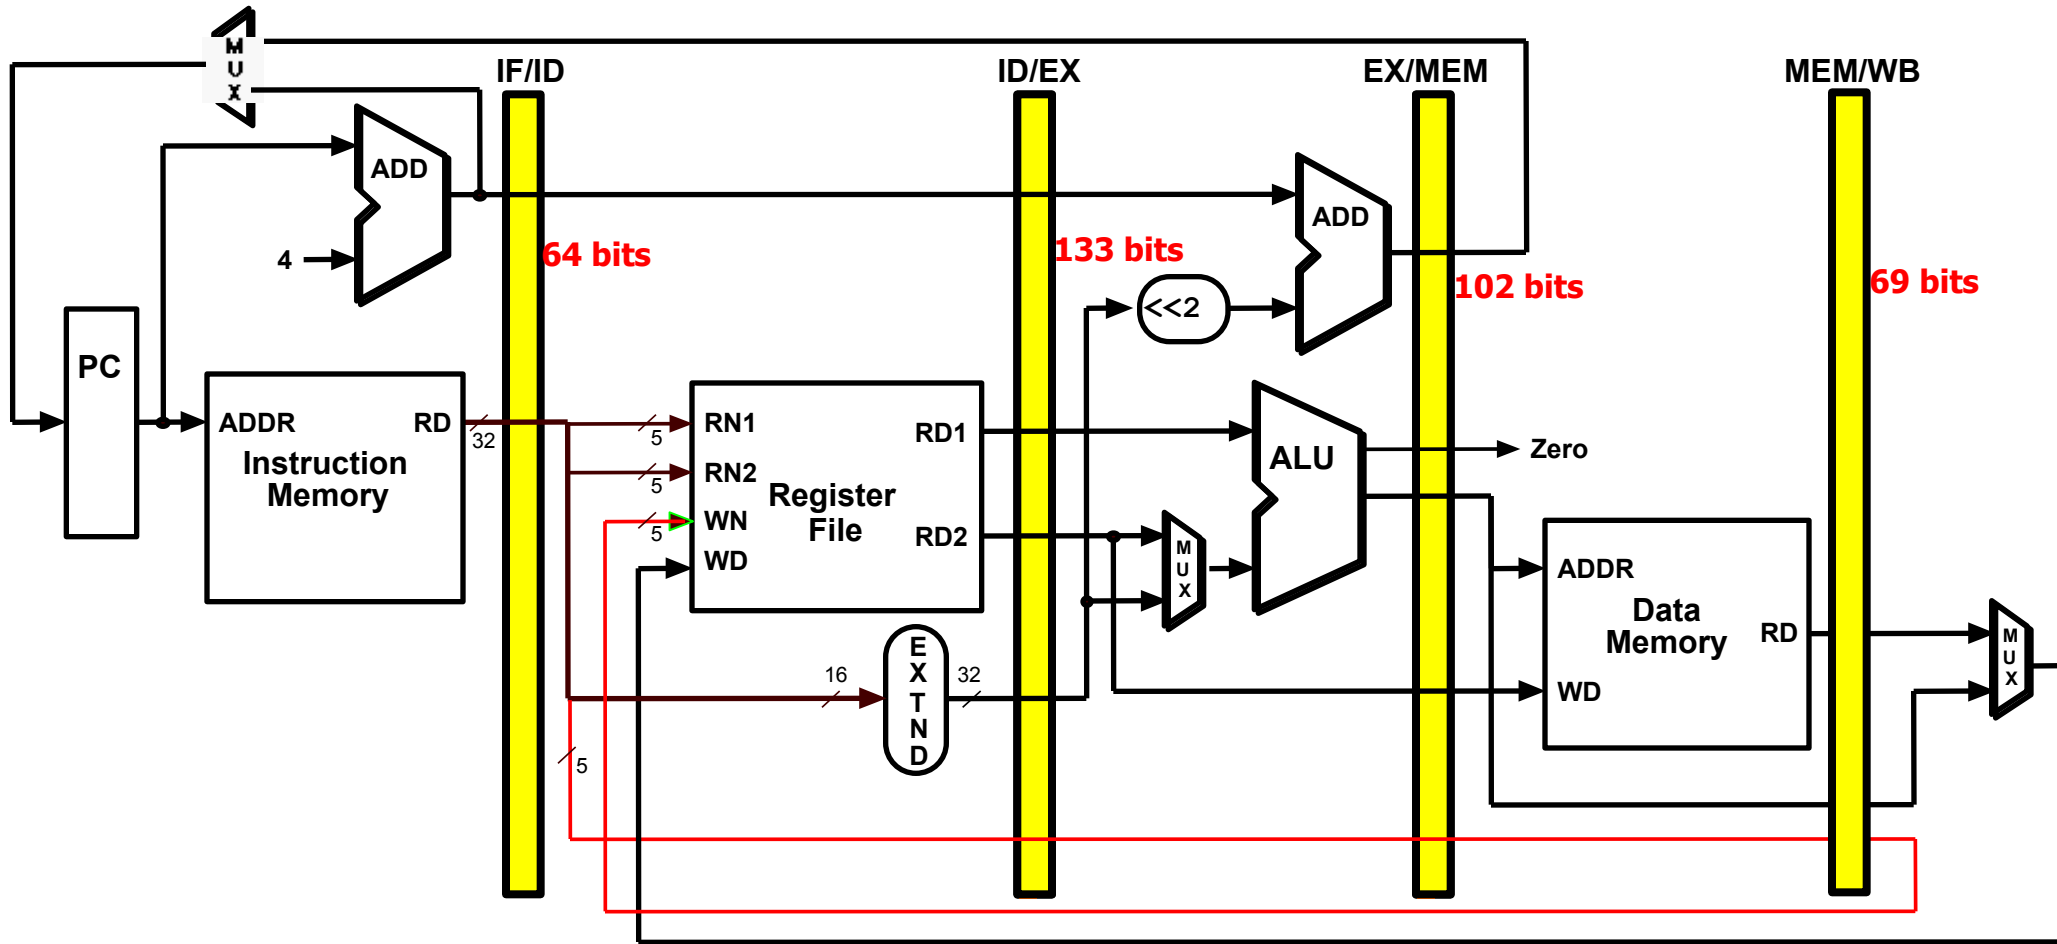


Bug in the Datapath



Write register number comes from another instruction in the latter stage!

Corrected Datapath



Destination register number is also passed through ID/EX, EX/MEM and MEM/WB registers, which are now wider by 5 bits

Pipelined Example

- ◆ Consider the following instruction sequence:

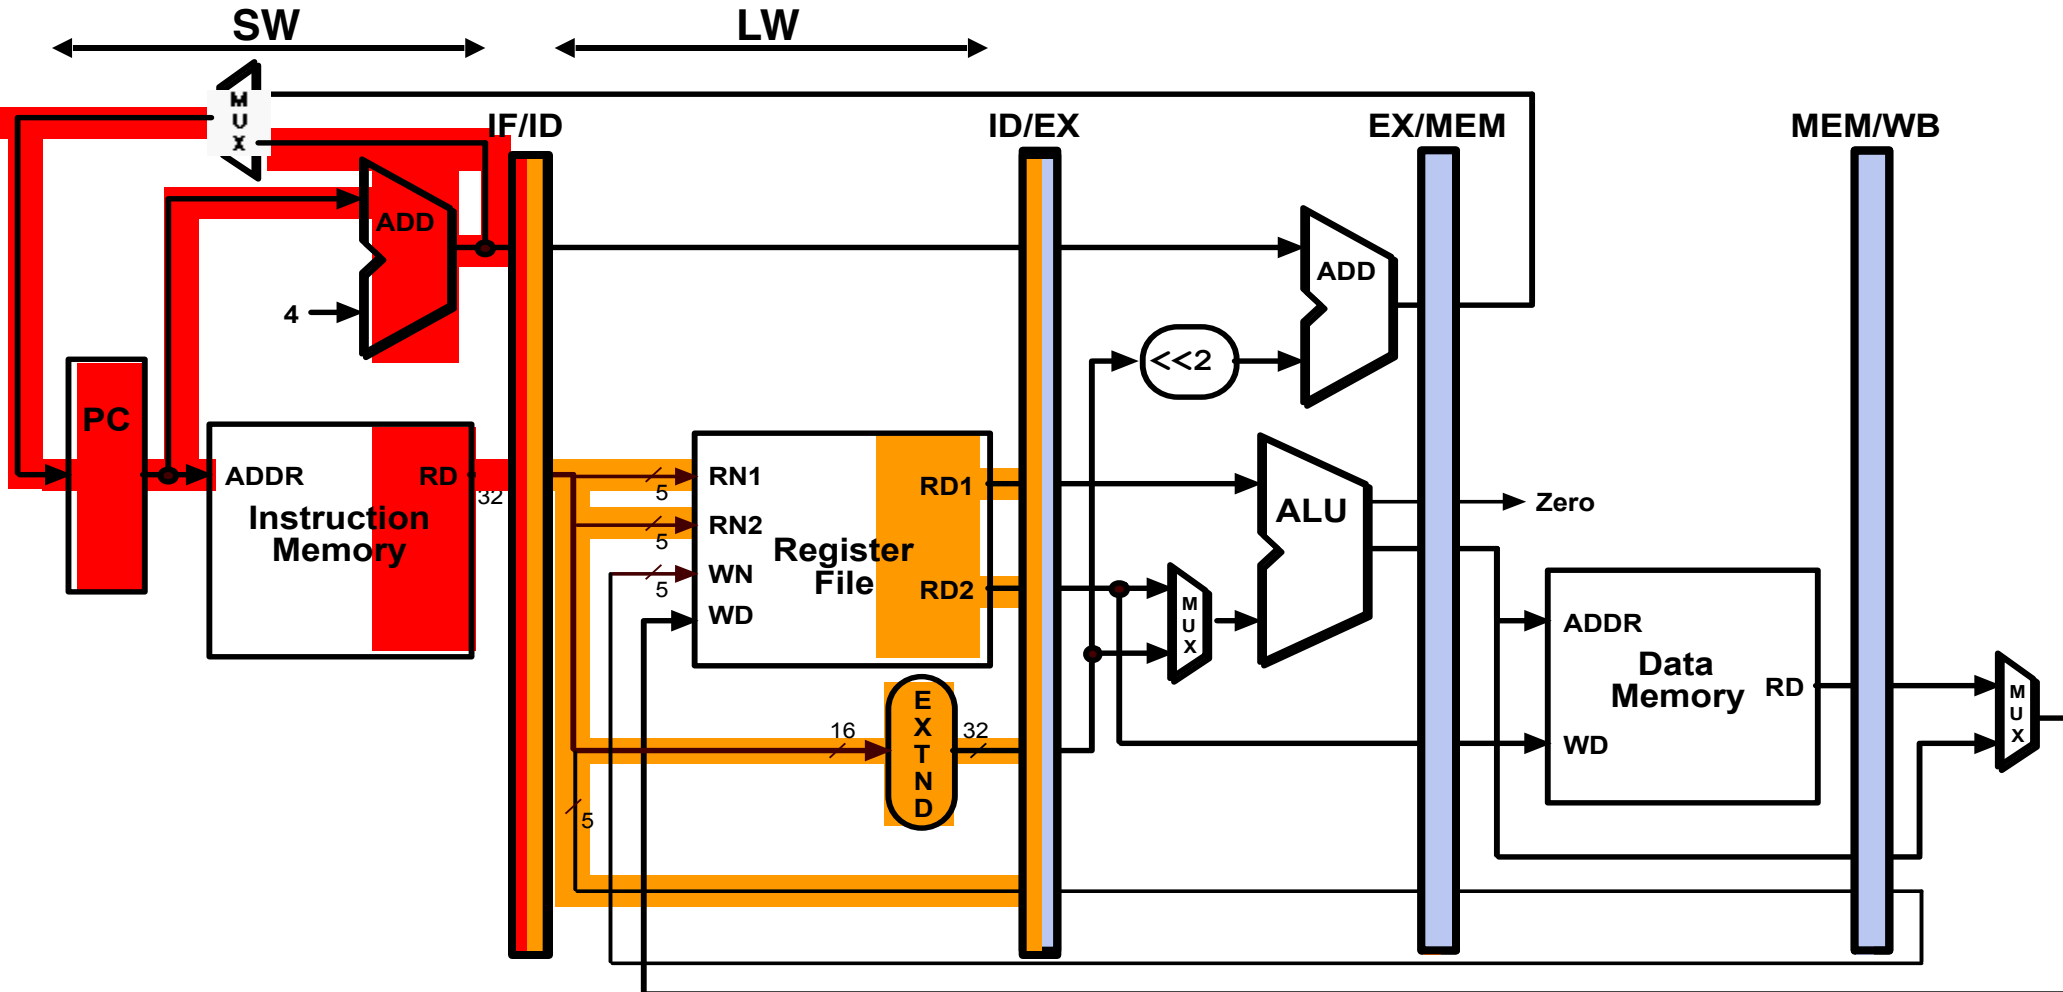
`lw $t0, 10($t1)`

`sw $t3, 20($t4)`

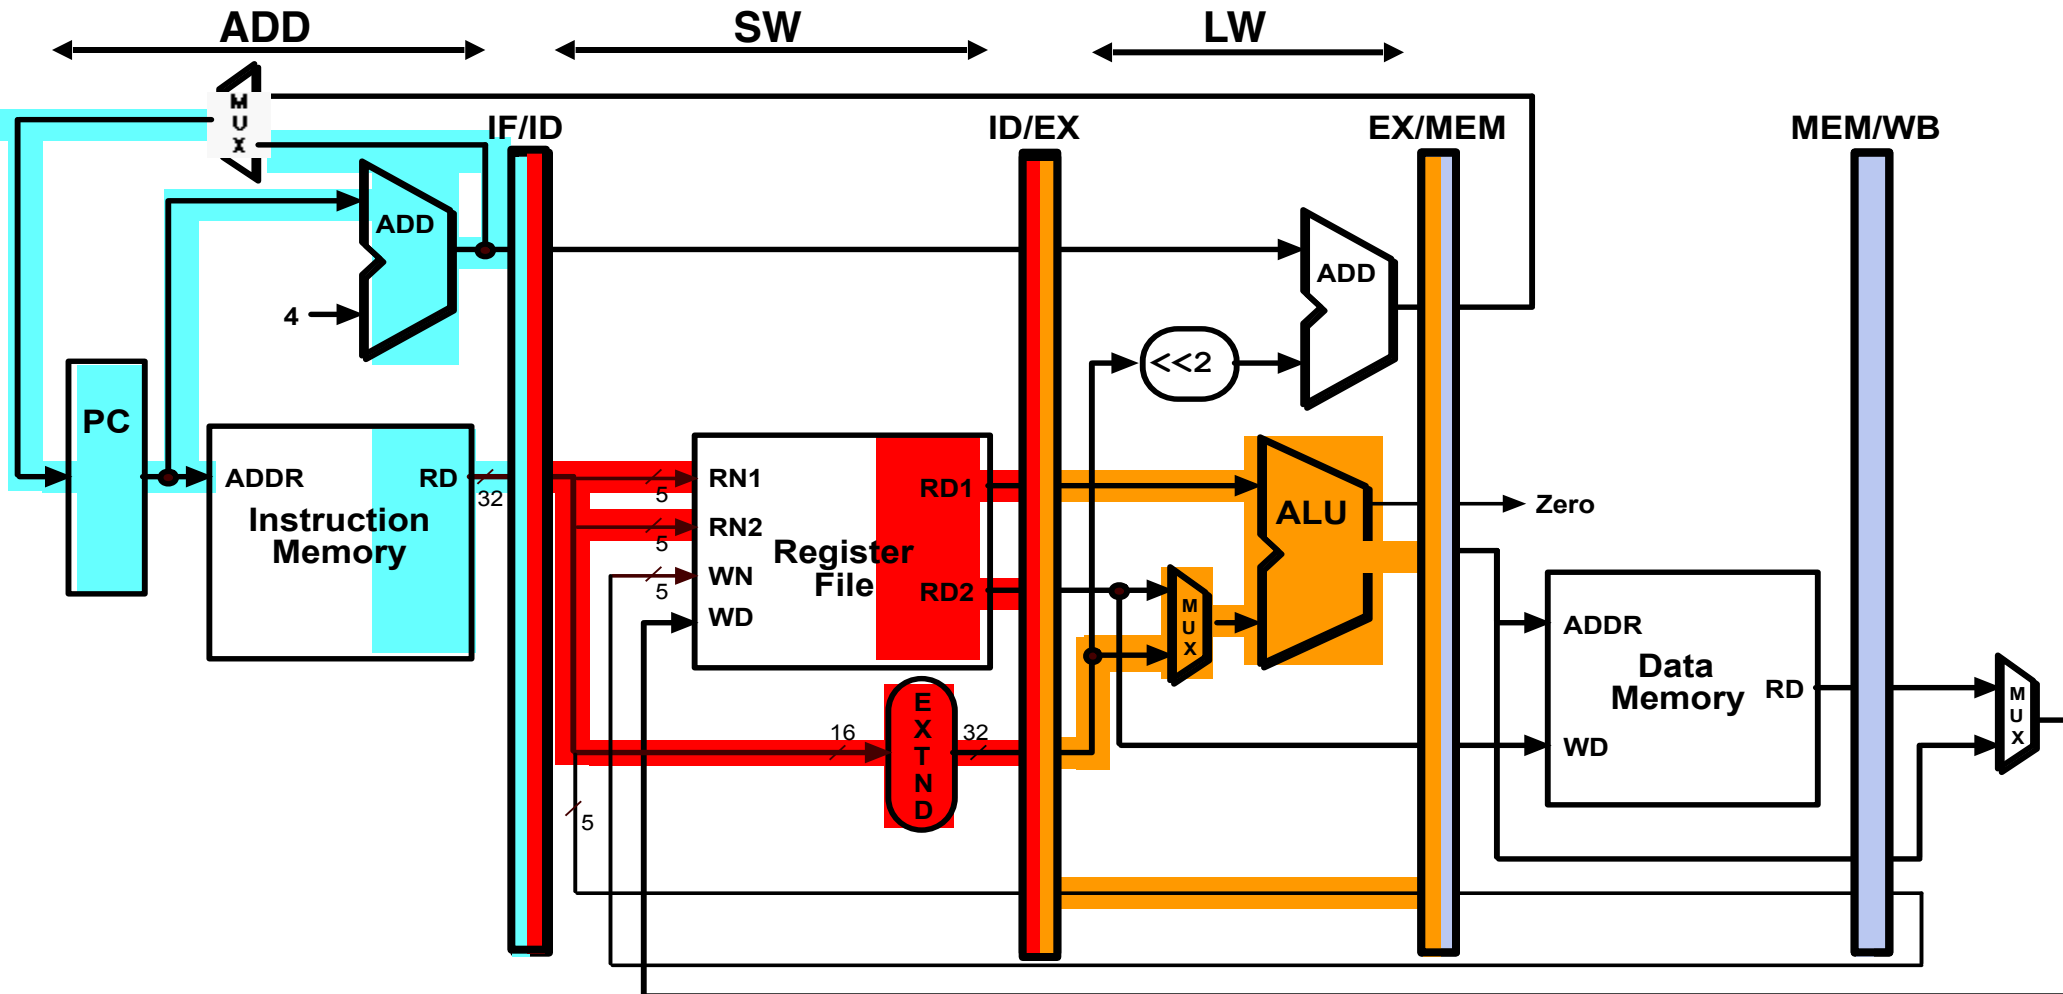
`add $t5, $t6, $t7`

`sub $t8, $t9, $t10`

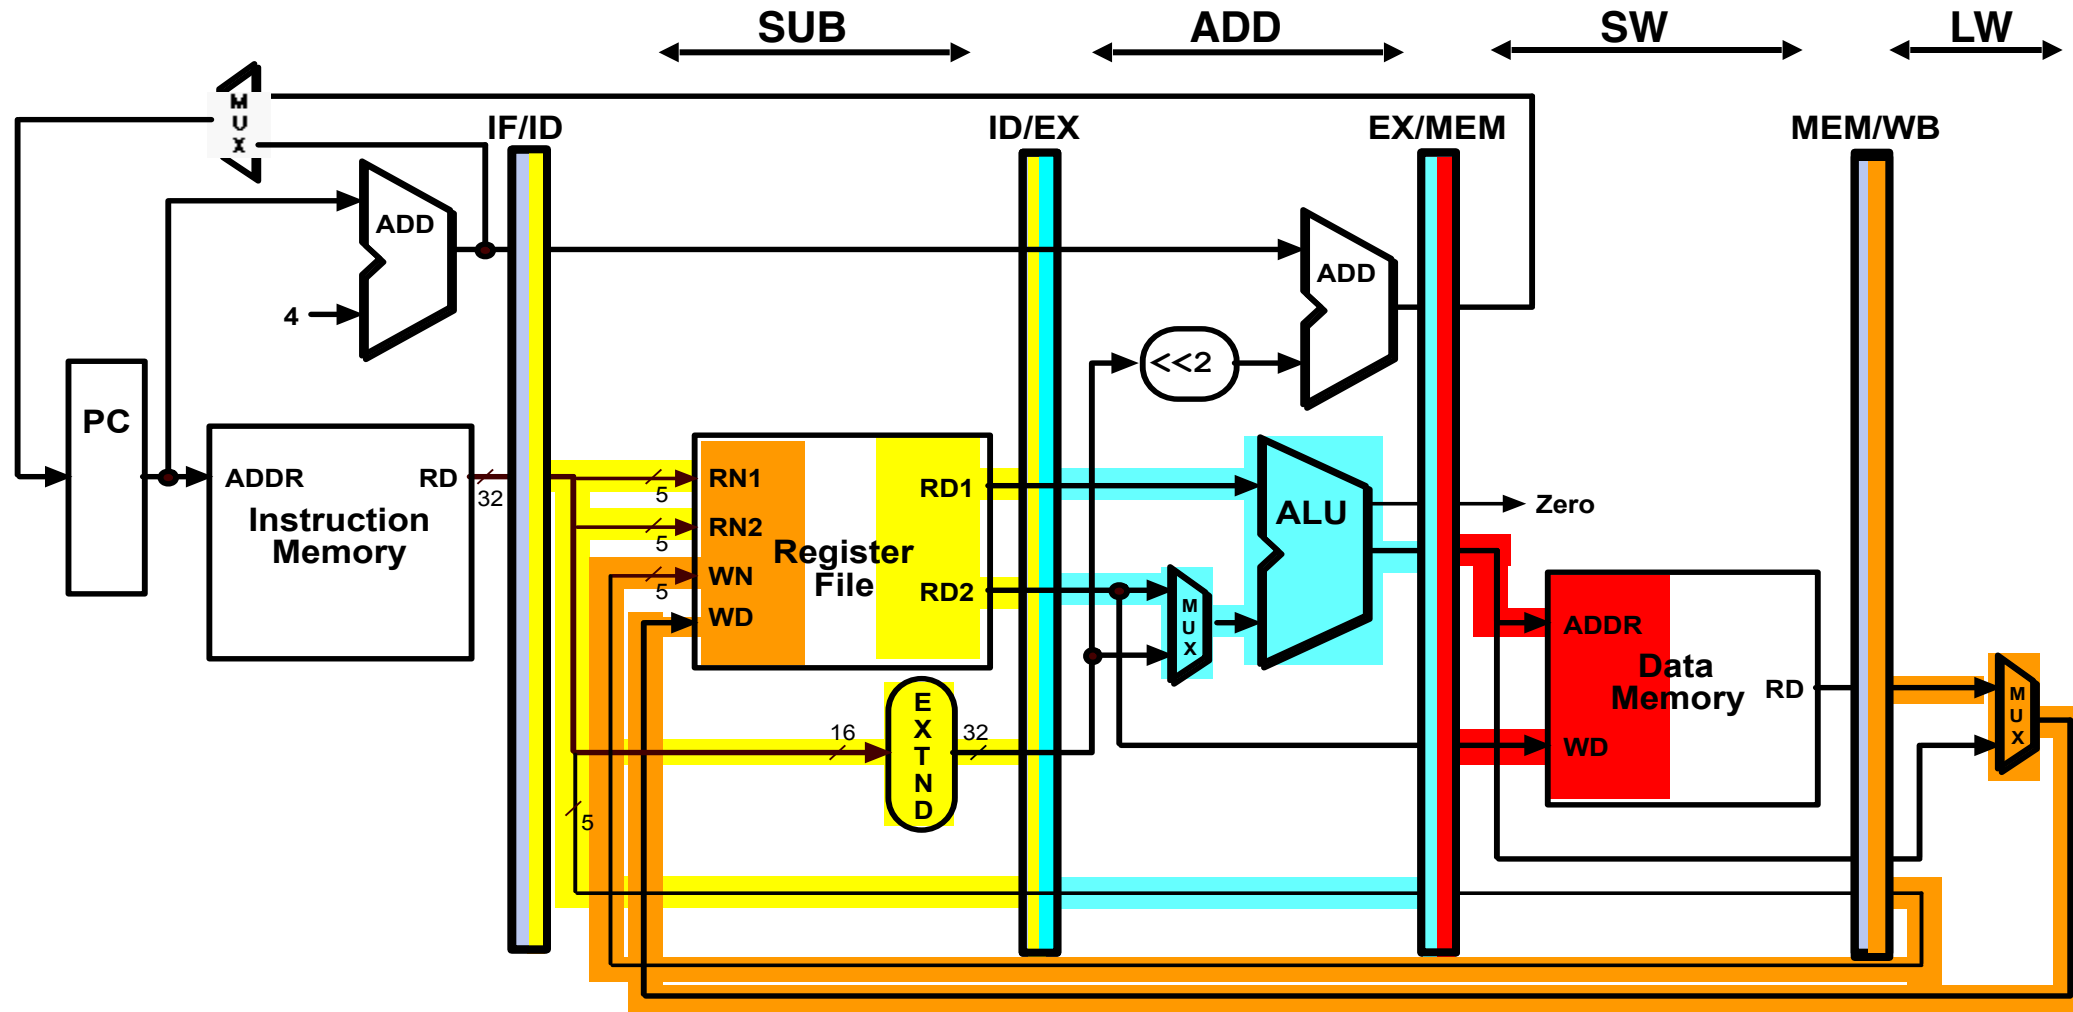
Single-Clock-Cycle Diagram: Clock Cycle 2



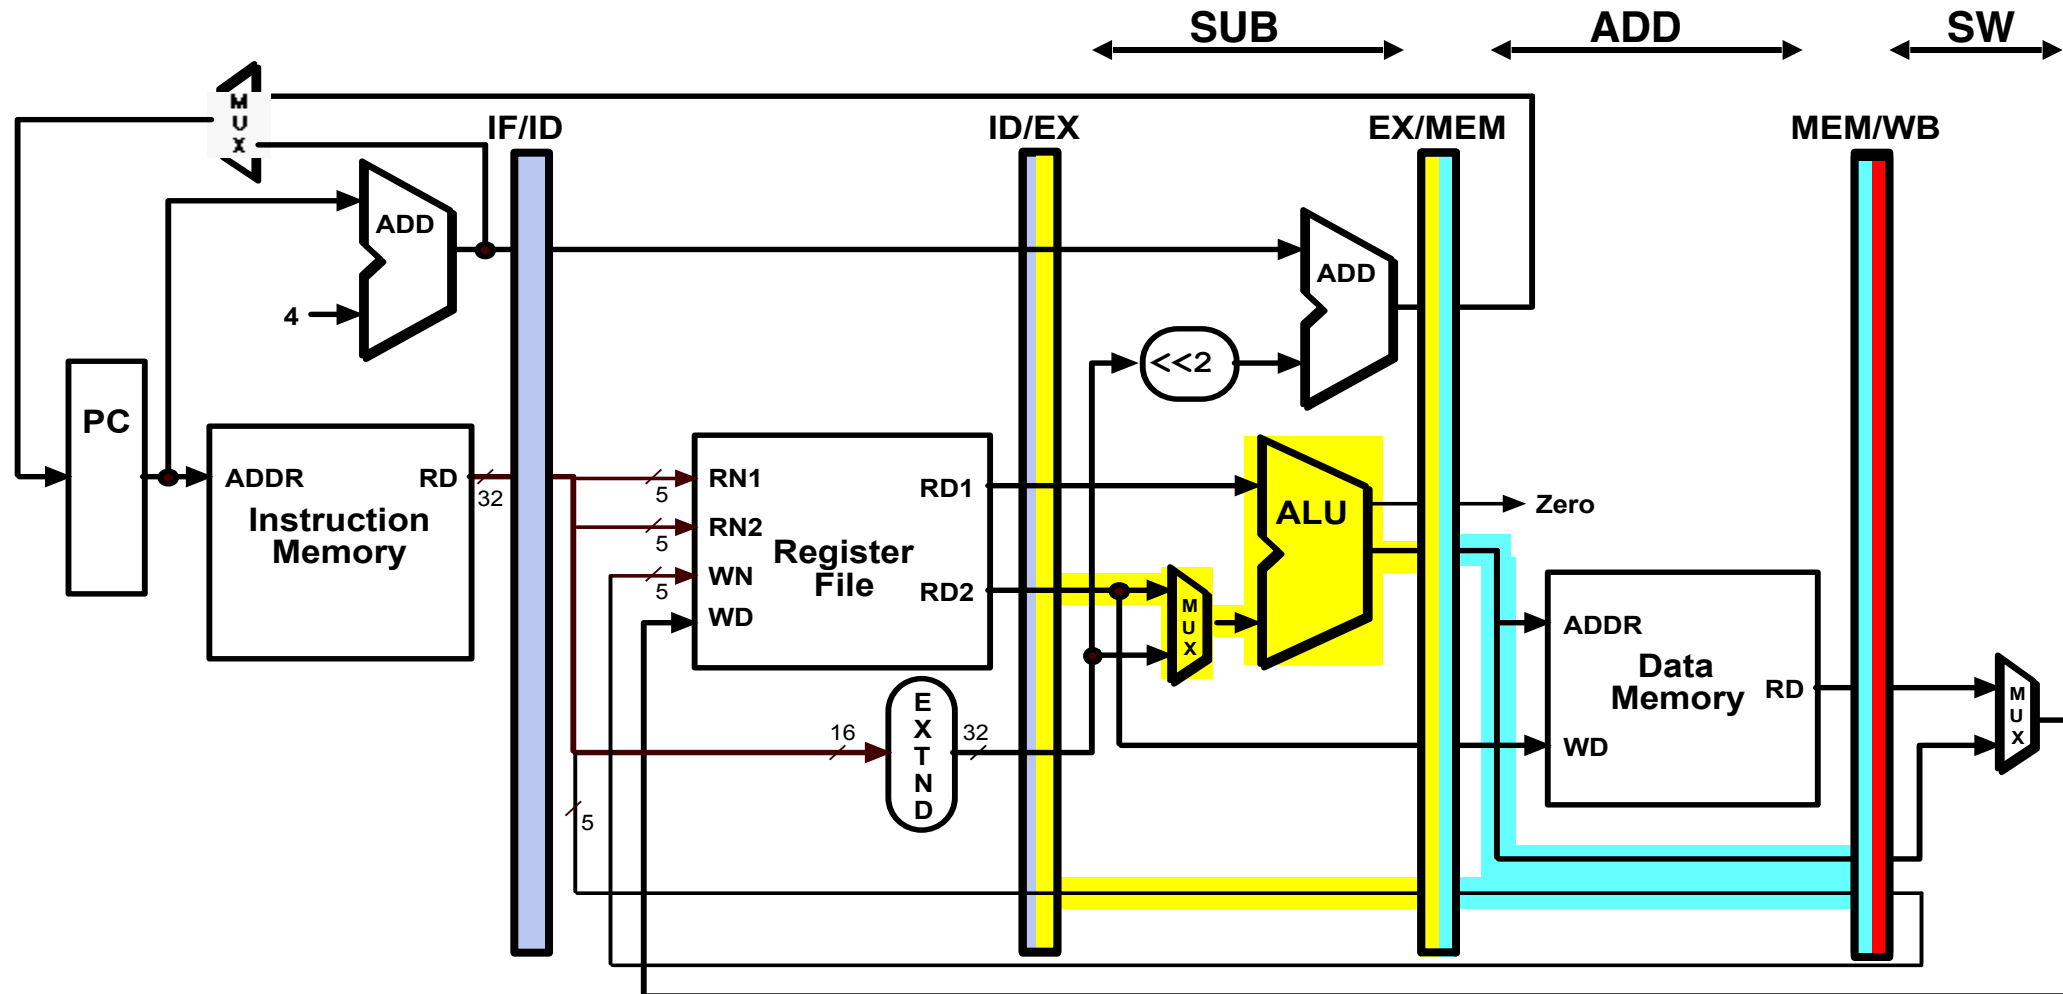
Single-Clock-Cycle Diagram: Clock Cycle 3



Single-Clock-Cycle Diagram: Clock Cycle 5



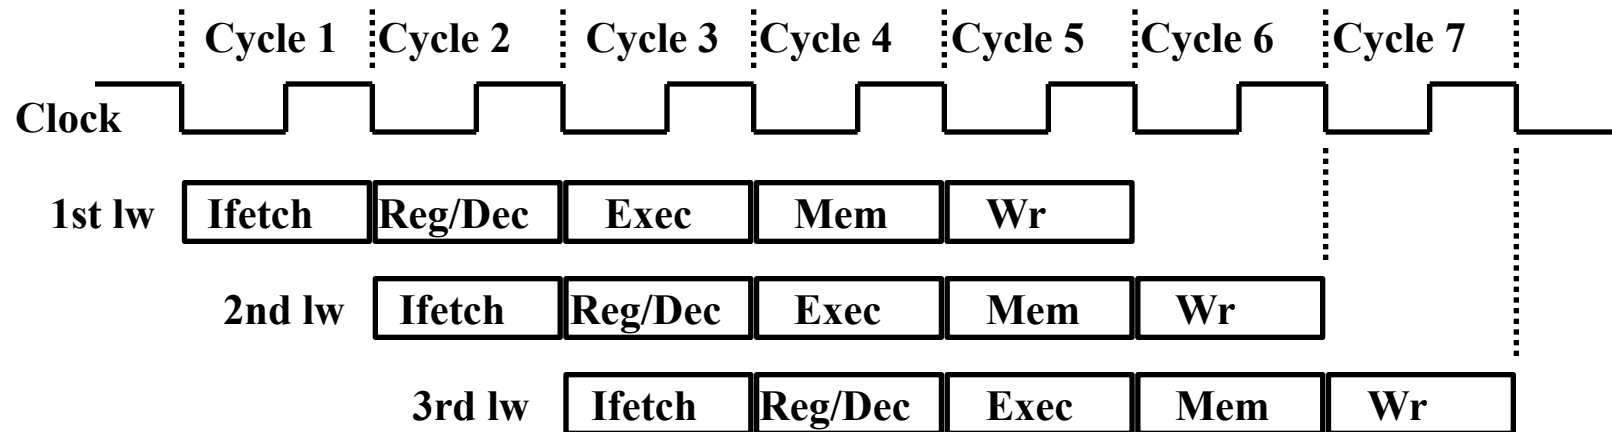
Single-Clock-Cycle Diagram: Clock Cycle 6



Observation


- ♦ The pipeline implementation inserts pipeline registers to *decouple the 5 stages*
- ♦ The CPI of an *ideal pipeline* (no stalls) is 1.
 - *What is the CPI value of multicycle design?*
 - *(discussed before)*
- ♦ One significant difference in the execution of an R-type instruction between multicycle and pipelined implementations:
 - **register write-back for the R-type instruction is the 5th (the last write-back) pipeline stage vs. the 4th stage for the multicycle implementation. Why? (see the below)**

Pipelining load



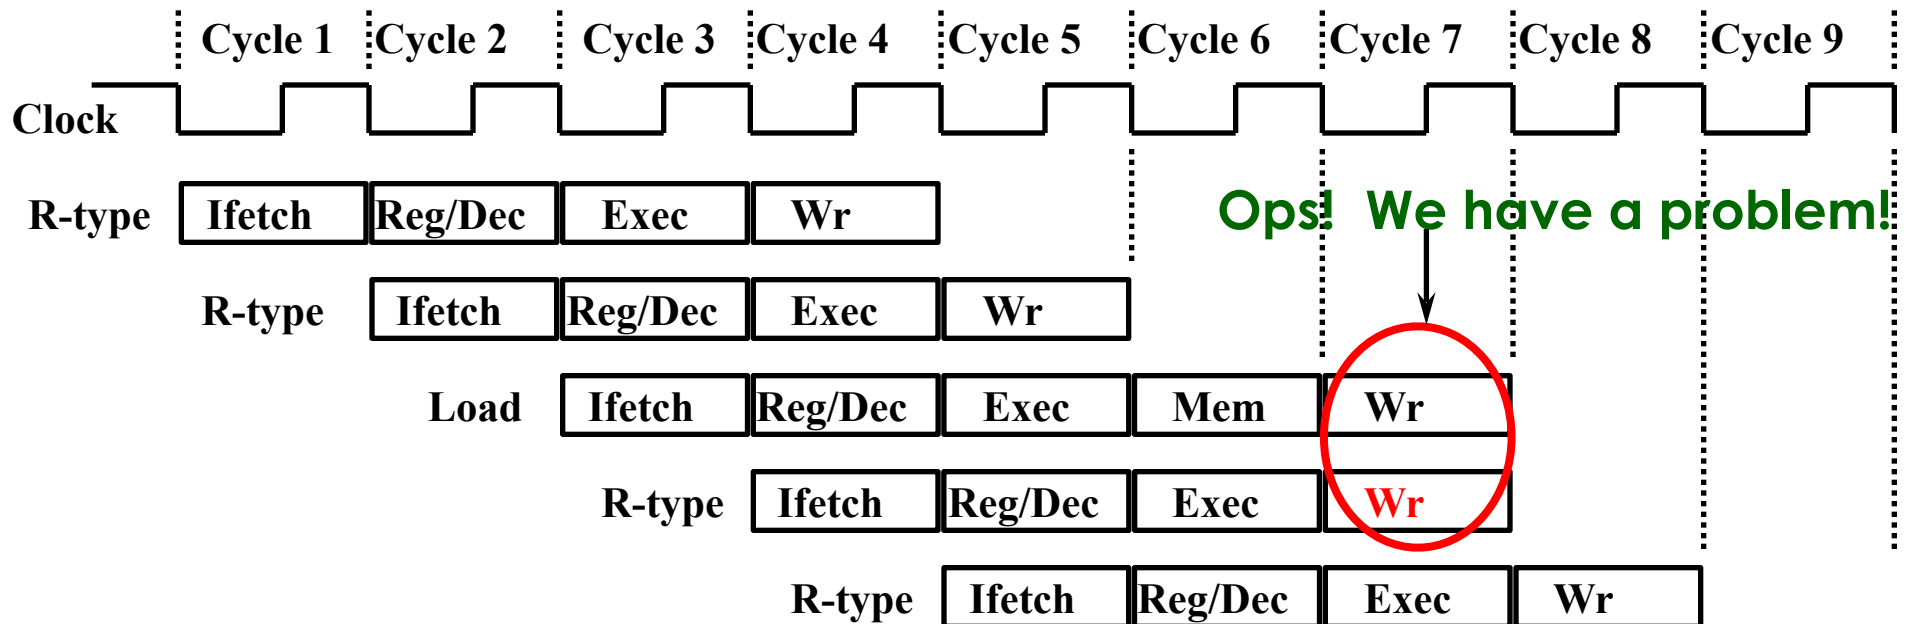
- ◆ 5 functional units in the pipeline datapath are:
 - **Instruction Memory** for the Fetch stage
 - Register File's **Read ports** (busA and busB) for the Reg/Dec stage
 - **ALU** for the Exec stage
 - **Data Memory** for the MEM stage
 - Register File's **Write port** (busW) for the WB stage

Pipelining R-type Instructions



- ♦ **IF:** fetch the instruction from the Instruction Memory
- ♦ **ID:** registers fetch and instruction decode
- ♦ **EX:** ALU operates on the two register operands
- ♦ **WB:** write ALU output back to the register file

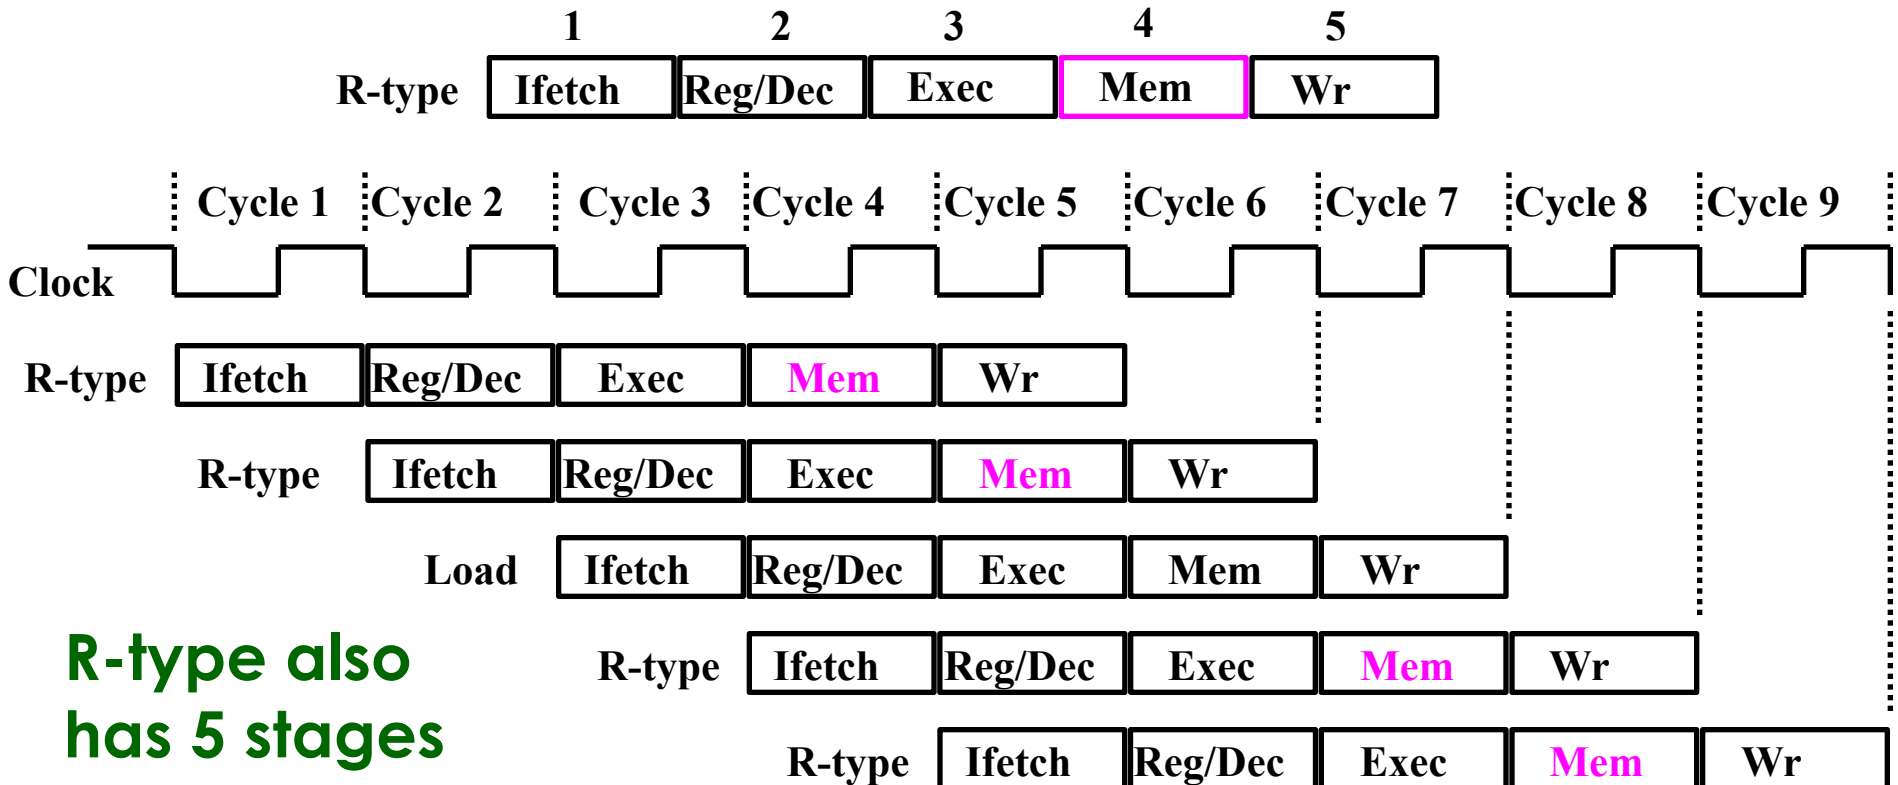
Pipelining R-type and load




- ♦ We have a **structural hazard**:
 - Two instructions try to write to the register file at the same time!
 - Only one write port

Solution: Delay R-type's Write

- ◆ Delay R-type's register write by one cycle:
 - R-type also use Reg File's write port at Stage 5
 - MEM is a NOP stage: nothing is being done.



SW




- ♦ **IF:** fetch the instruction from the Instruction Memory
- ♦ **ID:** registers fetch and instruction decode
- ♦ **EX:** calculate the memory address
- ♦ **MEM:** write the data into the Data Memory

Add an extra stage:

- ♦ **WB: NOP**

BEQ



- ♦ **IF:** fetch the instruction from the Instruction Memory
- ♦ **ID:** registers fetch and instruction decode
- ♦ **EX:**
 - compare the two register operand
 - select correct branch target address
 - latch into PC

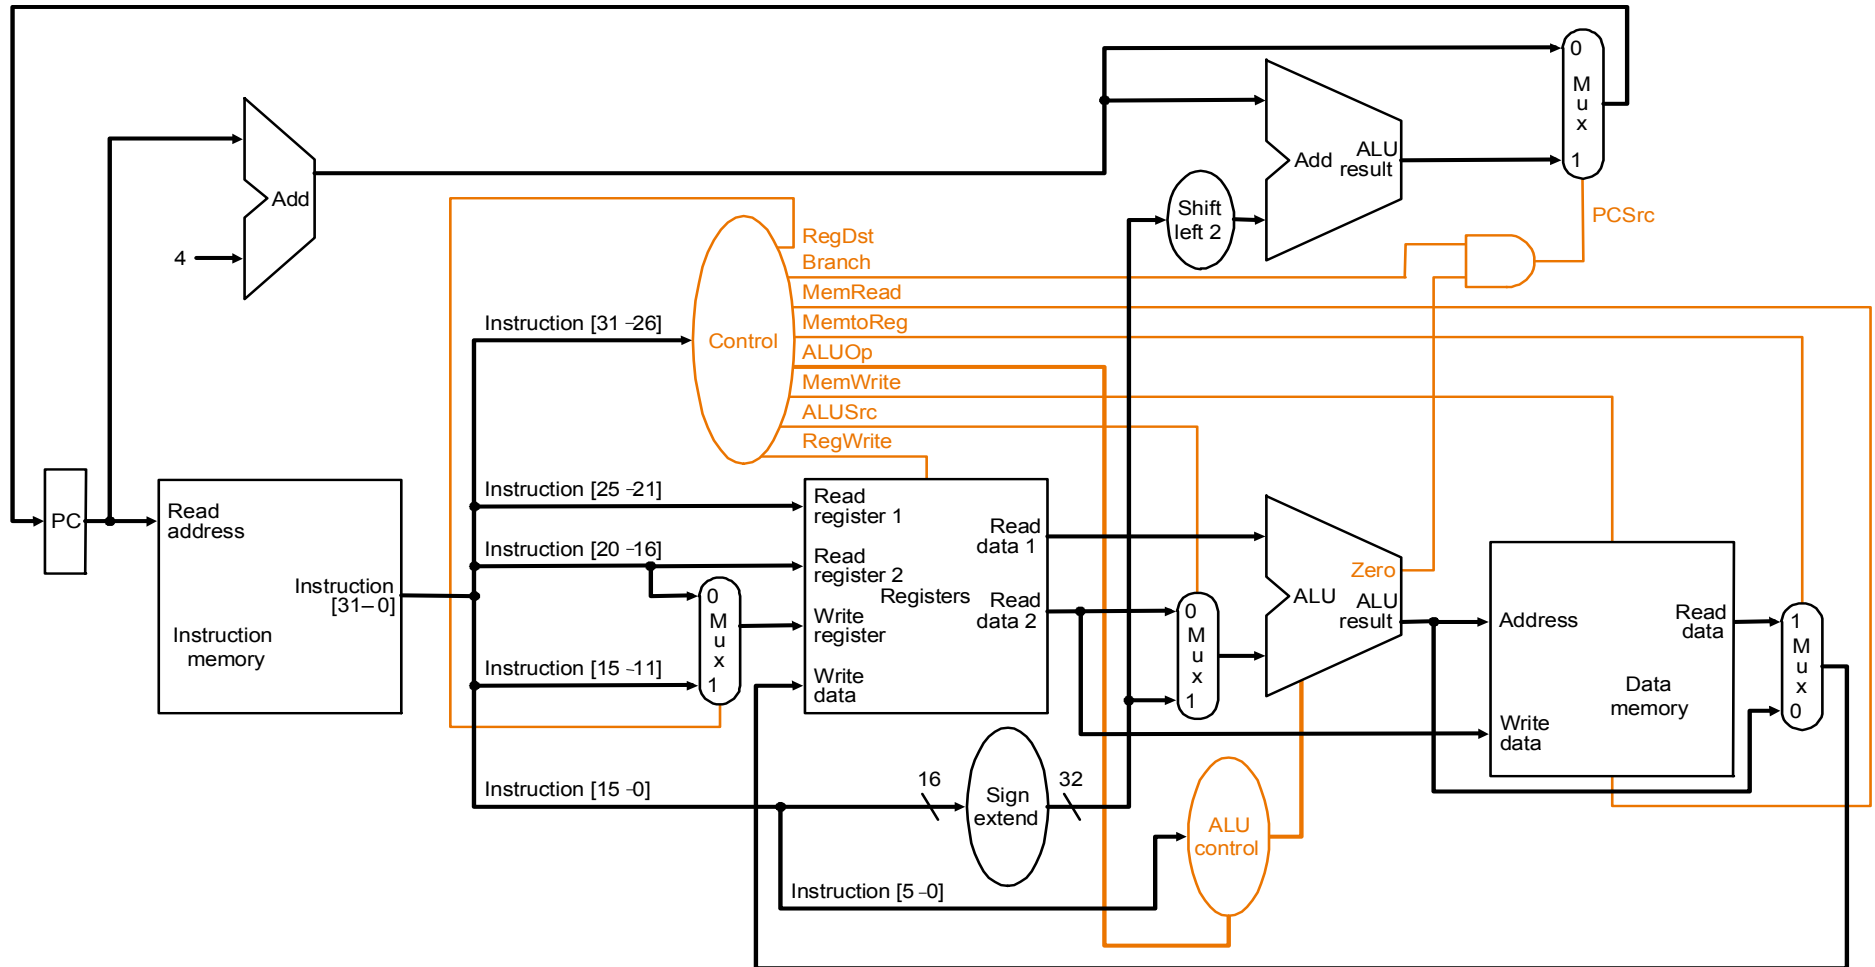
Add two extra stages:

- ♦ **MEM:** NOP
- ♦ **WB:** NOP

Outline

- ♦ A pipelined datapath
- ♦ **Pipelined control**
- ♦ Data hazards and forwarding
- ♦ Data hazards and stalls
- ♦ Branch hazards

Recall Single-Cycle Control – the Datapath



Recall Control for Single-Cycle Design

ALUOp		Funct field						Operation
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0	
0	0	X	X	X	X	X	X	010
0	1	X	X	X	X	X	X	110
1	X	X	X	0	0	0	0	010
1	X	X	X	0	0	1	0	110
1	X	X	X	0	1	0	0	000
1	X	X	X	0	1	0	1	001
1	X	X	X	1	0	1	0	111

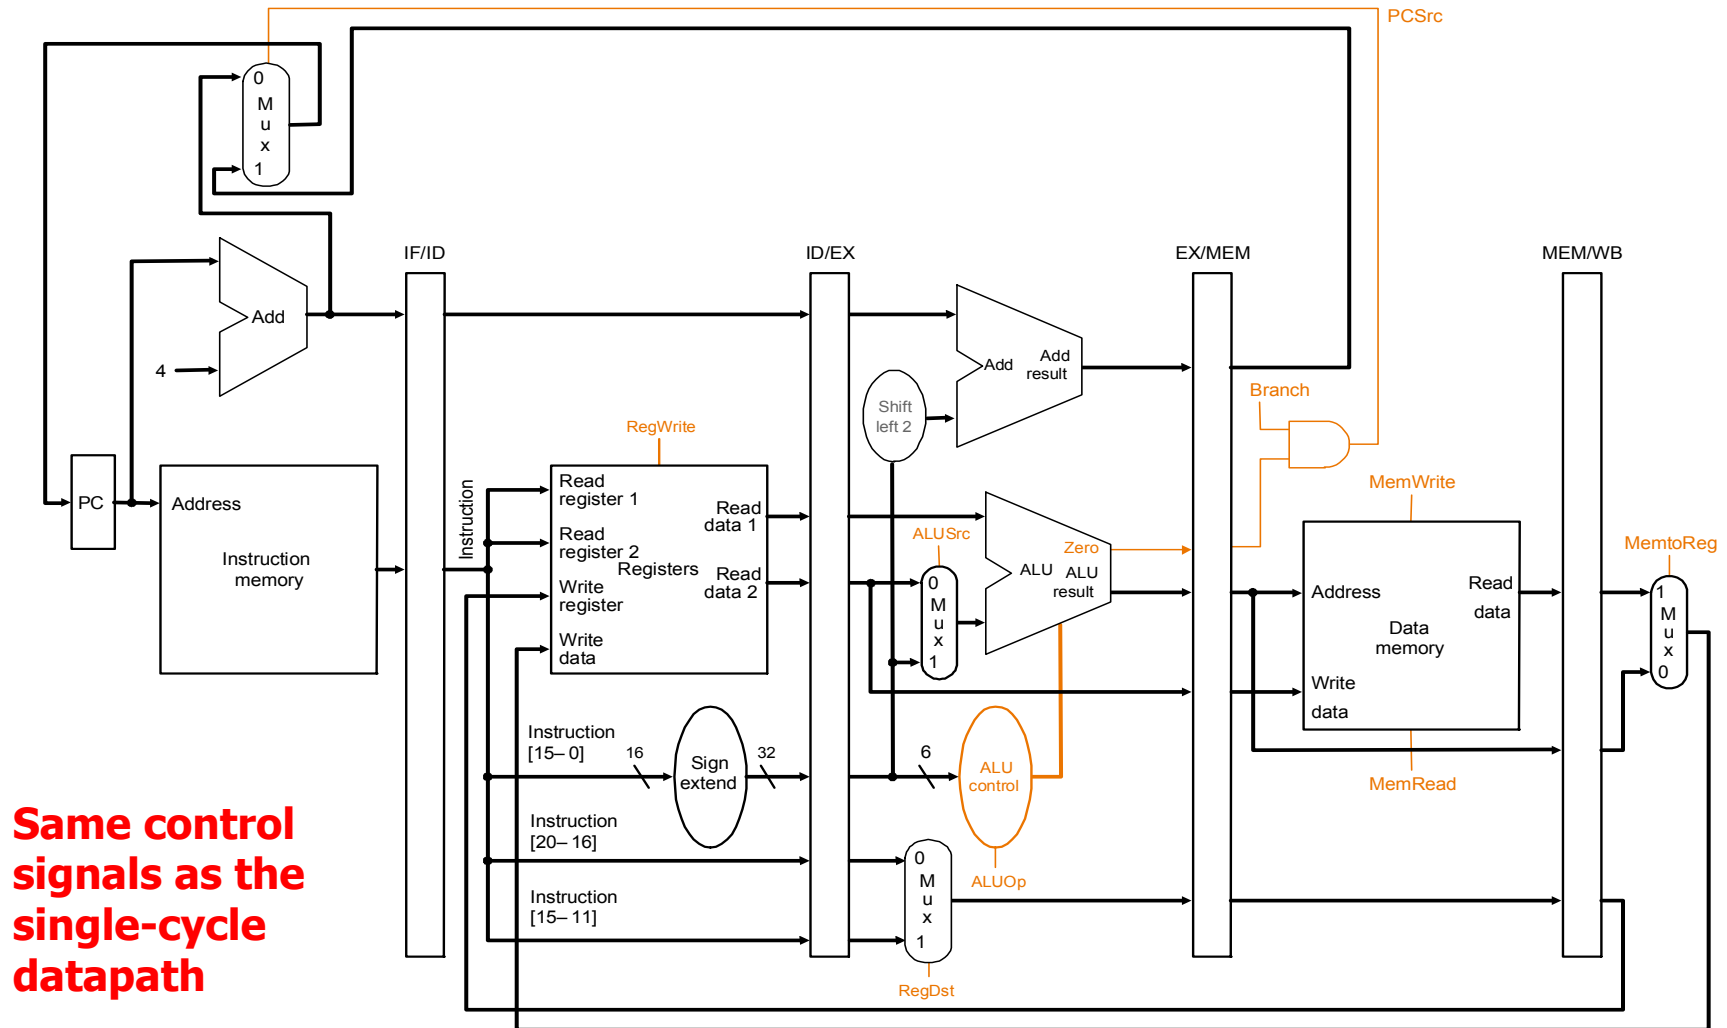
Instruction	RegDst	ALUSrc	Memto-Reg	Reg Write	Mem Read	Mem Write	Branch	ALUOp1	ALUp0
R-format	1	0	0	1	0	0	0	1	0
lw	0	1	1	1	1	0	0	0	0
sw	X	1	X	0	0	1	0	0	0
beq	X	0	X	0	0	0	1	0	1

Pipeline Control



- ♦ **Initial design – motivated by single-cycle datapath control – use the same control signals**
- ♦ **Since control signals are associated with components active during a single pipeline stage, can group control lines into five groups according to pipeline stage**
- ♦ **Need to set control signals during each pipeline stage**

Pipelined Datapath with Control I



Pipeline Control Signals

- instruction fetch / PC increment
- instruction decode / register fetch
- execution / address calculation (**EX**)
- memory access (**M**)
- write back (**WB**)

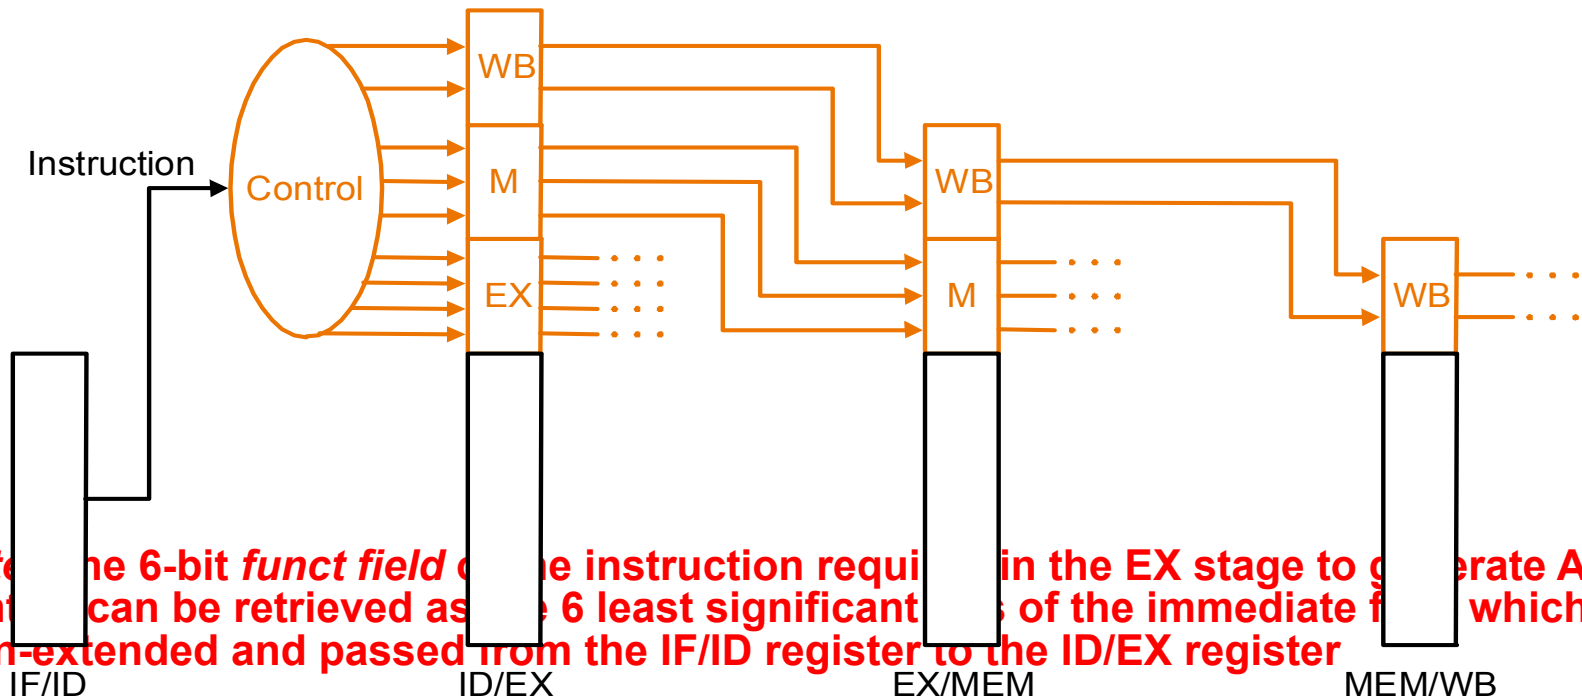
Nothing to control as instruction memory read and PC write are always enabled

(the following table)

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

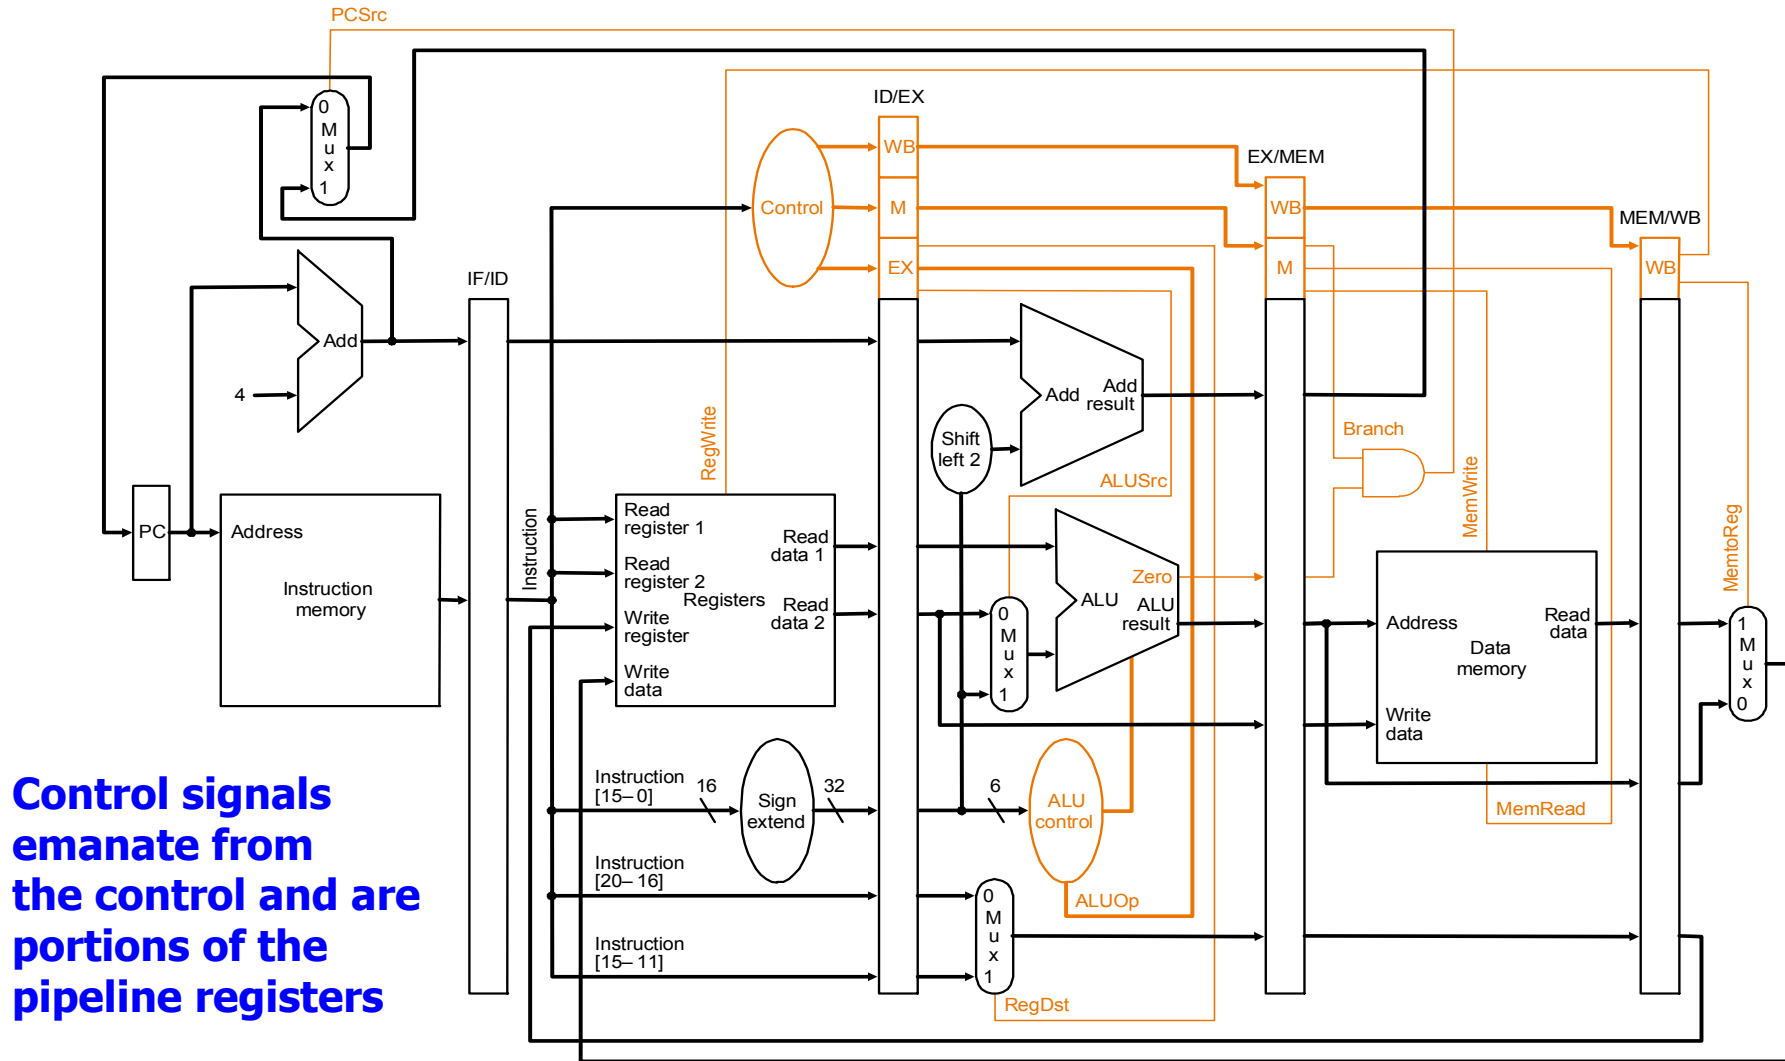
Pipeline Control Implementation

- ♦ *Pass control signals along just like the data* – extend each pipeline register to hold needed control bits for succeeding stages



- ♦ *Note: the 6-bit **funct** field of the instruction required in the EX stage to generate ALU control can be retrieved as the 6 least significant bits of the immediate field which is sign-extended and passed from the IF/ID register to the ID/EX register*

Pipelined Datapath with Control II



An Example



lw \$10, 20(\$1)

sub \$11, \$2, \$3

and \$12, \$4, \$5

or \$13, \$6, \$7

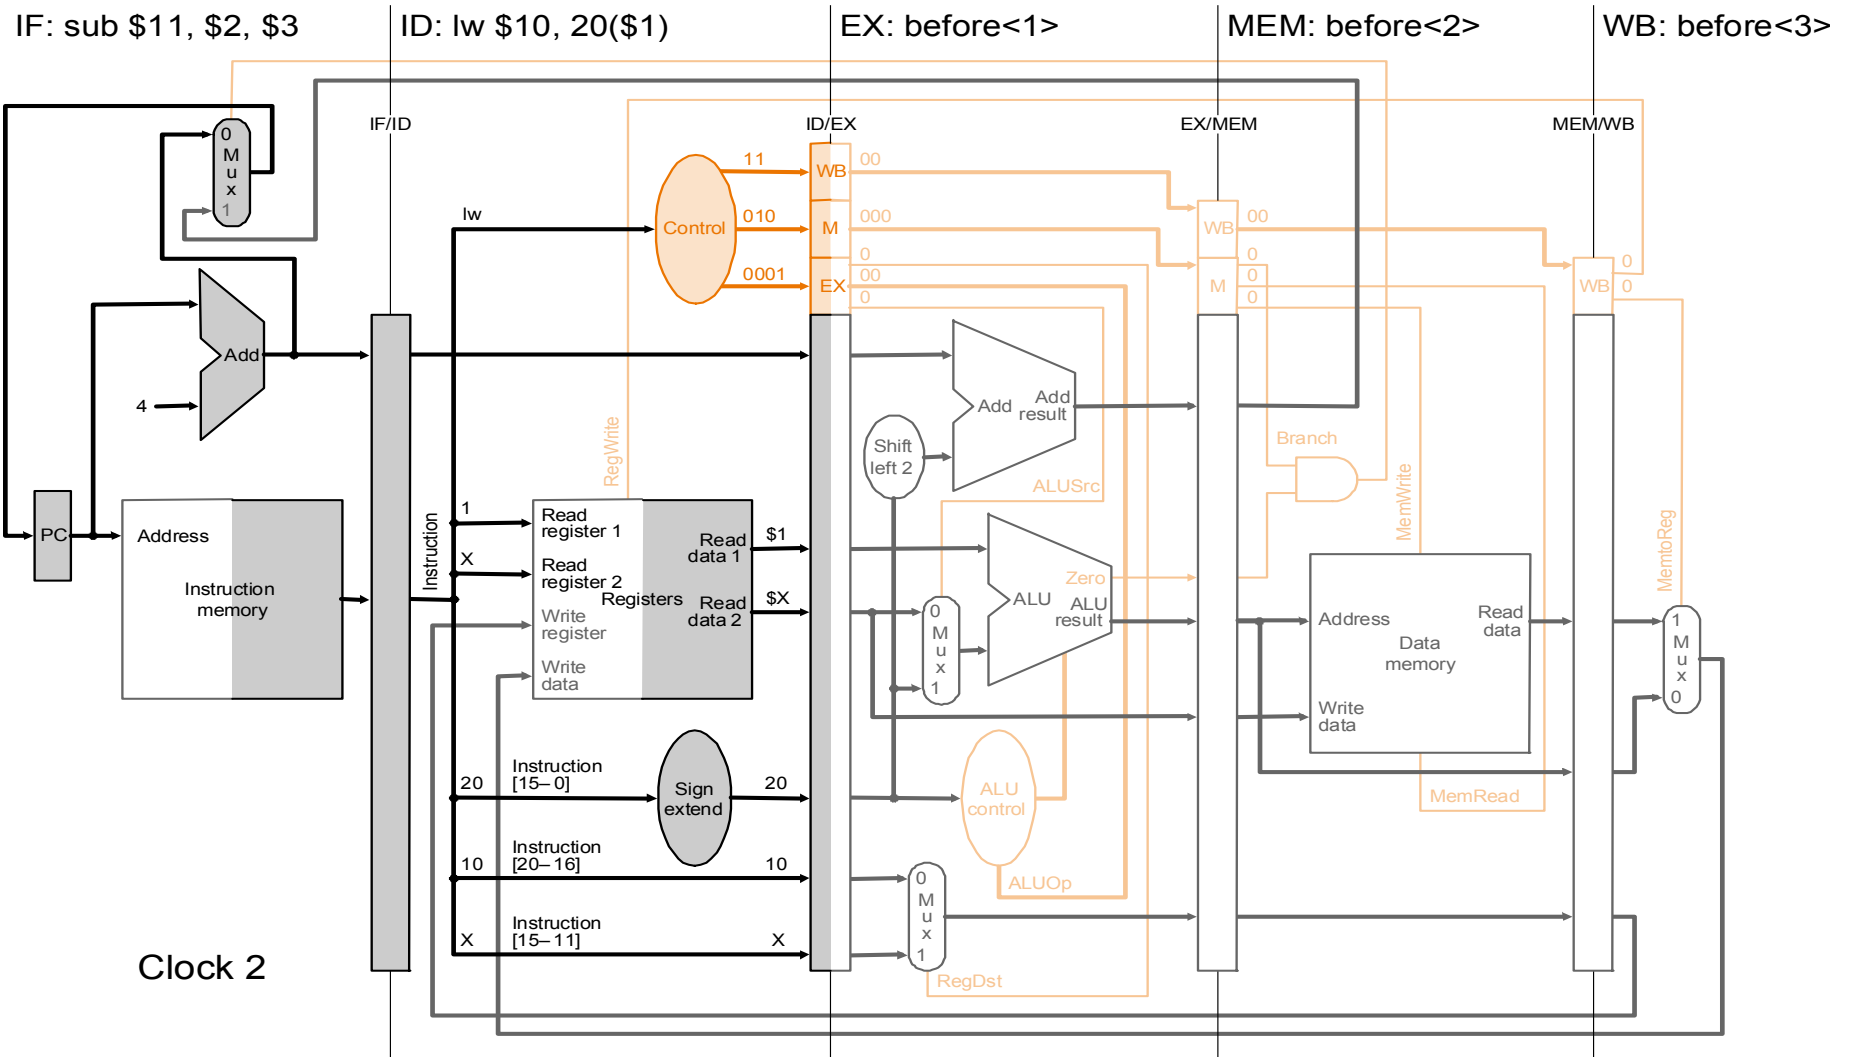
add \$14, \$8, \$9

```
lw $10, 20($1)
sub $11, $2, $3
and $12, $4, $5
or $13, $6, $7
add $14, $8, $9
```



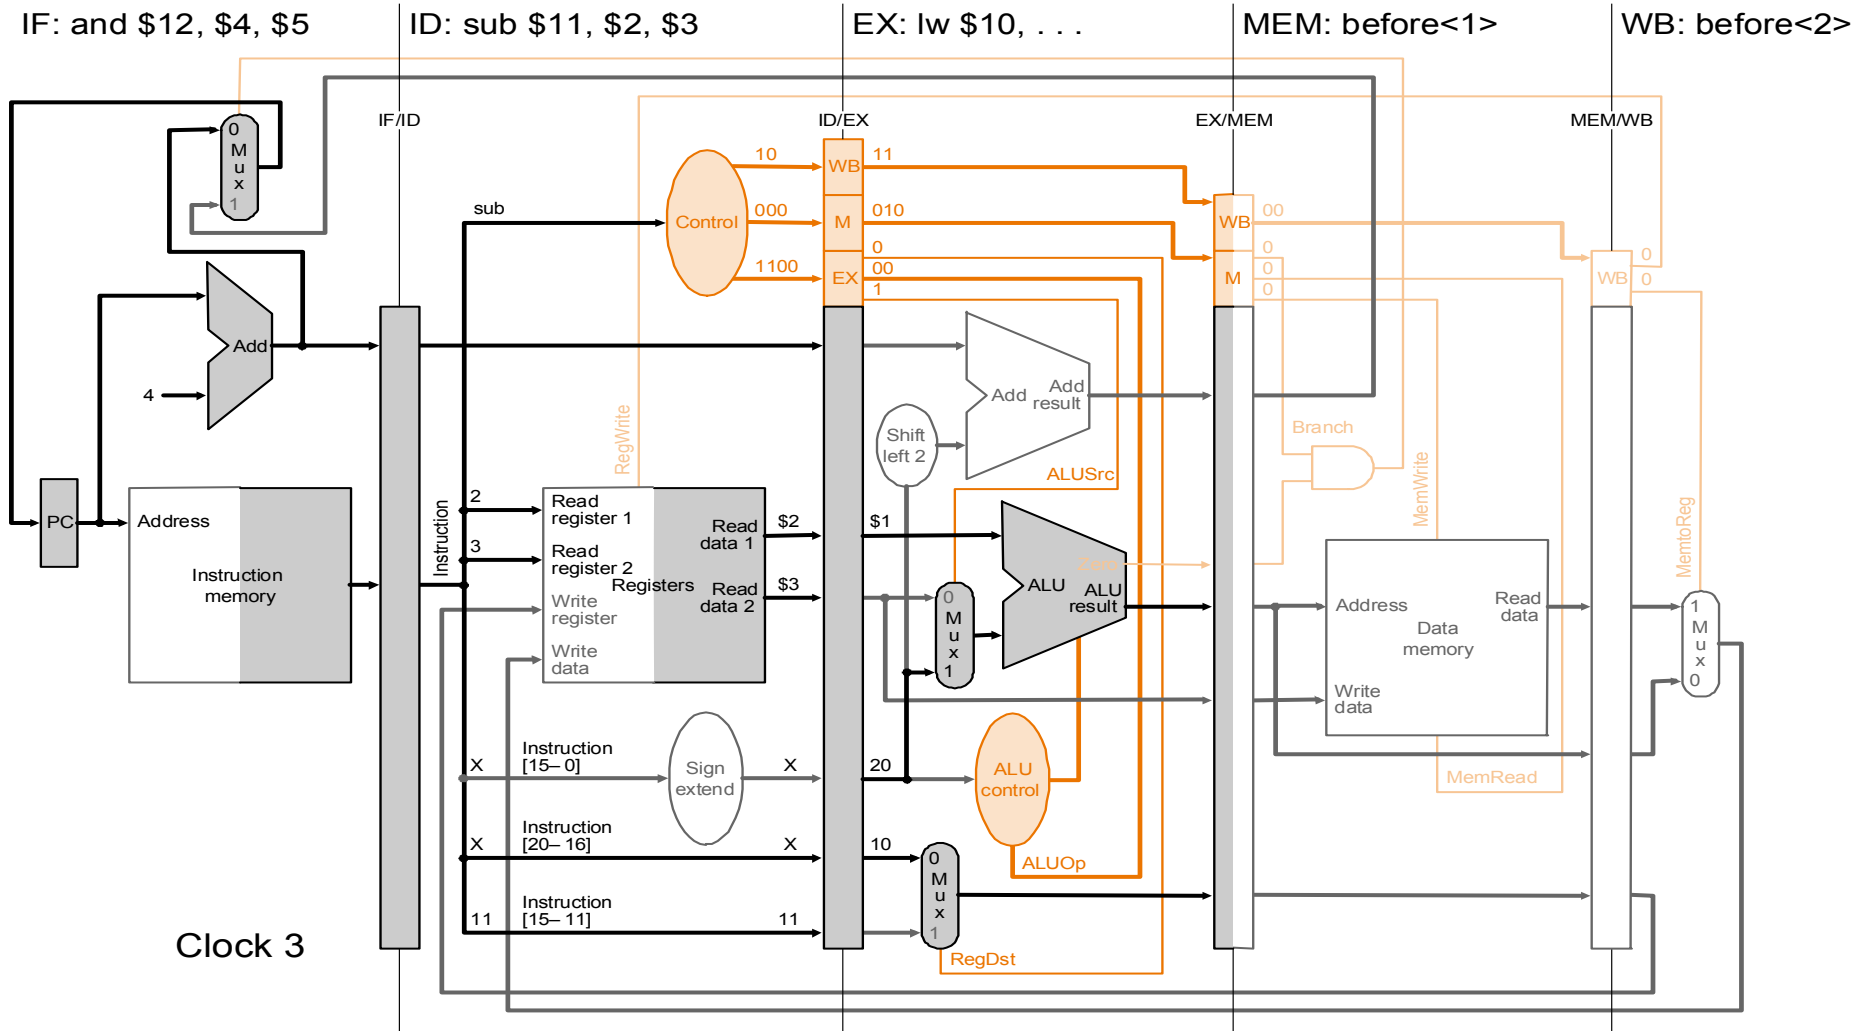
Cycle 2

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```




Cycle 3

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```




Cycle 4

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```




Cycle 5

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```




Cycle 6

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```



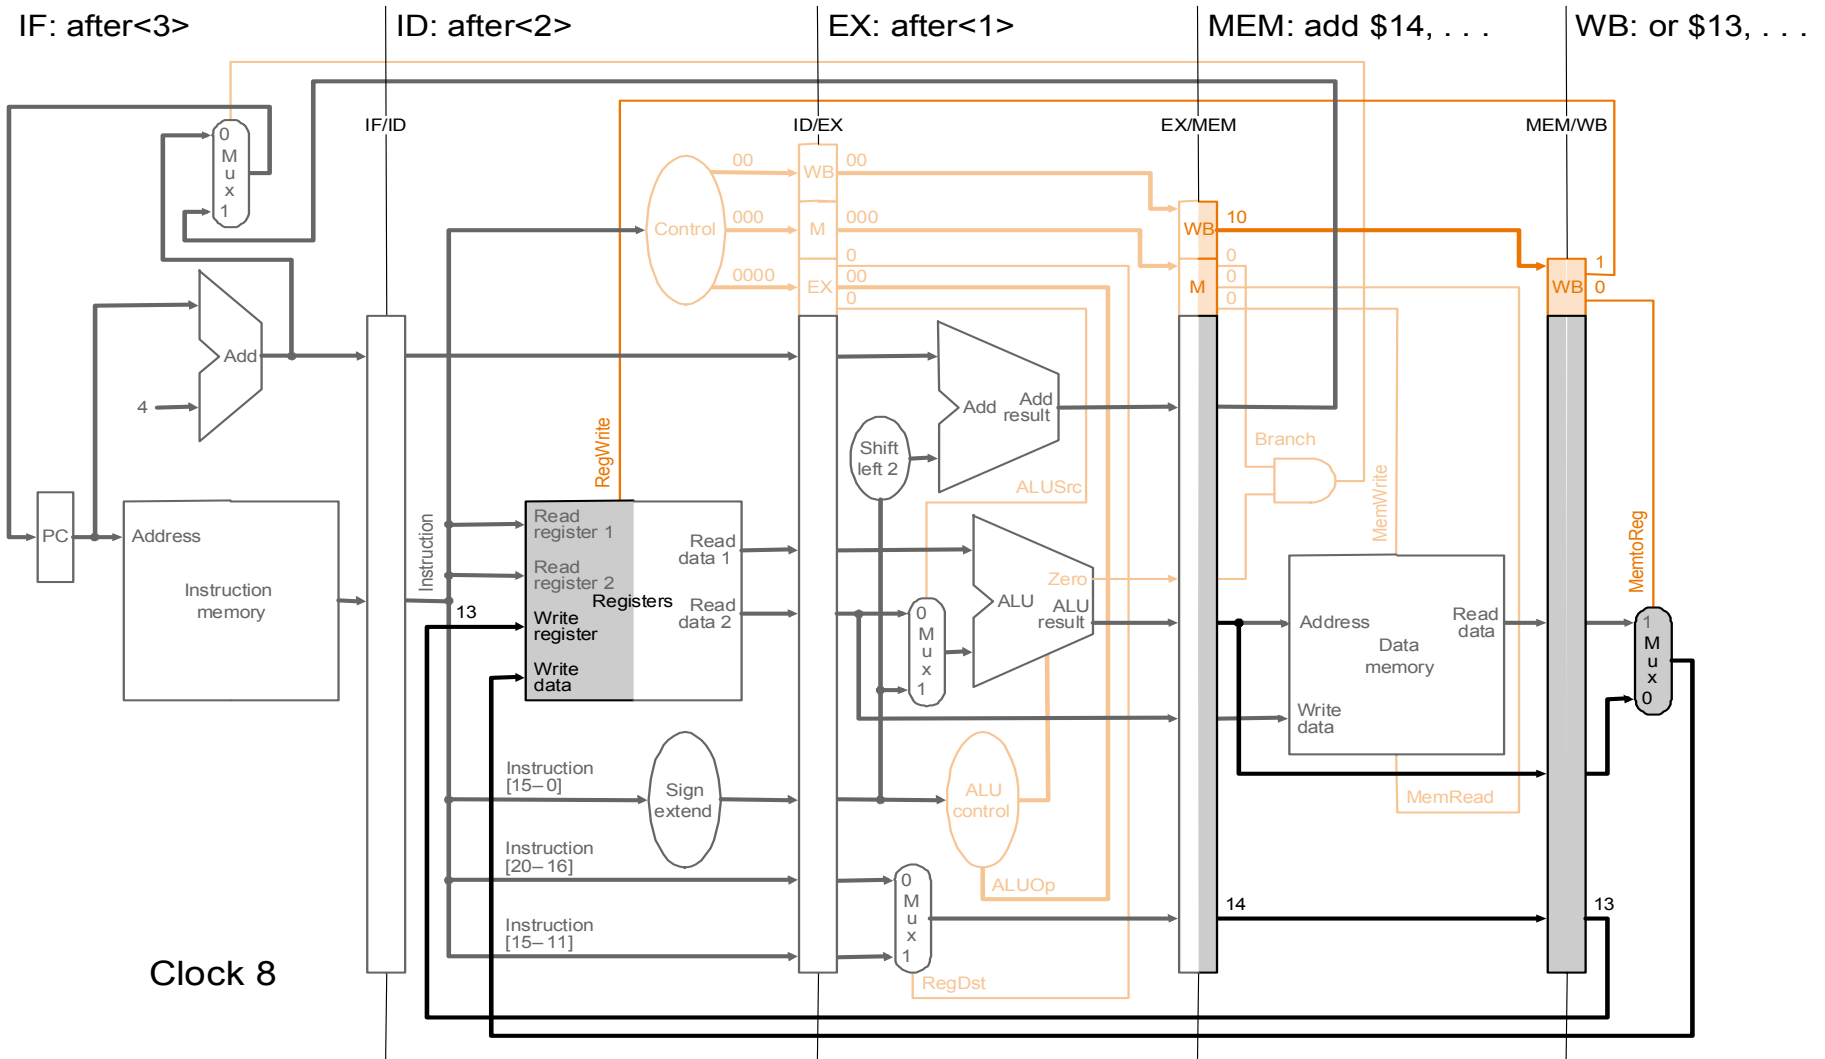
Cycle 7

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```



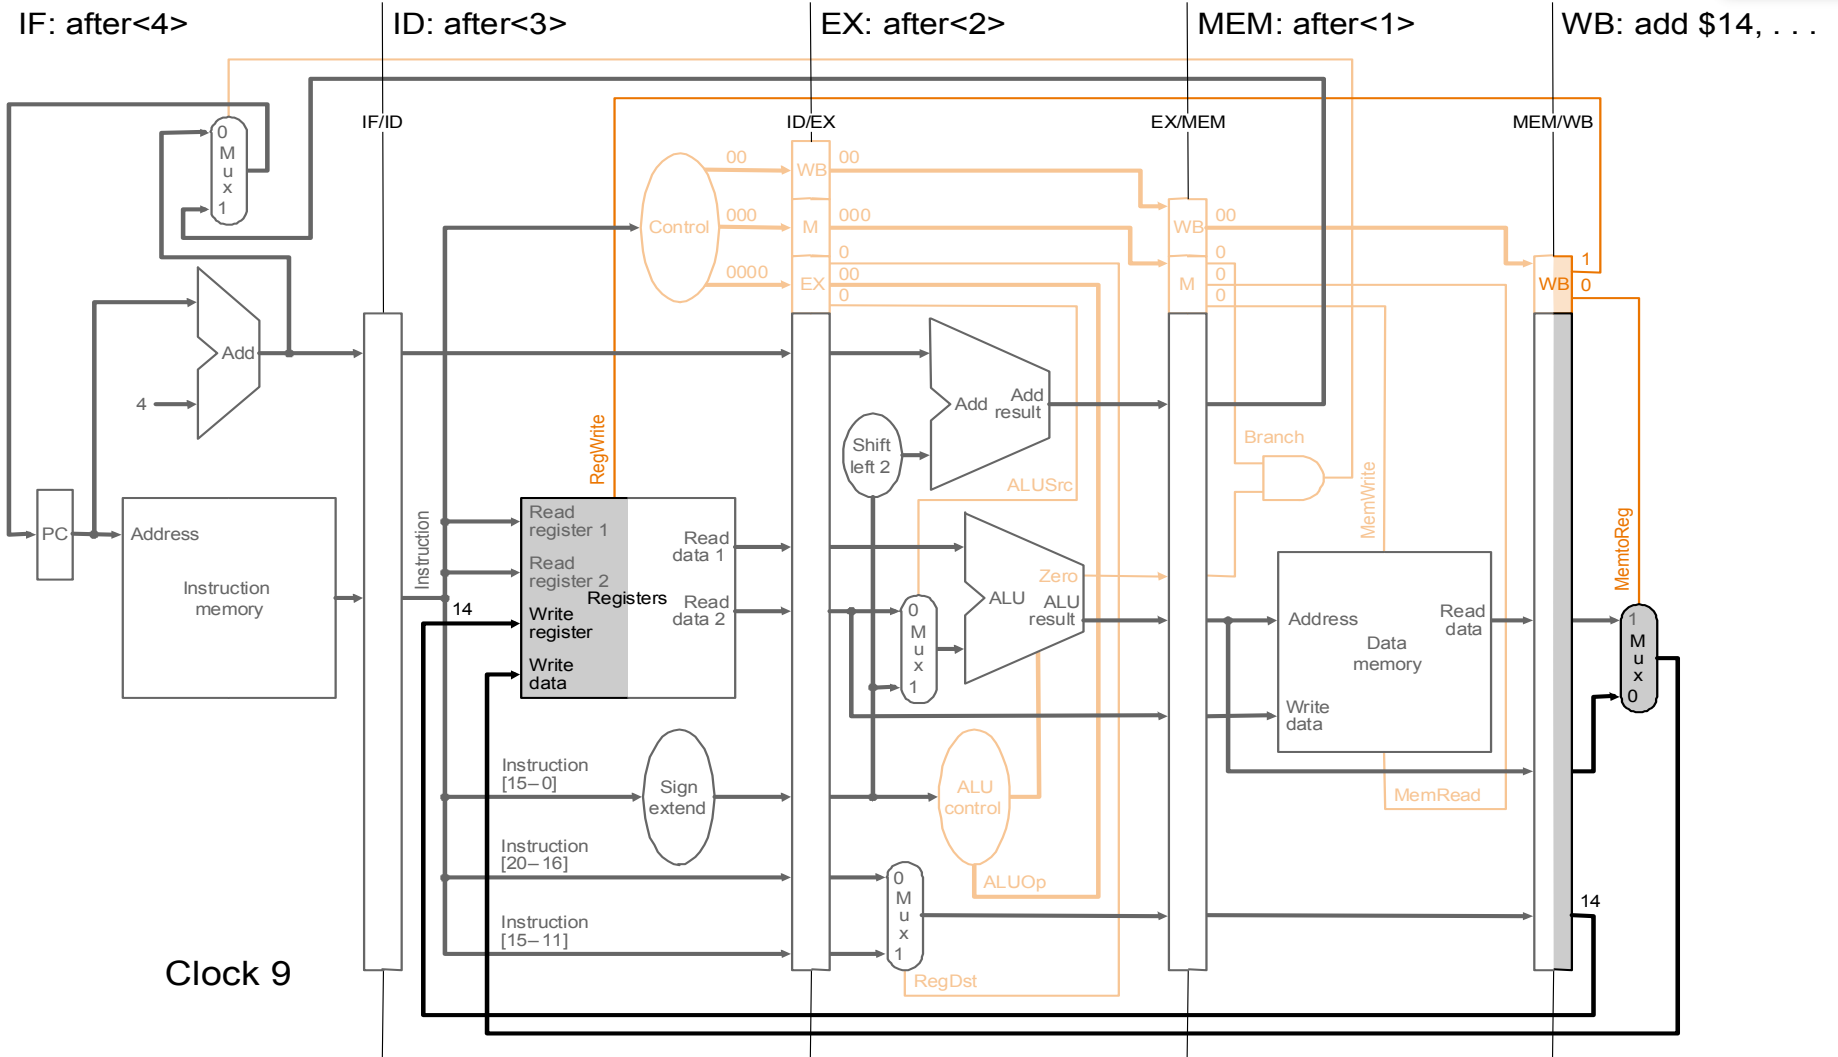
Cycle 8

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```



Cycle 9

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t5
or $t3, $t6, $t7
add $t4, $t8, $t9
```



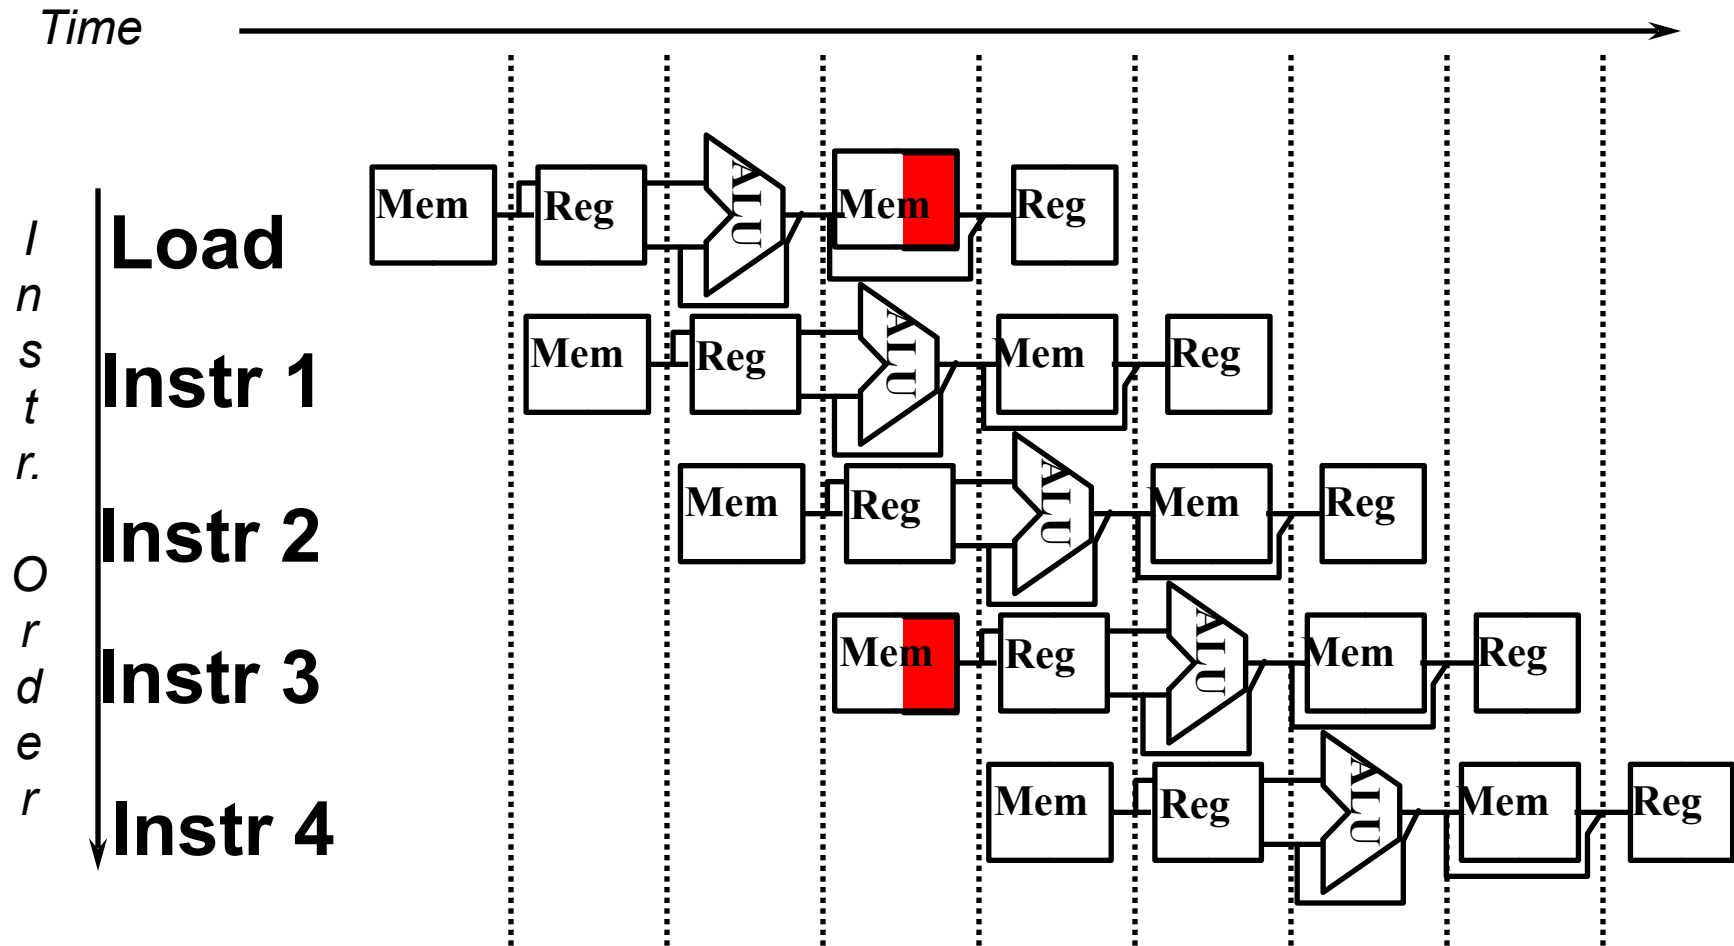
Outline

- ♦ A pipelined datapath
- ♦ Pipelined control
- ♦ Data hazards and forwarding
- ♦ Data hazards and stalls
- ♦ Branch hazards

Pipeline Hazards

- ♦ Pipeline Hazards:
 - **Structural hazards**: attempt to use the same resource in two different ways at the same time
 - **Data hazards**: attempt to use data item before ready
 - Instruction depends on result of prior instruction still in the pipeline
 - **Control hazards**: attempt to make decision before condition is evaluated
 - Branch instructions
- ♦ Can always resolve hazards by **waiting**
 - pipeline control must detect the hazard
 - take action (or delay action) to resolve hazards

Structural Hazard: Single Memory



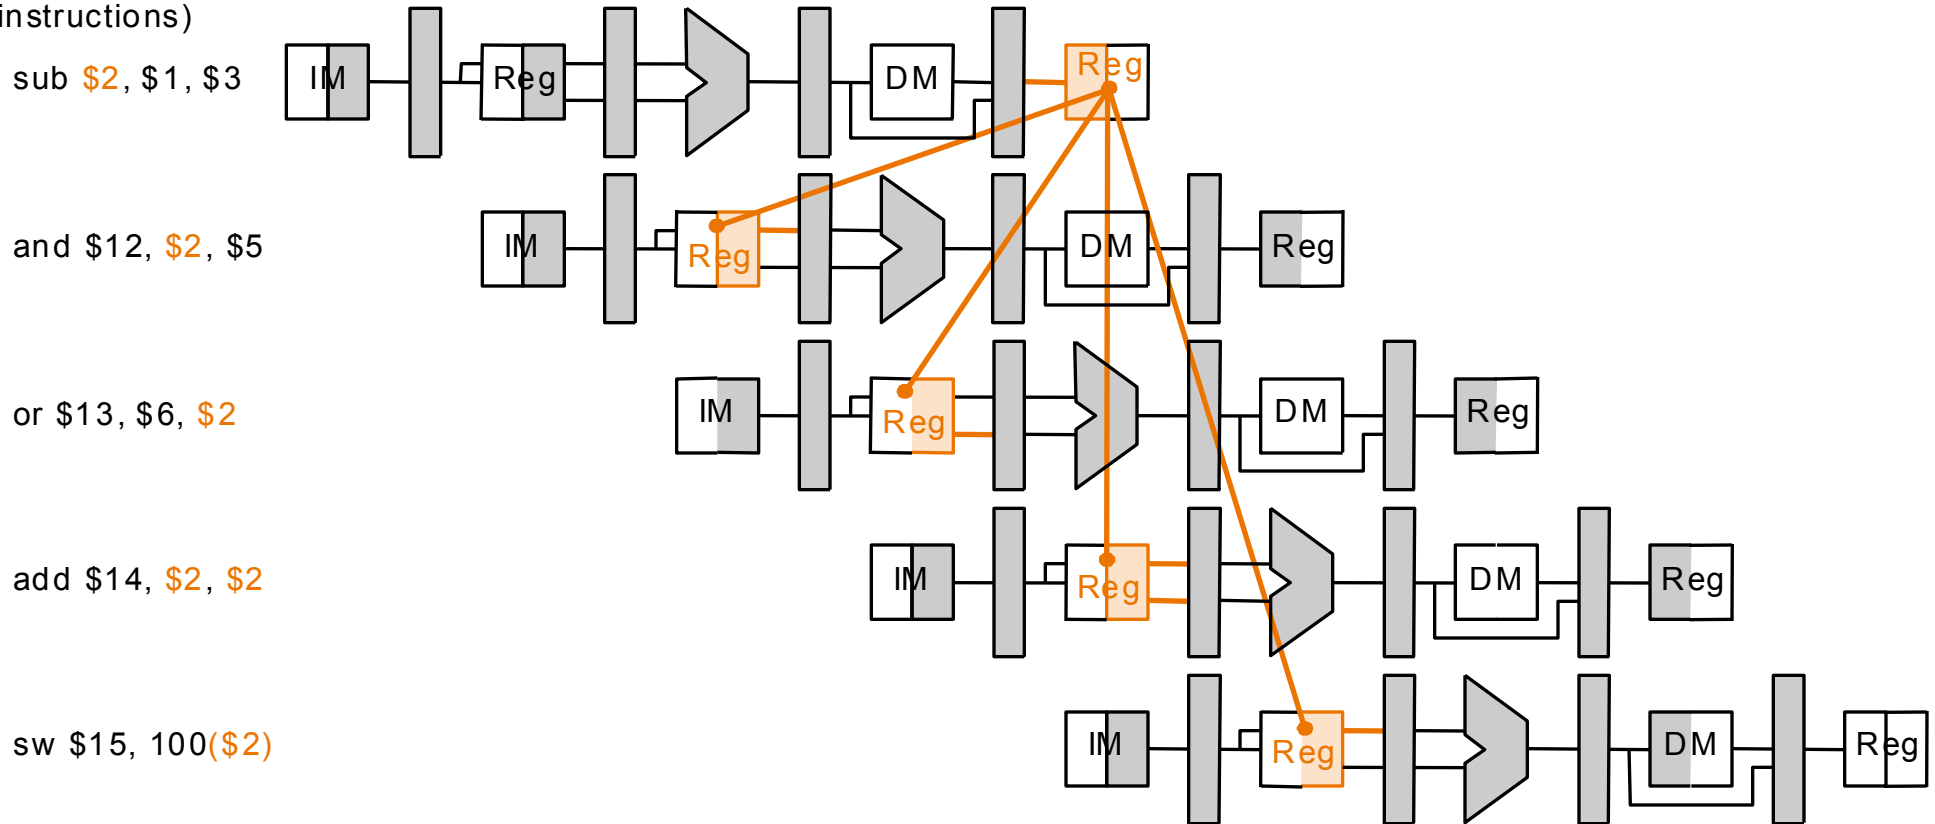
Use 2 memory: data memory and instruction memory

Data Hazards

Time (in clock cycles)

	CC 1	CC 2	CC 3	CC 4	CC 5	CC 6	CC 7	CC 8	CC 9
Value of register \$2:	10	10	10	10	10 / -20	-20	-20	-20	-20

Program
execution
order
(in instructions)

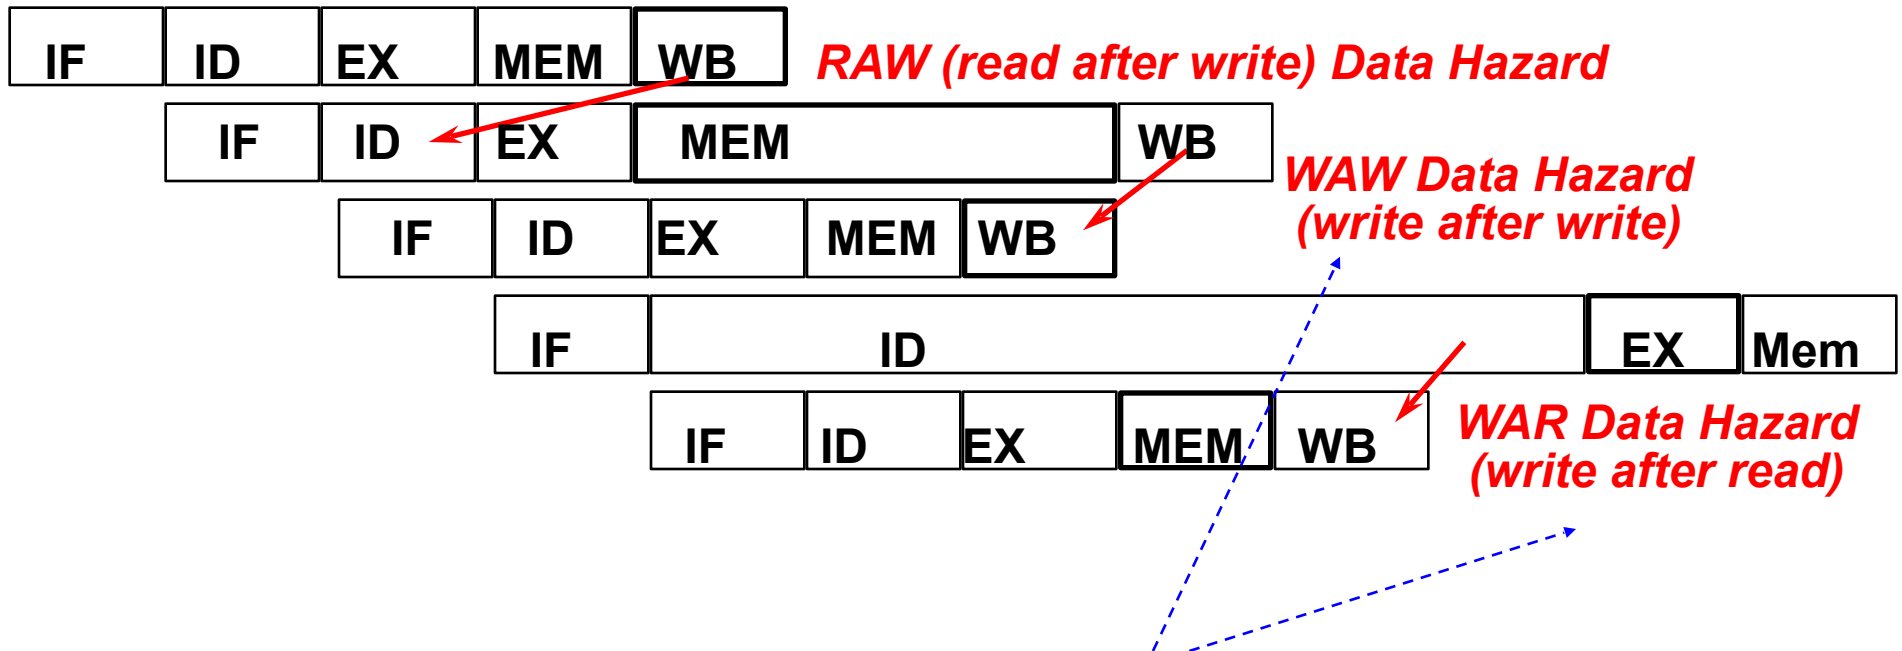


Types of Data Hazards

Three types: (inst. **i1** followed by inst. **i2**)

- ♦ **RAW (read after write):**
i2 tries to read operand before **i1** writes it
- ♦ **WAR (write after read):**
i2 tries to write operand before **i1** reads it
 - Can't happen in MIPS 5-stage pipeline because all instructions take 5 stages, and reads are always in stage 2, and writes are always in stage 5
- ♦ **WAW (write after write):**
i2 tries to write operand before **i1** writes it
 - Can't happen in MIPS 5-stage pipeline because all instructions take 5 stages, and writes are always in stage 5

Pipeline Hazards Illustrated



Cannot appear in the MIPS discussed in Chap. 4

Hardware Solution: Forwarding

- ♦ Idea: fetch “fresh” data as early as possible

- ♦ Two steps:

Step 1: **Detect** data hazard:

- Is the datum just produced required by the following inst.?

Step 2: **Forward** intermediate data to resolve hazard

- If yes, then forward the requested datum to the requesting inst. immediately.

The diagram illustrates the internal architecture of a processor, showing the flow of data and control signals through five stages: IF/ID, ID/EX, EX/MEM, MEM/WB, and WB. Key components include the Program Counter (PC), Instruction Memory, Register File, ALU, Data Memory, and various multiplexers and control logic blocks. Orange lines highlight specific control paths: RegWrite, ALUSrc, Zero, ALUOp, RegDst, Branch, MemWrite, MemRead, and MemtoReg.

Hazard Detection

sub	\$2, \$1, \$3
and	\$12, \$2, \$5
or	\$13, \$6, \$2
add	\$14, \$2, \$2
sw	\$15, 100(\$2)

♦ Hazard conditions:

1a. EX/MEM.RegisterRd = ID/EX.RegisterRs

1b. EX/MEM.RegisterRd = ID/EX.RegisterRt

2a. MEM/WB.RegisterRd = ID/EX.RegisterRs

2b. MEM/WB.RegisterRd = ID/EX.RegisterRt

- E.g., in the earlier example, first hazard between `sub $2, $1, $3` and `and $12, $2, $5` is detected when the `and` is in EX stage and the `sub` is in MEM stage because
 - EX/MEM.RegisterRd = ID/EX.RegisterRs = \$2 (1a)

♦ Whether to forward also depends on:

- if the later instruction is going to write the same register – no need to forward, even if there is register number match as in conditions above
- if the destination register of the later instruction is \$0 – in which case there is no need to forward value (\$0 is always 0 and never overwritten)

Data Forwarding

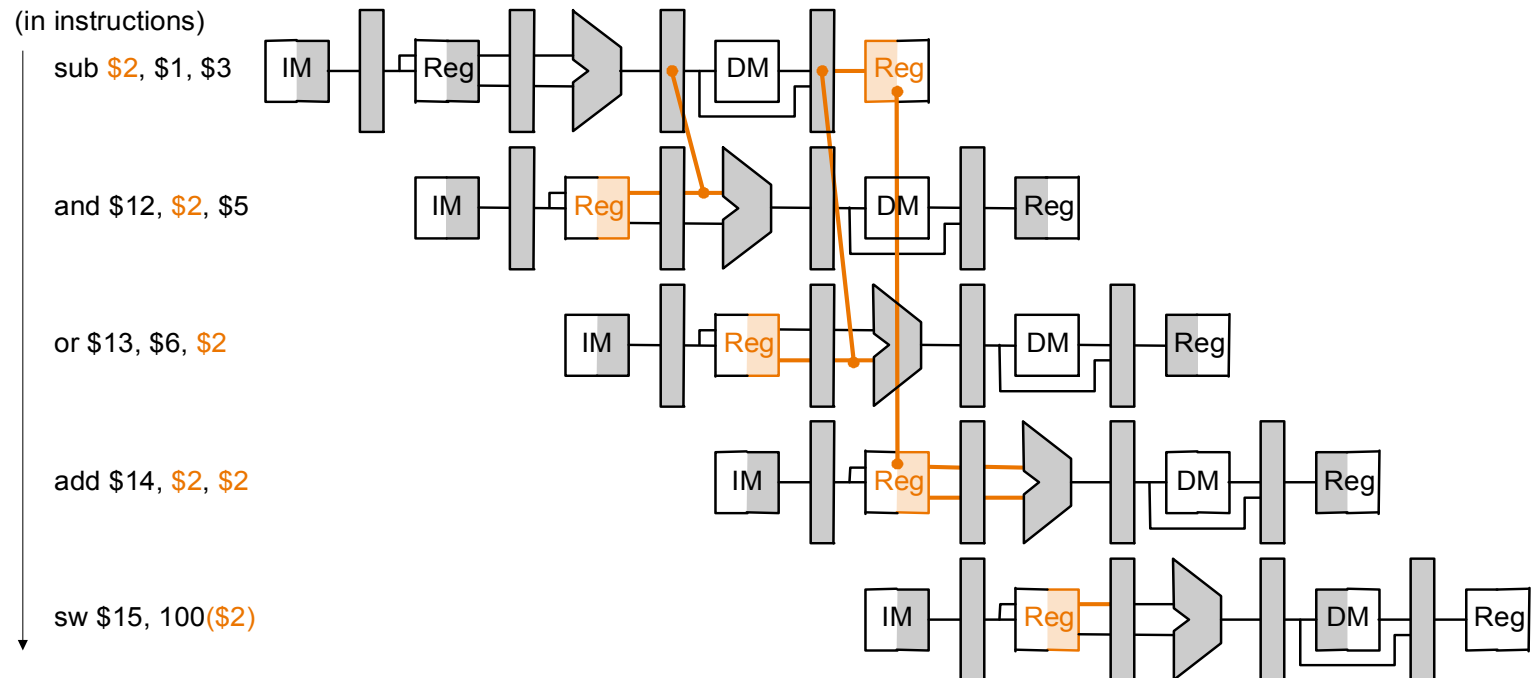
- ♦ Allow inputs to ALU not just from ID/EX, but also later pipeline registers
- ♦ Use multiplexors and control signals to choose appropriate inputs to ALU

Time (in clock cycles) →

	CC 1	CC 2	CC 3	CC 4	CC 5	CC 6	CC 7	CC 8	CC 9
Value of register \$2 :	10	10	10	10	10/-20	-20	-20	-20	-20
Value of EX/MEM :	X	X	X	-20	X	X	X	X	X
Value of MEM/WB :	X	X	X	X	-20	X	X	X	X

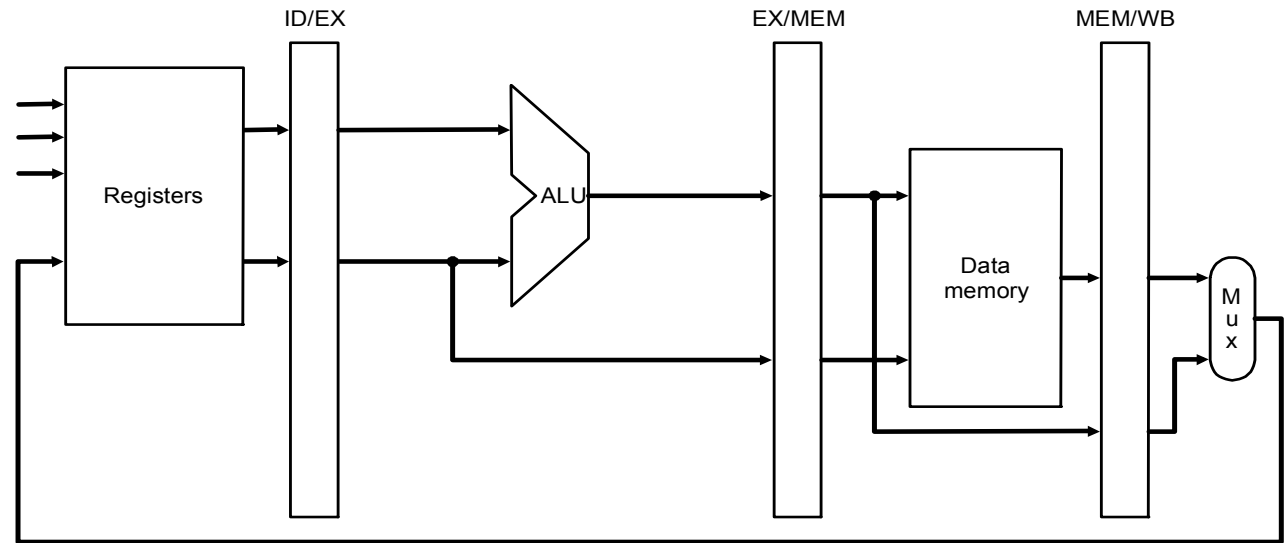
Program
execution order
(in instructions)

```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

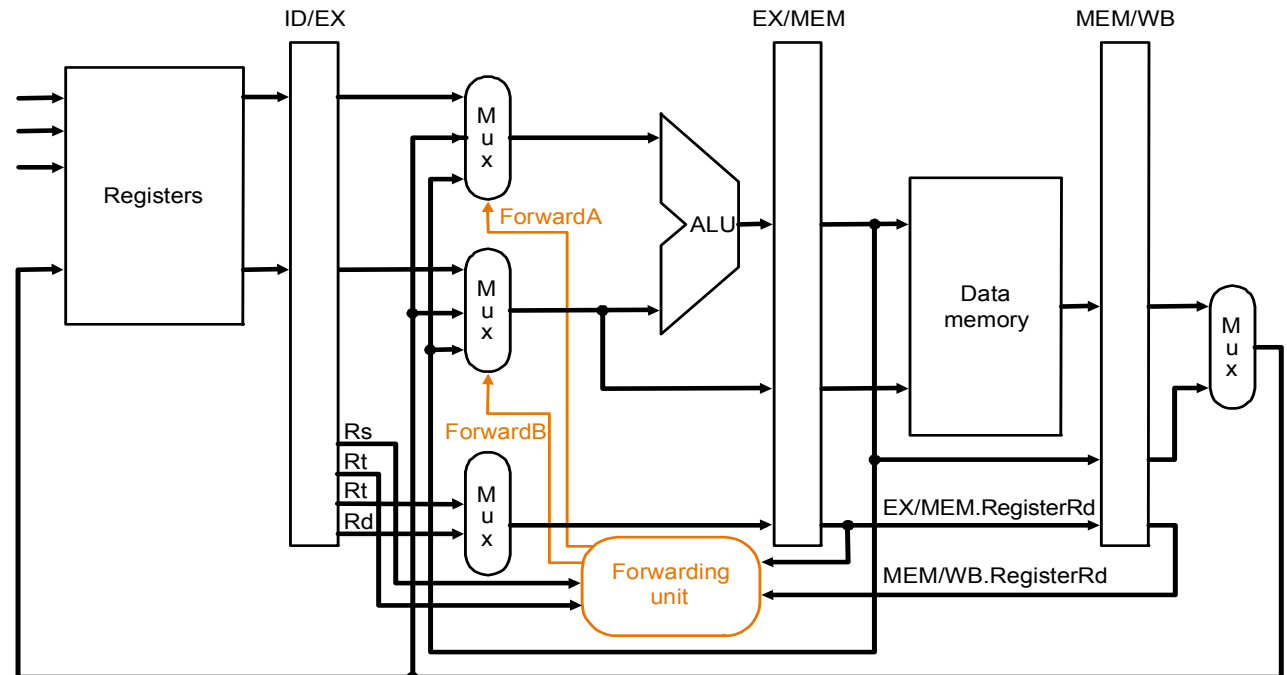


Dependencies between pipelines move forward in time

Forwarding Hardware



Datapath before adding forwarding hardware



Datapath after adding forwarding hardware

Forwarding Hardware: Multiplexor Control

Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from prior ALU result
ForwardA = 01	MEM/WB	*The first ALU operand is forwarded from data memory or an earlier ALU result
ForwardB = 00	ID/EX	The second ALU operand comes from the register file
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from prior ALU result
ForwardB = 01	MEM/WB	*The second ALU operand is forwarded from data memory or an earlier ALU result

* Depending on the selection in the rightmost multiplexor (see datapath with control diagram)

Data Hazard: Detection and Forwarding

- ♦ Forwarding unit determines multiplexor control according to the following rules:

1. *EX hazard*

```
if (      EX/MEM.RegWrite                // if there is a write...
    and ( EX/MEM.RegisterRd  $\neq$  0 )      // to a non-$0 register...
    and ( EX/MEM.RegisterRd = ID/EX.RegisterRs ) ) // which matches, then...
    ForwardA = 10
```

```
if (      EX/MEM.RegWrite                // if there is a write...
    and ( EX/MEM.RegisterRd  $\neq$  0 )      // to a non-$0 register...
    and ( EX/MEM.RegisterRd = ID/EX.RegisterRt ) ) // which matches, then...
    ForwardB = 10
```

Data Hazard: Detection and Forwarding

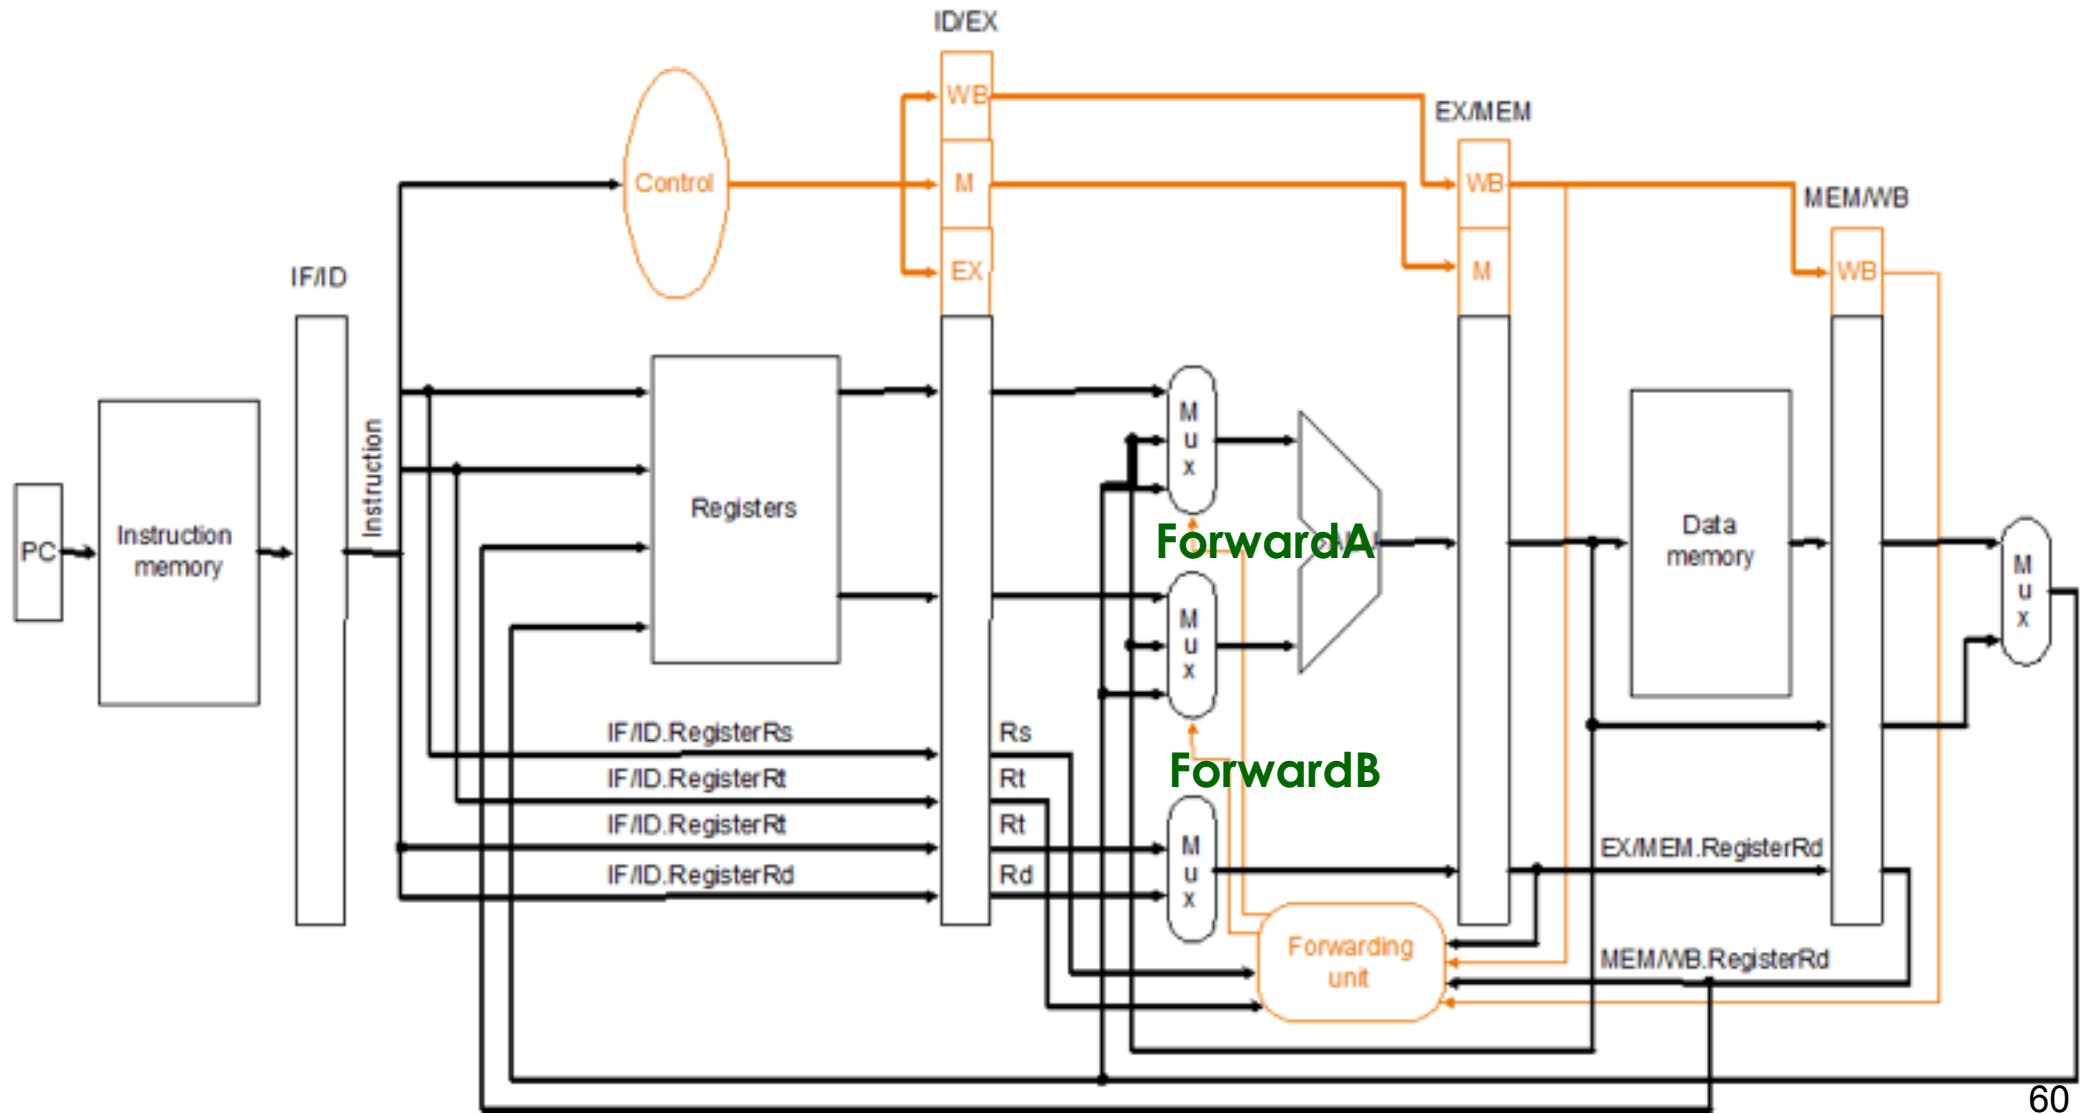
2. MEM hazard

```
if (    MEM/WB.RegWrite                // if there is a write...
    and ( MEM/WB.RegisterRd  $\neq$  0 )    // to a non-$0 register...
    and ( EX/MEM.RegisterRd  $\neq$  ID/EX.RegisterRs * ) // and not already a register match
                                                // with earlier pipeline register...
    and ( MEM/WB.RegisterRd = ID/EX.RegisterRs ) // but match with later pipeline register,
then...
    ForwardA = 01

if (    MEM/WB.RegWrite                // if there is a write...
    and ( MEM/WB.RegisterRd  $\neq$  0 )    // to a non-$0 register...
    and ( EX/MEM.RegisterRd  $\neq$  ID/EX.RegisterRt * ) // and not already a register match
                                                // with earlier pipeline register...
    and ( MEM/WB.RegisterRd = ID/EX.RegisterRt ) // but match with later pipeline register,
then...
    ForwardB = 01
```

* This check is necessary, e.g., for sequences such as `add $1, $1, $2; add $1, $1, $3; add $1, $1, $4`, where an earlier pipeline (EX/MEM) register has more recent data

Pipeline with Forwarding



Cycle 3

sub \$2, \$1, \$3
and \$4, \$2, \$5
or \$4, \$4, \$2
add \$9, \$4, \$2

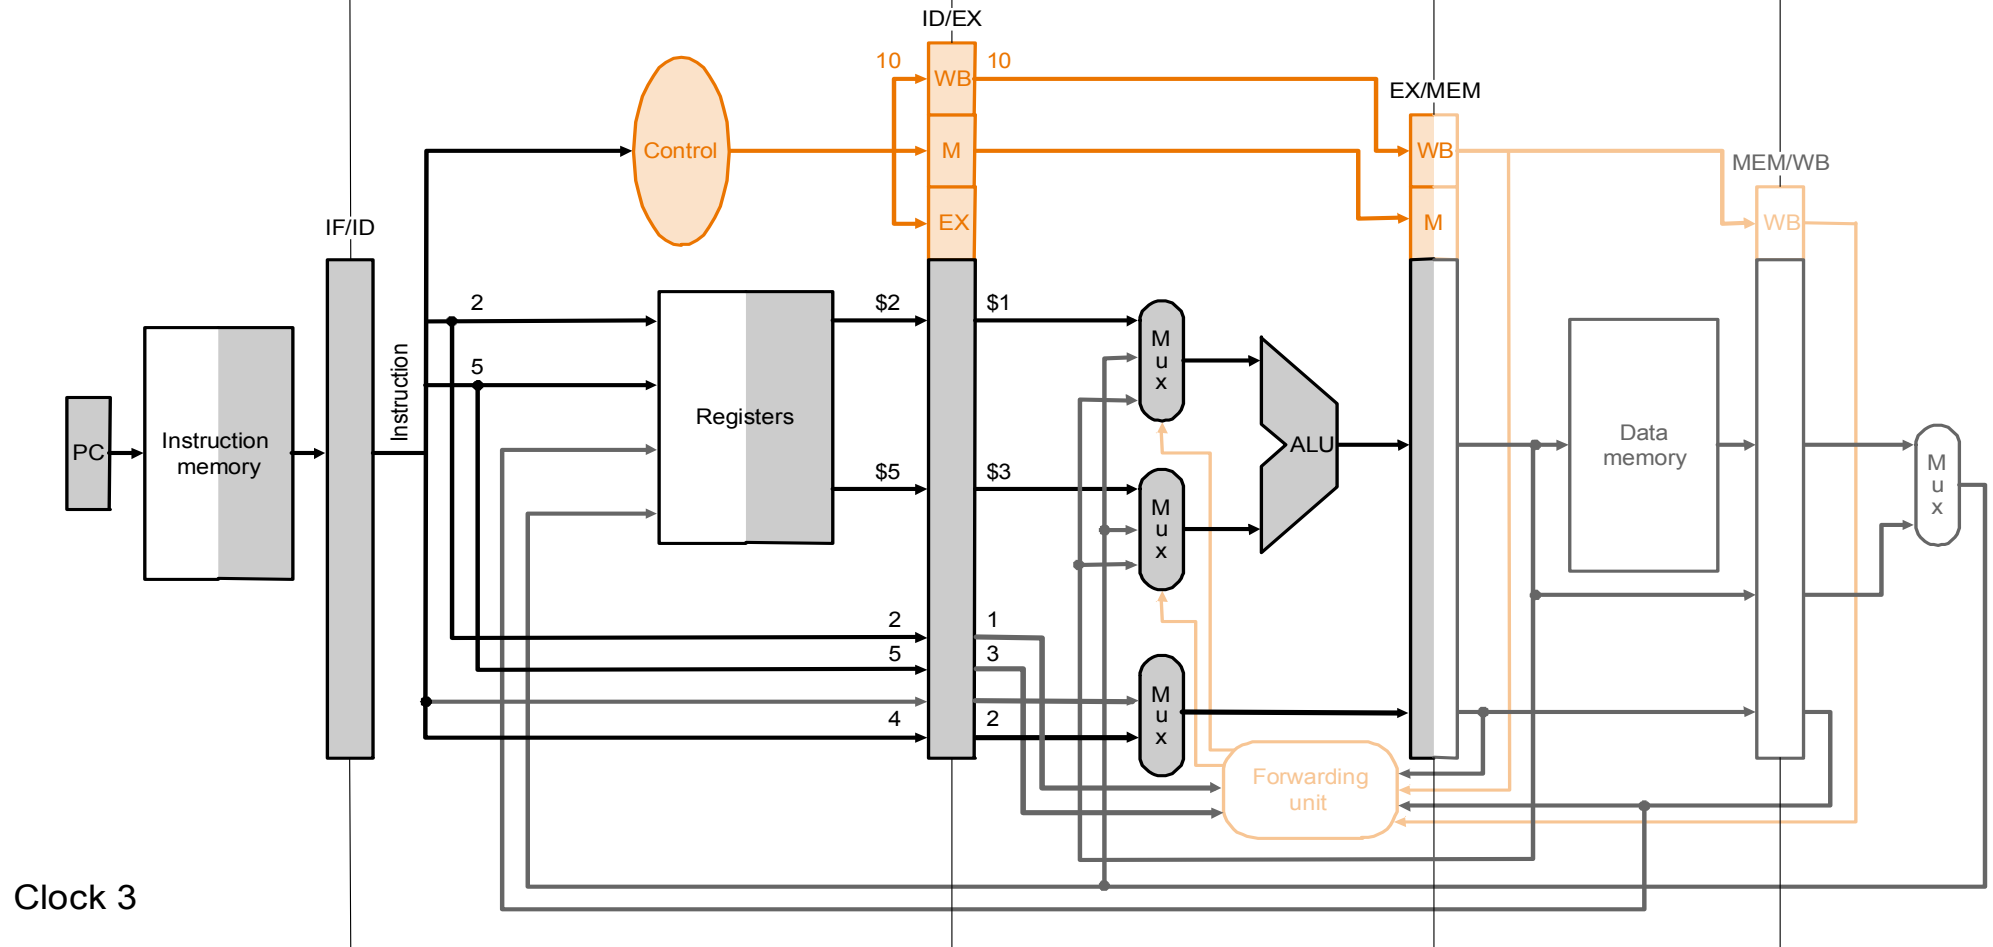
or \$4, \$4, \$2

and \$4, \$2, \$5

sub \$2, \$1, \$3

before<1>

before<2>



Cycle 4

sub \$2, \$1, \$3
and \$4, \$2, \$5
or \$4, \$4, \$2
add \$9, \$4, \$2


add \$9, \$4, \$2

or \$4, \$4, \$2

and \$4, \$2, \$5

sub \$2, ...

before<1>



Clock 4

Cycle 5

sub \$2, \$1, \$3
and \$4, \$2, \$5
or \$4, \$4, \$2
add \$9, \$4, \$2


after<1>

add \$9, \$4, \$2

or \$4, \$4, \$2

and \$4, ...


sub \$2, ...



Clock 5

Cycle 6

sub \$2, \$1, \$3
and \$4, \$2, \$5
or \$4, \$4, \$2
add \$9, \$4, \$2



Questions?



- ♦ The data path is incomplete (in Slide 59) as we cannot compute the target address for `lw` and `sw`.
- ♦ How about `lw` followed by `sw`?

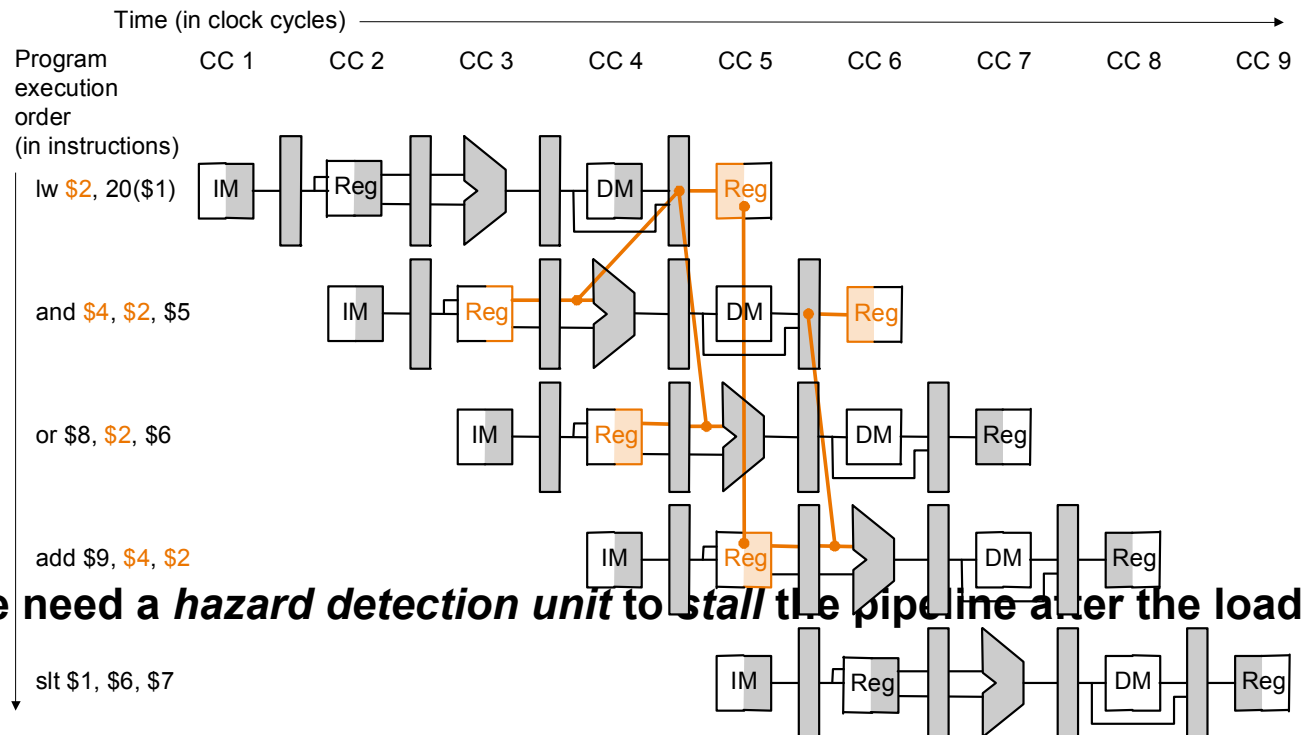
Data Hazards and Stalls

- ♦ Load word can still cause a hazard:
 - an instruction tries to read a register following a load instruction that writes to the same register

```
lw    $2, 20($1)
and   $4, $2, $5
or    $8, $2, $6
add   $9, $4, $2
slt   $1, $6, $7
```

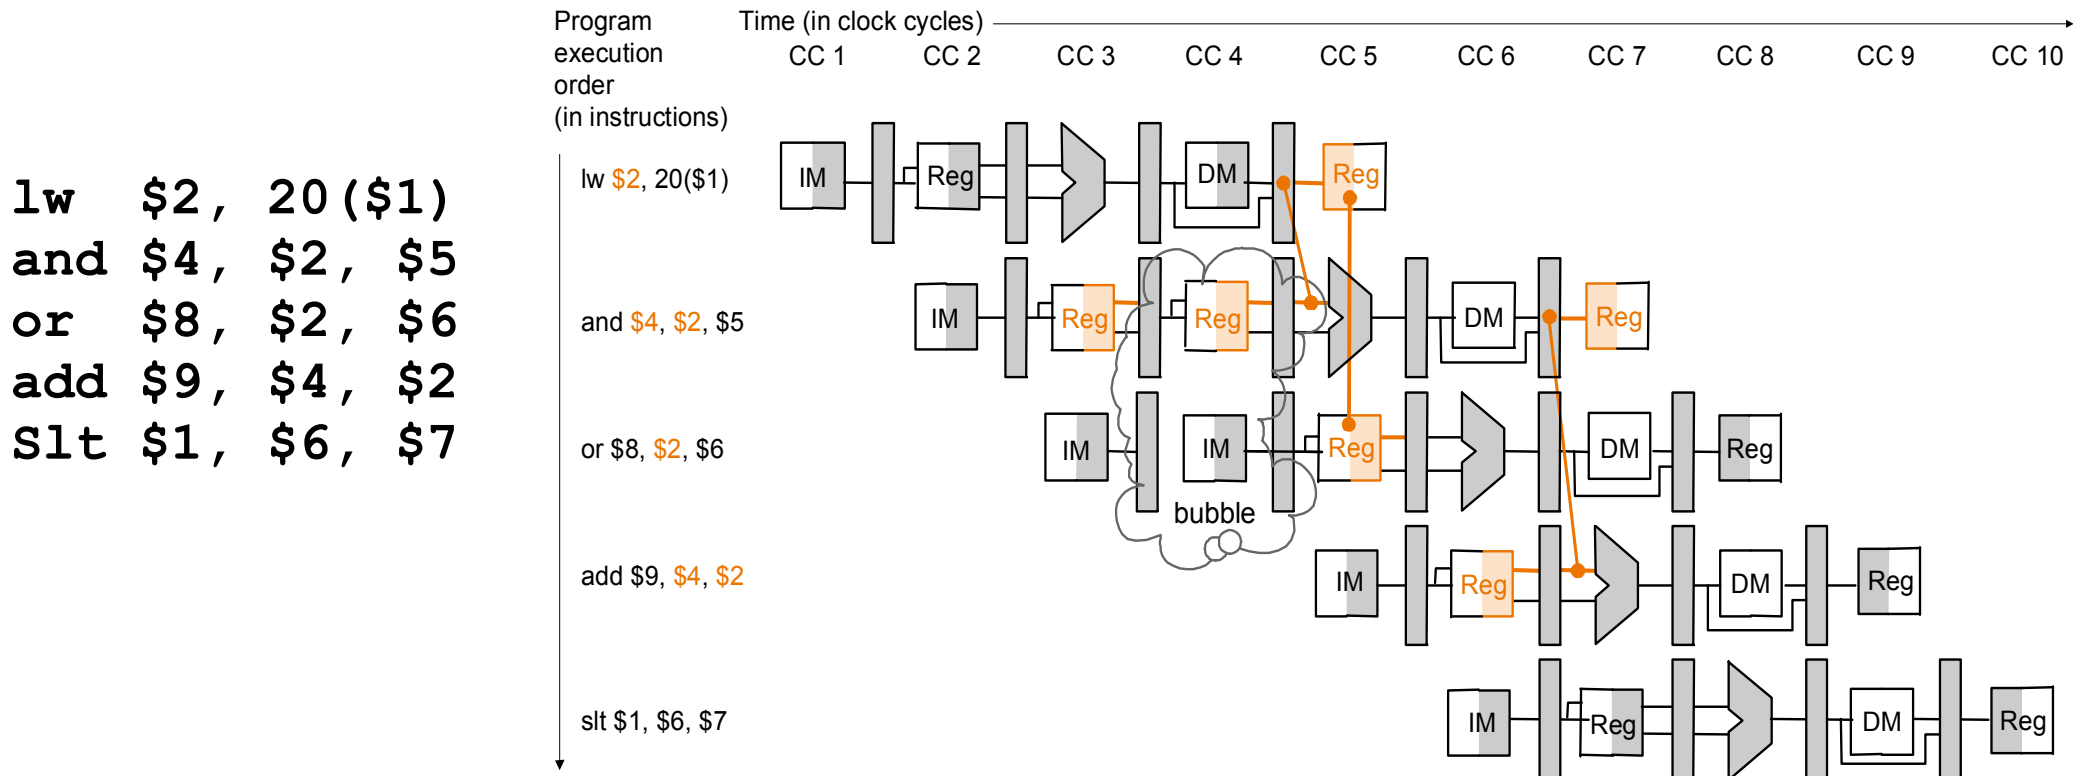
As even a pipeline dependency goes backward in time forwarding will not solve the hazard

therefore, we need a hazard detection unit to stall the pipeline after the load instruction



Stalling Resolves a Hazard

- Same instruction sequence as before for which forwarding by itself could not resolve the hazard:



Hazard detection unit inserts a 1-cycle bubble in the pipeline, after which all pipeline register dependencies go forward so then the forwarding unit can handle them and there are no more hazards

Hazard Detection Logic to Stall

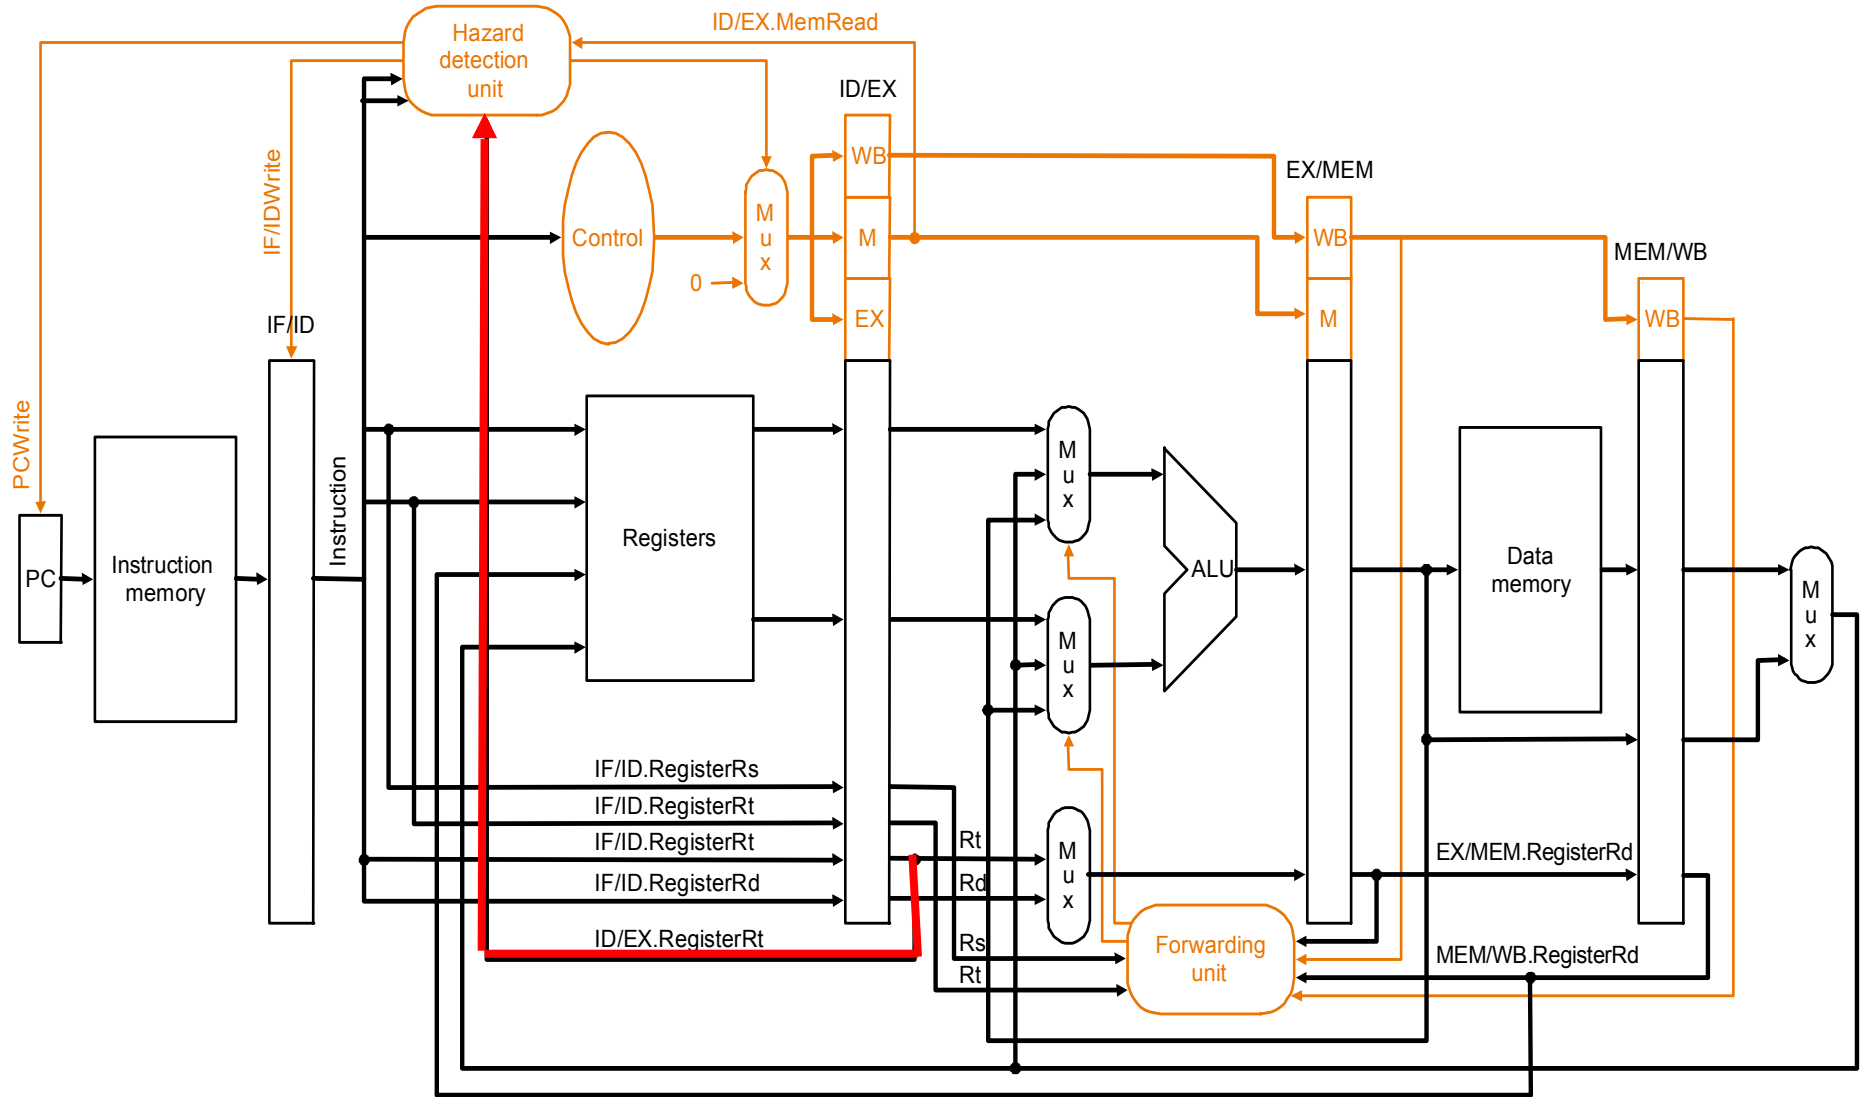
- ♦ Hazard detection unit implements the following check if to stall

```
if ( ID/EX.MemRead           // if the instruction in the EX stage is a load...
    and ( ( ID/EX.RegisterRt = IF/ID.RegisterRs )      // and the destination register
        or ( ID/EX.RegisterRt = IF/ID.RegisterRt ) ) ) // matches either source register
    // of the instruction in the ID stage, then...
    stall the pipeline
```

Mechanics of Stalling

- ♦ If the check to stall verifies, then the *pipeline needs to stall only 1 clock cycle* after the load as after that the forwarding unit can resolve the dependency
- ♦ How the hardware stalls the pipeline 1 cycle:
 - **does NOT let the IF/ID register change (disable write!)** – this will cause the instruction in the ID stage to repeat, i.e., *stall*
 - therefore, the instruction, just behind, in the IF stage must be stalled as well – so hardware **does NOT let the PC change (disable write!)** – this will cause the instruction in the IF stage to repeat, i.e., *stall*
 - **changes all the EX, MEM and WB control fields in the ID/EX pipeline register to 0**, so effectively the instruction just behind the load becomes a nop – a **bubble** is said to have been inserted into the pipeline

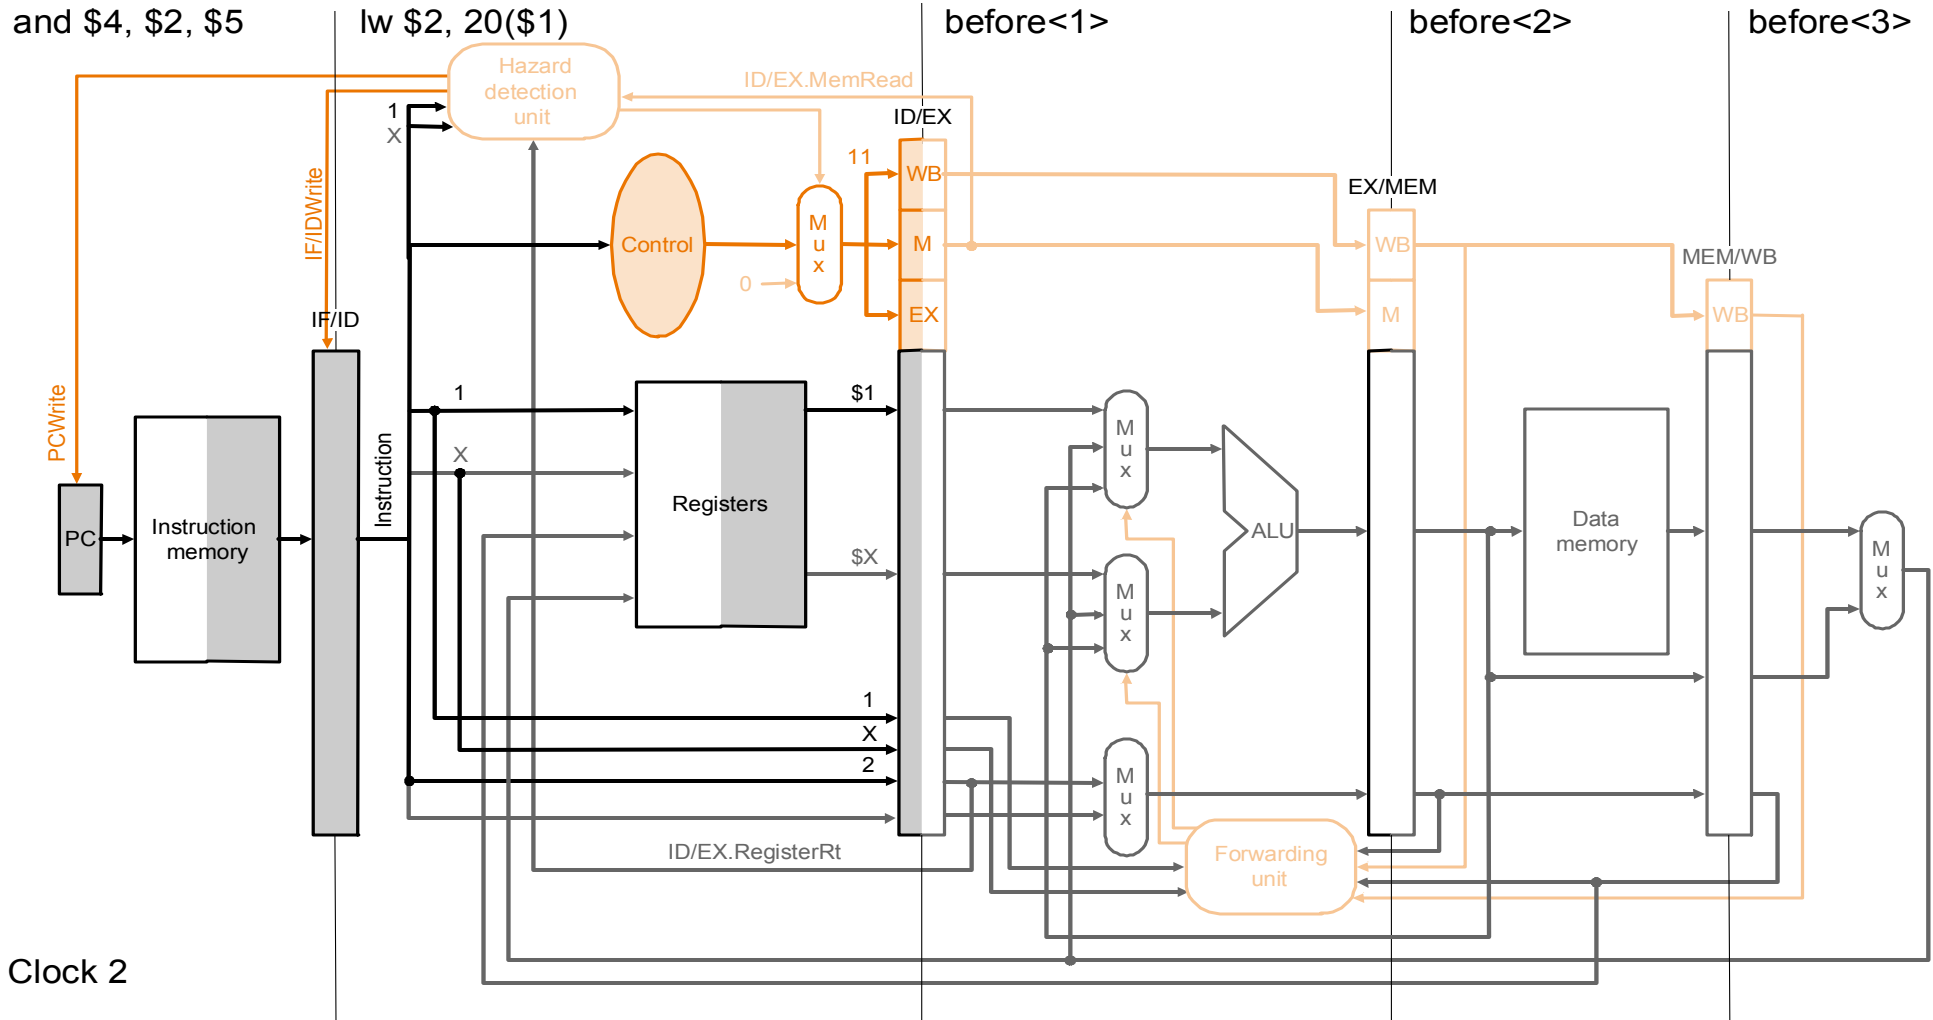
Adding Hazard Detection Unit



Datapath with forwarding hardware, the hazard detection unit and controls wires – certain details, e.g., branching hardware are omitted to simplify the drawing


Cycle 2

```
lw    $2, 20($1)
and    $4, $2, $5
or     $4, $4, $2
add    $9, $4, $2
```



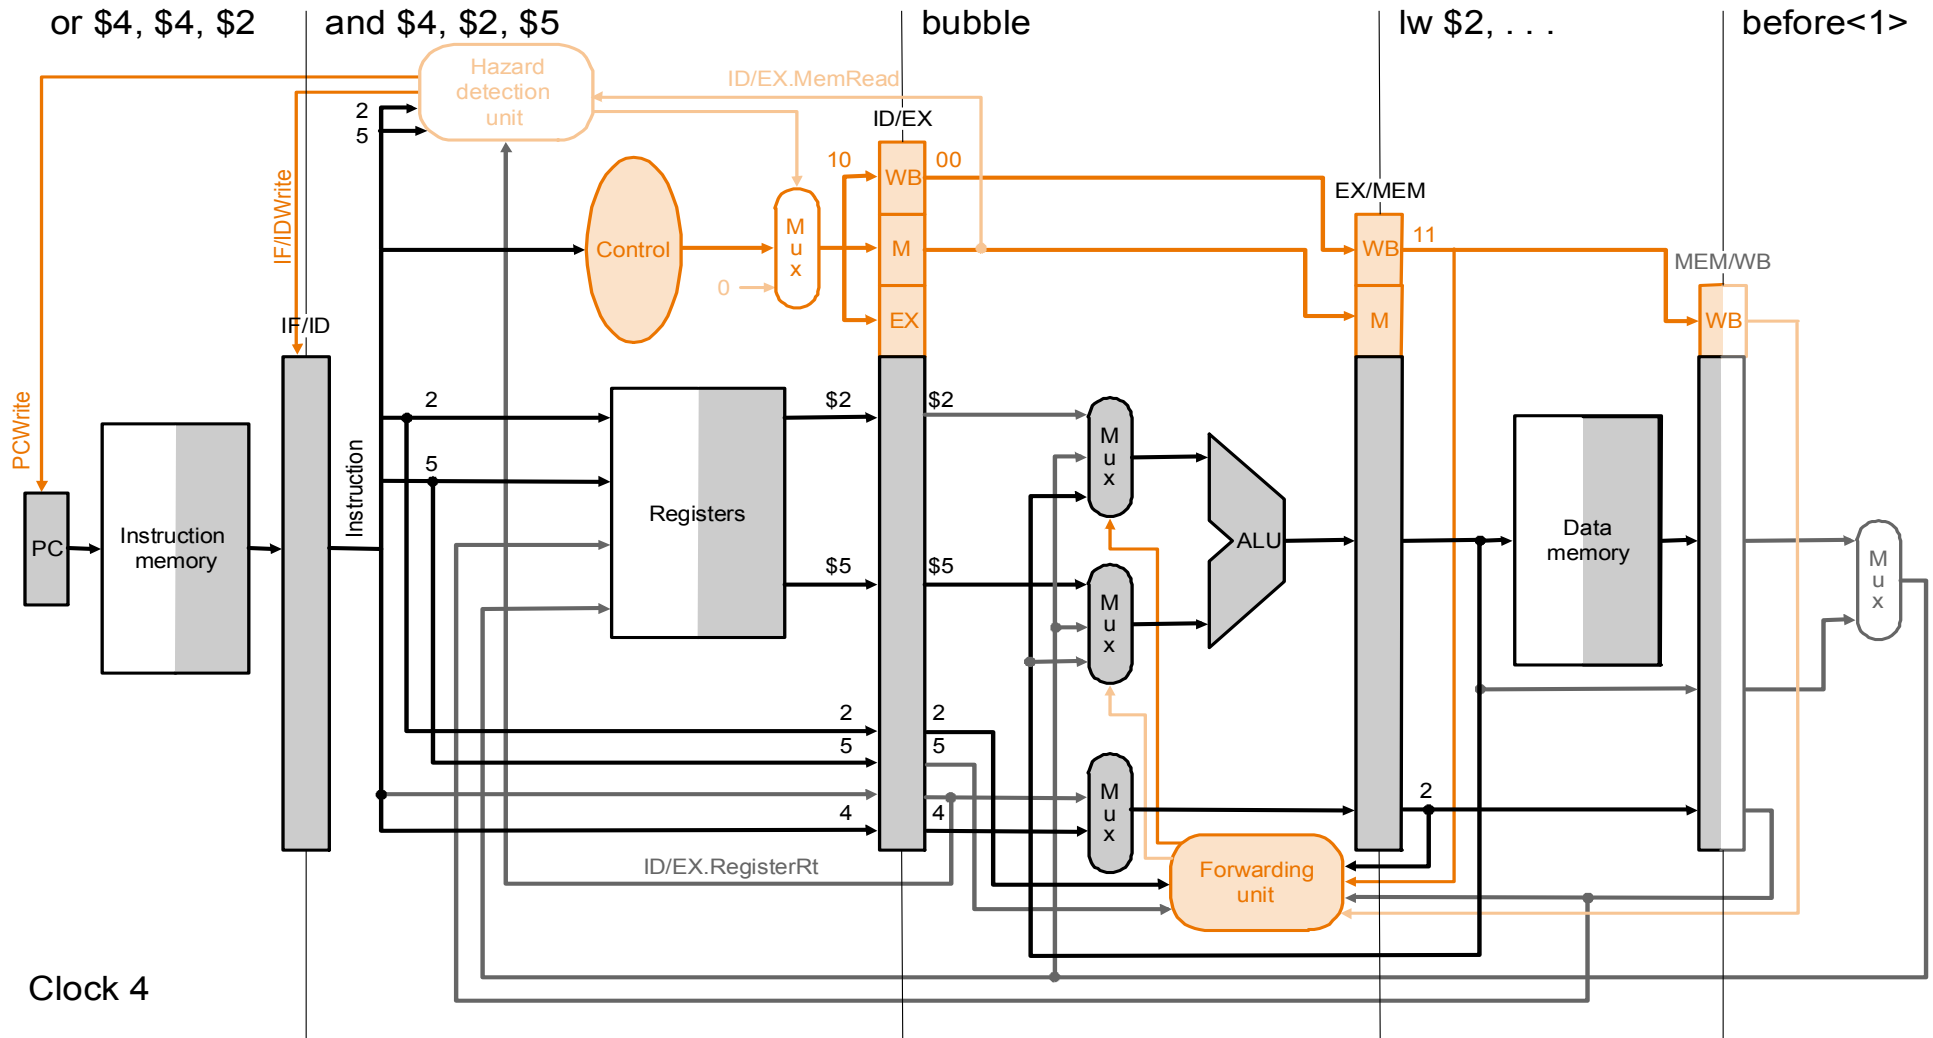
Cycle 3

```
lw    $2, 20($1)
and    $4, $2, $5
or     $4, $4, $2
add    $9, $4, $2
```



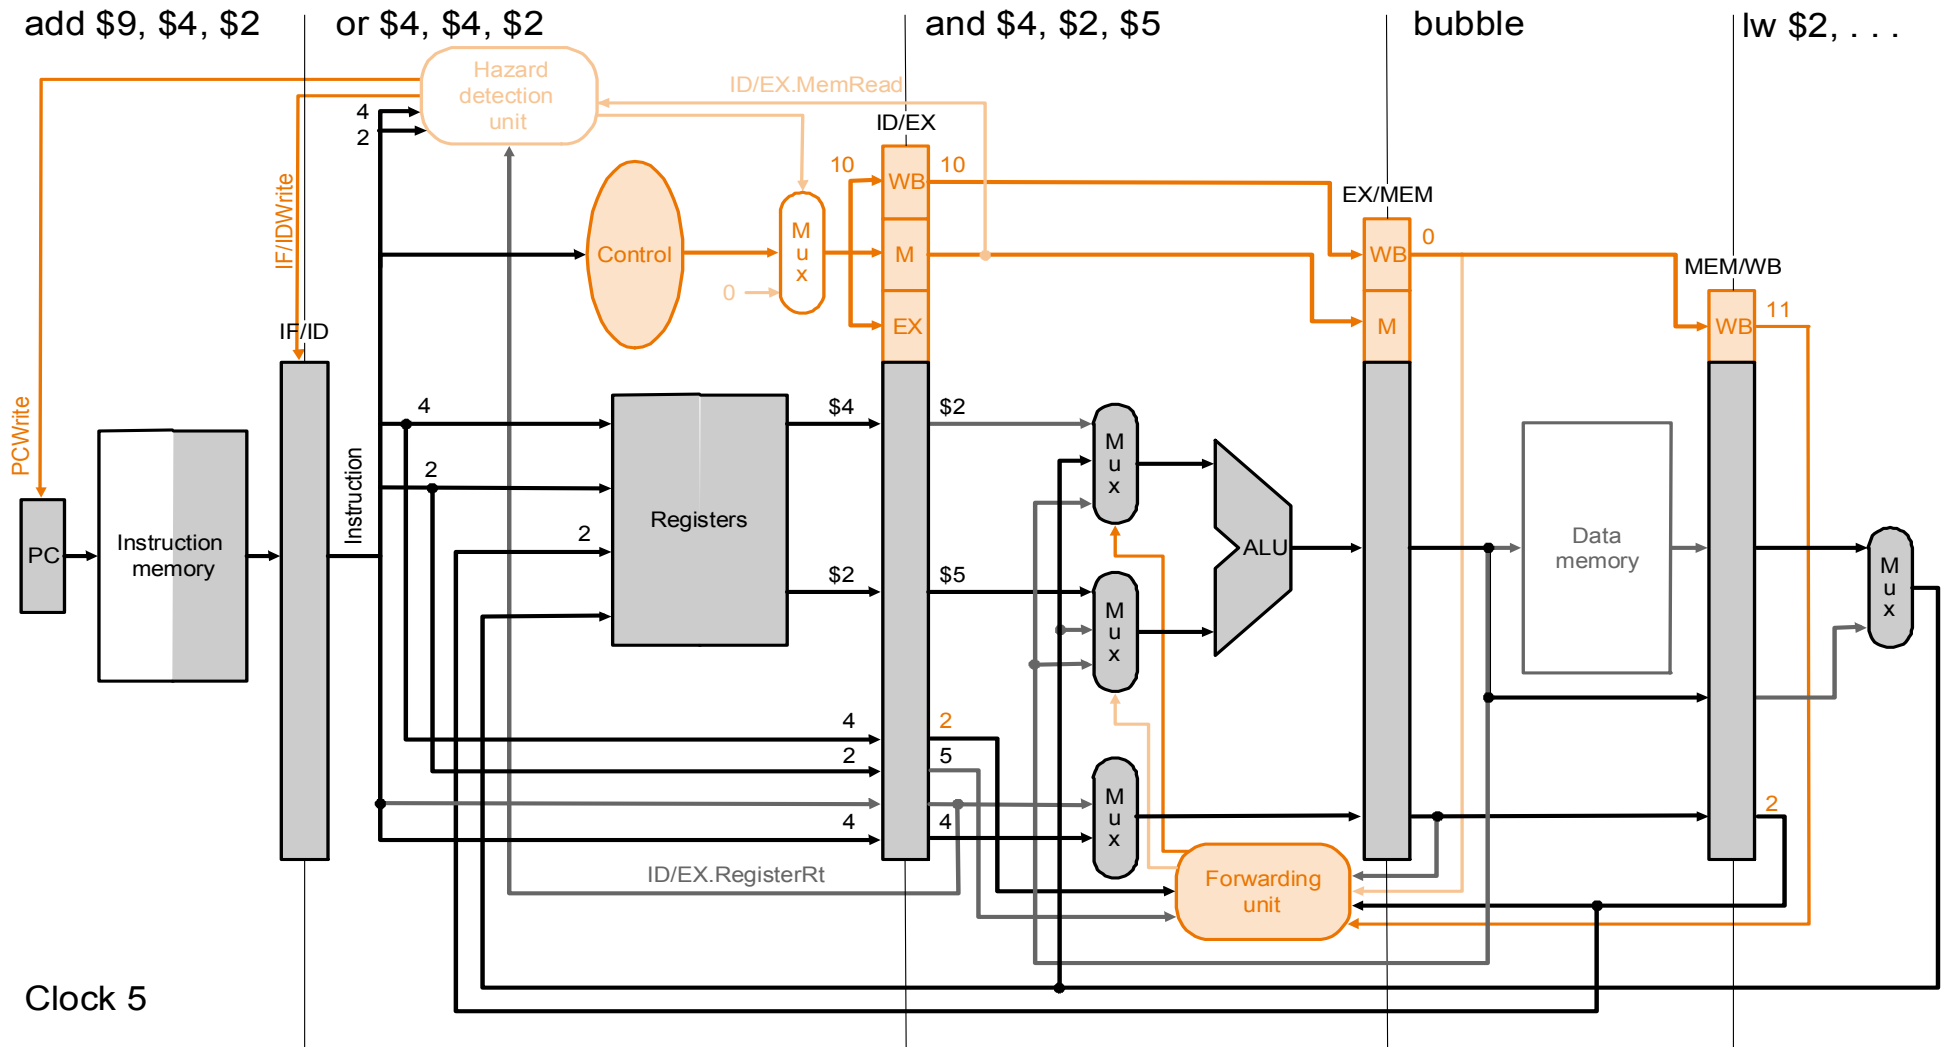
Cycle 4

lw \$2, 20(\$1)
and \$4, \$2, \$5
or \$4, \$4, \$2
add \$9, \$4, \$2



Cycle 5

```
lw $2, 20($1)
and $4, $2, $5
or $4, $4, $2
add $9, $4, $2
```




```
lw    $2, 20($1)
and   $4, $2, $5
or    $4, $4, $2
add   $9, $4, $2
```



Cycle 7

```
lw    $2, 20($1)
and    $4, $2, $5
or     $4, $4, $2
add    $9, $4, $2
```

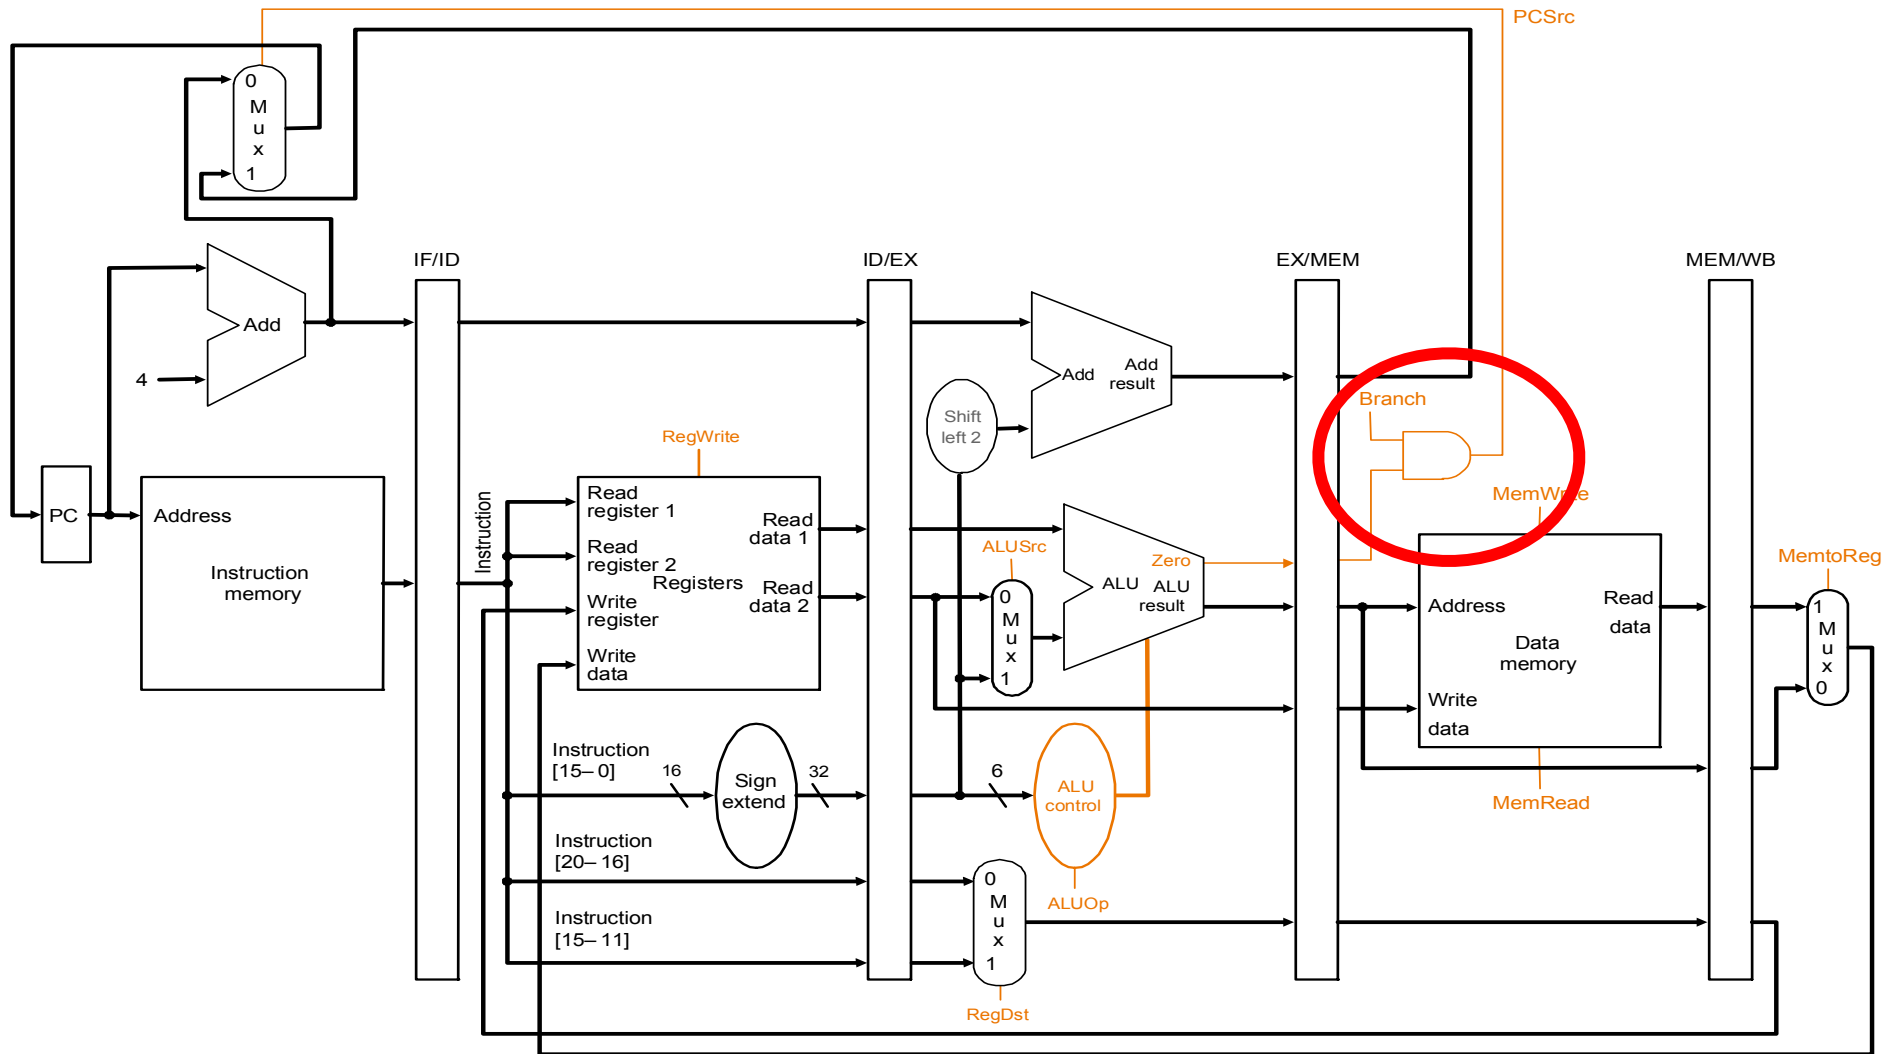


Clock 7

Outline

- ♦ A pipelined datapath
- ♦ Pipelined control
- ♦ Data hazards and forwarding
- ♦ Data hazards and stalls
- ♦ **Branch hazards**

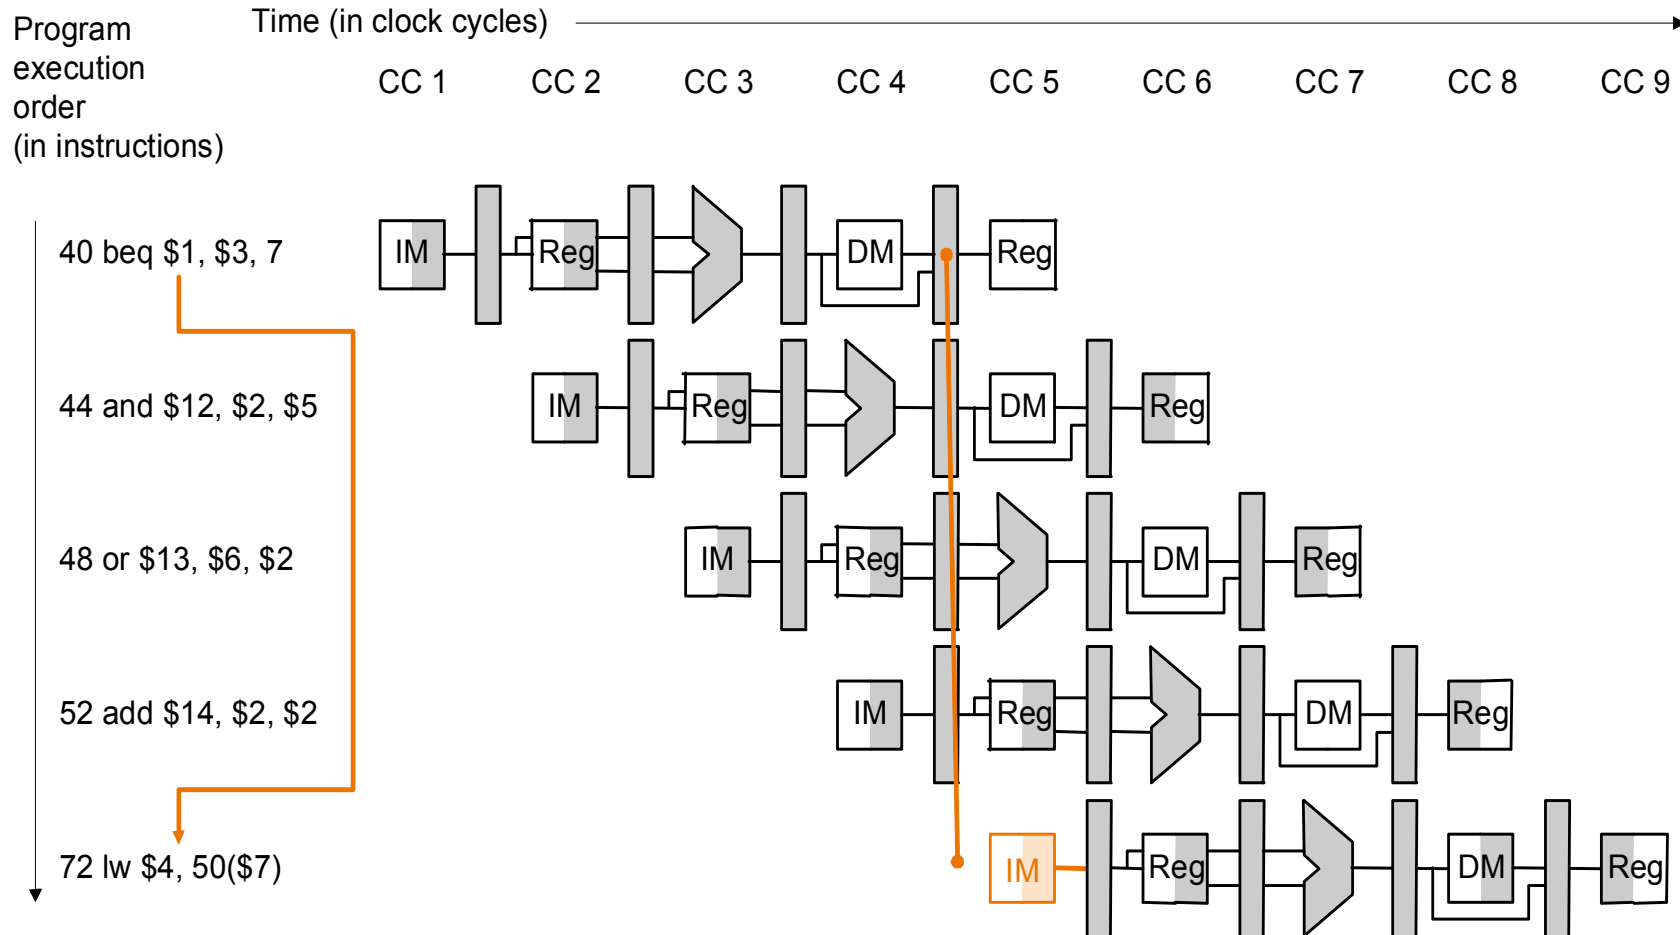
Pipeline Datapath with Control Signals



Control (or Branch) Hazards

- ♦ Problem with branches in the pipeline we have so far is that the branch decision is not made **till the MEM stage** – so what instructions, if at all, should we insert into the pipeline following the branch instructions?
- ♦ Possible solution: *stall* the pipeline till branch decision is known
 - not efficient, slow the pipeline significantly!
- ♦ Another solution: *predict* the branch outcome
 - e.g., always predict *branch-not-taken* – *continue with next sequential instructions*
 - if the prediction is wrong, have to **flush** the pipeline **behind** the branch
 - **Need to flush 3 instructions (too many to degrade the pipeline performance)**

Predicting Branch-not-taken



The outcome of branch taken (prediction wrong) is decided only when `beq` is in the MEM stage, so the following three sequential instructions already in the pipeline have to be flushed and execution resumes at `lw`

Flushing on Misprediction

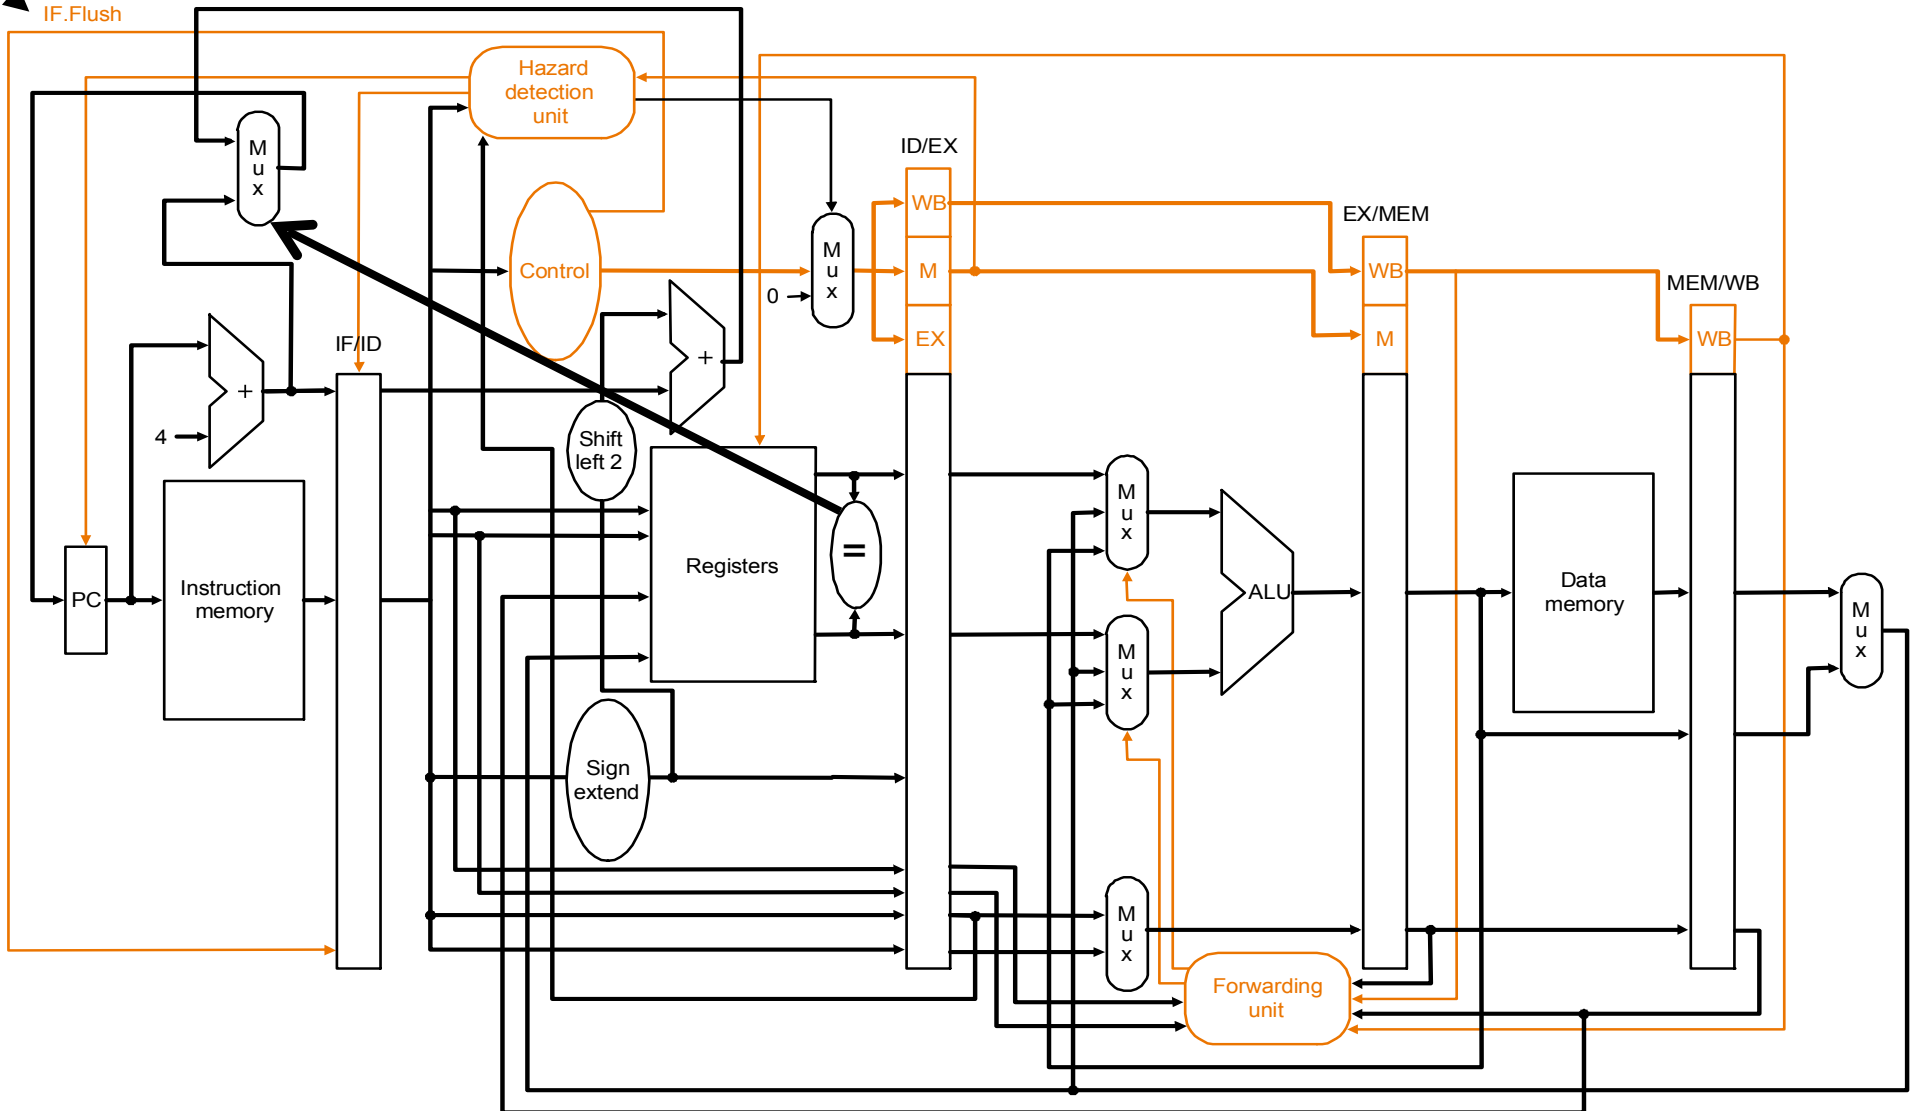
- ♦ Similar to the strategy as for stalling on load-use data hazard (RAW) ...
- ♦ Flush: Zero out all the control values (or the instruction itself) in pipeline registers for the instructions following the branch that are already in the pipeline
- ♦ If branch decision made **in the ID stage (discussed later)**, we have to flush only one instruction in the IF stage – the branch delay penalty is then only one clock cycle

Optimizing the Pipeline to Reduce Branch Delay

- ♦ *Move the branch decision from the MEM stage (as in our current pipeline) earlier to the ID stage*
 - **calculating the branch target address** involves moving the branch adder from the MEM stage to the ID stage – **inputs to this adder, the PC value and the immediate fields are already available in the IF/ID pipeline register**
 - **calculating the branch decision** is efficiently done, e.g., for equality test, by XORing respective bits and then ORing all the results and inverting, rather than using the ALU to subtract and then test for zero (when there is a carry delay)
 - **we must correspondingly make additions to the forwarding and hazard detection units** to forward to or stall the branch at the ID stage in case the branch decision depends on an earlier result

Optimized Datapath for Branch

IF.Flush control zeros out the instruction (which follows the branch) in the IF/ID pipeline register



Branch decision is moved from the MEM stage to the ID stage – simplified drawing not showing enhancements to the forwarding and hazard detection units


Pipelined Branch (predict not taken)

```

36 sub $10, $4, $8
40 beq $1, $3, 7
44 and $12, $2, $5
48 or $13, $2, $6
52 add $14, $4, $2
56 slt $15, $6, $7
...
72 lw $4, 50($7)

```

Optimized pipeline with
only one bubble as a result
of the taken branch



Summary

- ♦ **Pipelined processor design**
 - An overview of pipelining
 - A pipelined datapath
 - Pipelined control
- ♦ **Problems with pipelined processor**
 - Data hazards and forwarding
 - Data hazards and stalls
 - Branch hazards